



Introduction to  
Simultaneous  
Localization and Mapping

Spring Semester 2019

CS284

Course Convener:

Prof Laurent Kneip

lkneip@ shanghaitech.edu.cn

# Disclaimer



- Some of the lecturing material is naturally taken from publically available online material from other groups. Sources include:
  - Autonomous Systems Lab ETH Zurich
  - Research School of Engineering, ANU
  - Informatik Department, Uni Freiburg
  - MIT
  - ...
- By using the present material you confirm that
  - You use it only internally and for the purpose of education
  - You are aware that the lecture material may originate from other sources, even if explicit reference is occasionally missing

# Price question

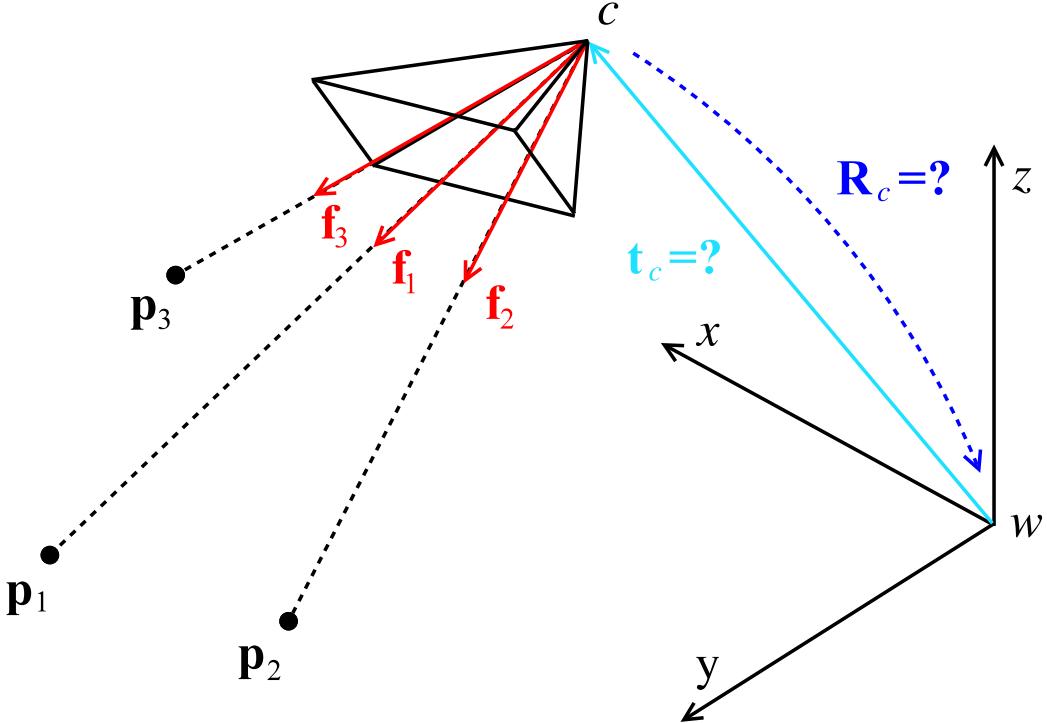
- Example: Absolute Pose

- Objective:

$$\operatorname{argmax}_{\mathbf{t}_c, \mathbf{R}_c} \sum_i \mathbf{f}_i^t \cdot (\mathbf{R}_c^T \cdot (\mathbf{p}_i - \mathbf{t}_c))$$

- What's a better objective?

$$\operatorname{argmax}_{\mathbf{t}_c, \mathbf{R}_c} \sum_i \mathbf{f}_i^t \cdot \frac{\mathbf{R}_c^T \cdot (\mathbf{p}_i - \mathbf{t}_c)}{\|\mathbf{R}_c^T \cdot (\mathbf{p}_i - \mathbf{t}_c)\|}$$



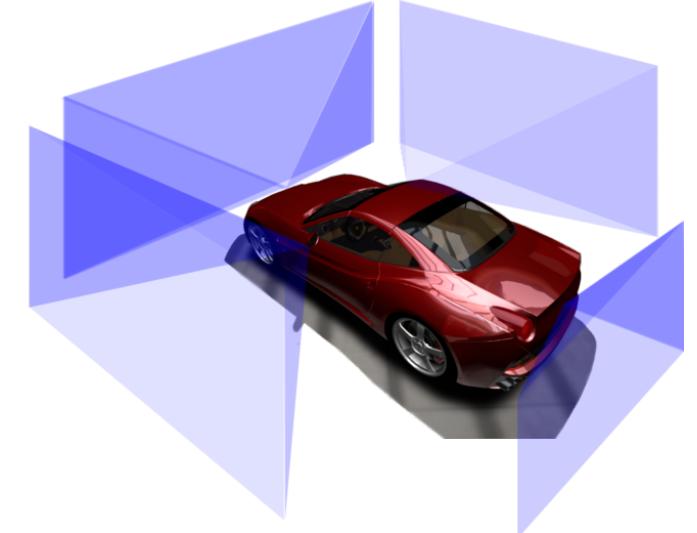
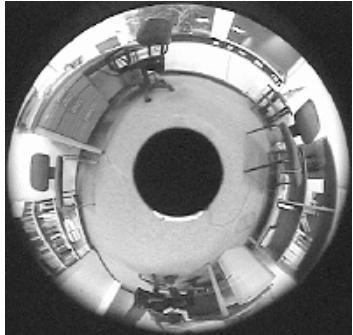
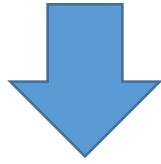
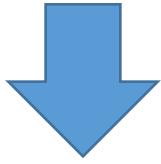
# Where are we at?



- Today: Front-end choices, SLAM as a graphical optimization problem

Week	Tuesday course	Thursday course
18th February	Introduction (what is SLAM, curriculum, course structure (homeworks, projects), course webpage)	Homogeneous coordinates, Transformations, etc.
25th February	Filtering methods (GF, PF)	Filtering methods (GF, PF)
4th March	Pose-only SLAM, Graph SLAM	Introduction to the projects
11th March	Camera as a sensor, camera calibration	Introduction into Visual SLAM, feature extraction, tracking, and matching
18th March	Feature extraction, tracking, and matching	Case study: MonoSLAM (why filter?)
25th March	Bootstrapping: The Relative Pose Problem	Full visual odometry pipeline, Non-linear optimization, Bundle adjustment
1st April	Non-linear optimization, Bundle adjustment	Place recognition, loop closure
8th April	Place recognition, loop closure, the absolute pose problem	Case study: ORB SLAM
15th April	Dense Tracking and Mapping (DTAM), Photometric vs geometric errors, semi-dense opt.	RGBD SLAM, Iterative closest points (Laser-point cloud registration), TSDF
22nd April	SLAM with Stereo and Multi-Perspective cameras	Midterm exam
29th April	Visual-Inertial SLAM	Project discussions/buffer
6th May	Dynamic and Multi-body SLAM	Project discussions/buffer
13th May	Semantic SLAM	Project Presentations

# SLAM: different sensors

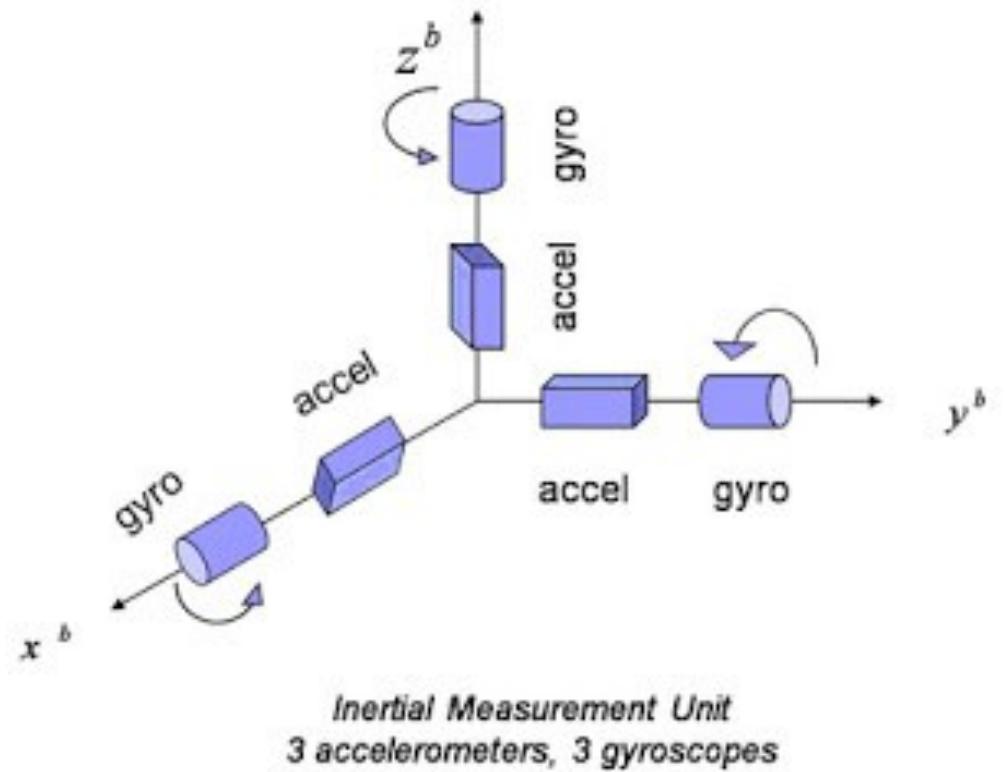
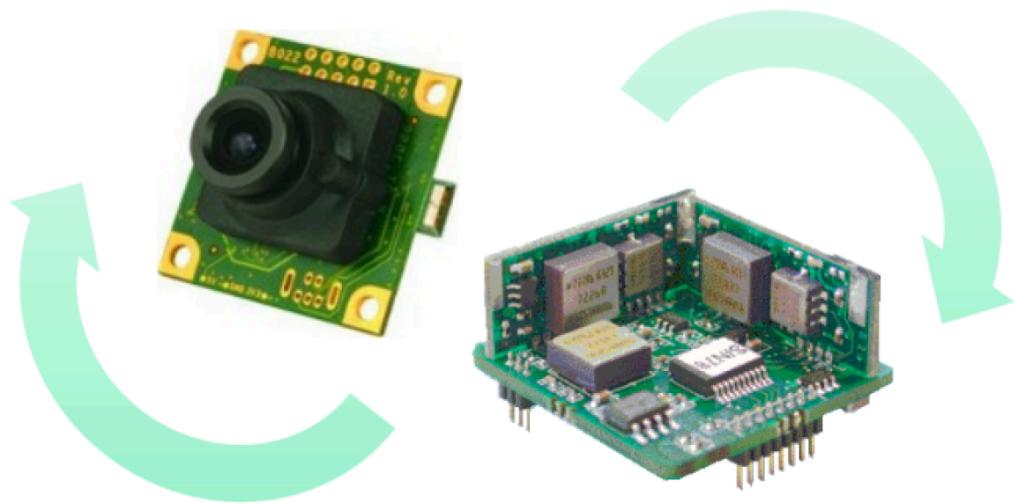


# SLAM: different sensors

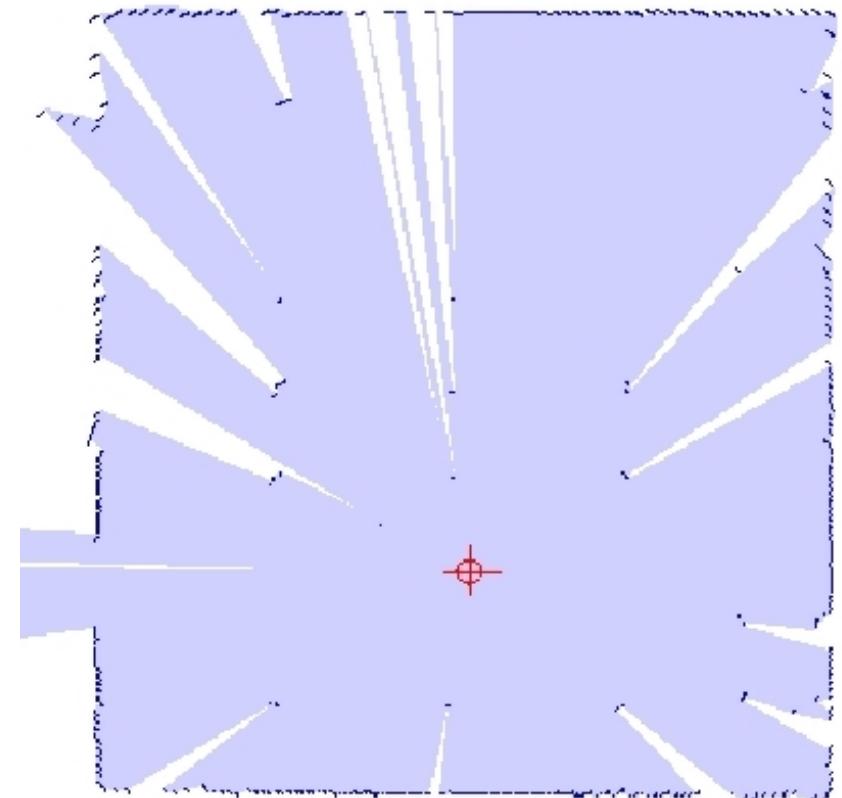
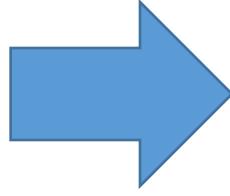


# SLAM: different sensors

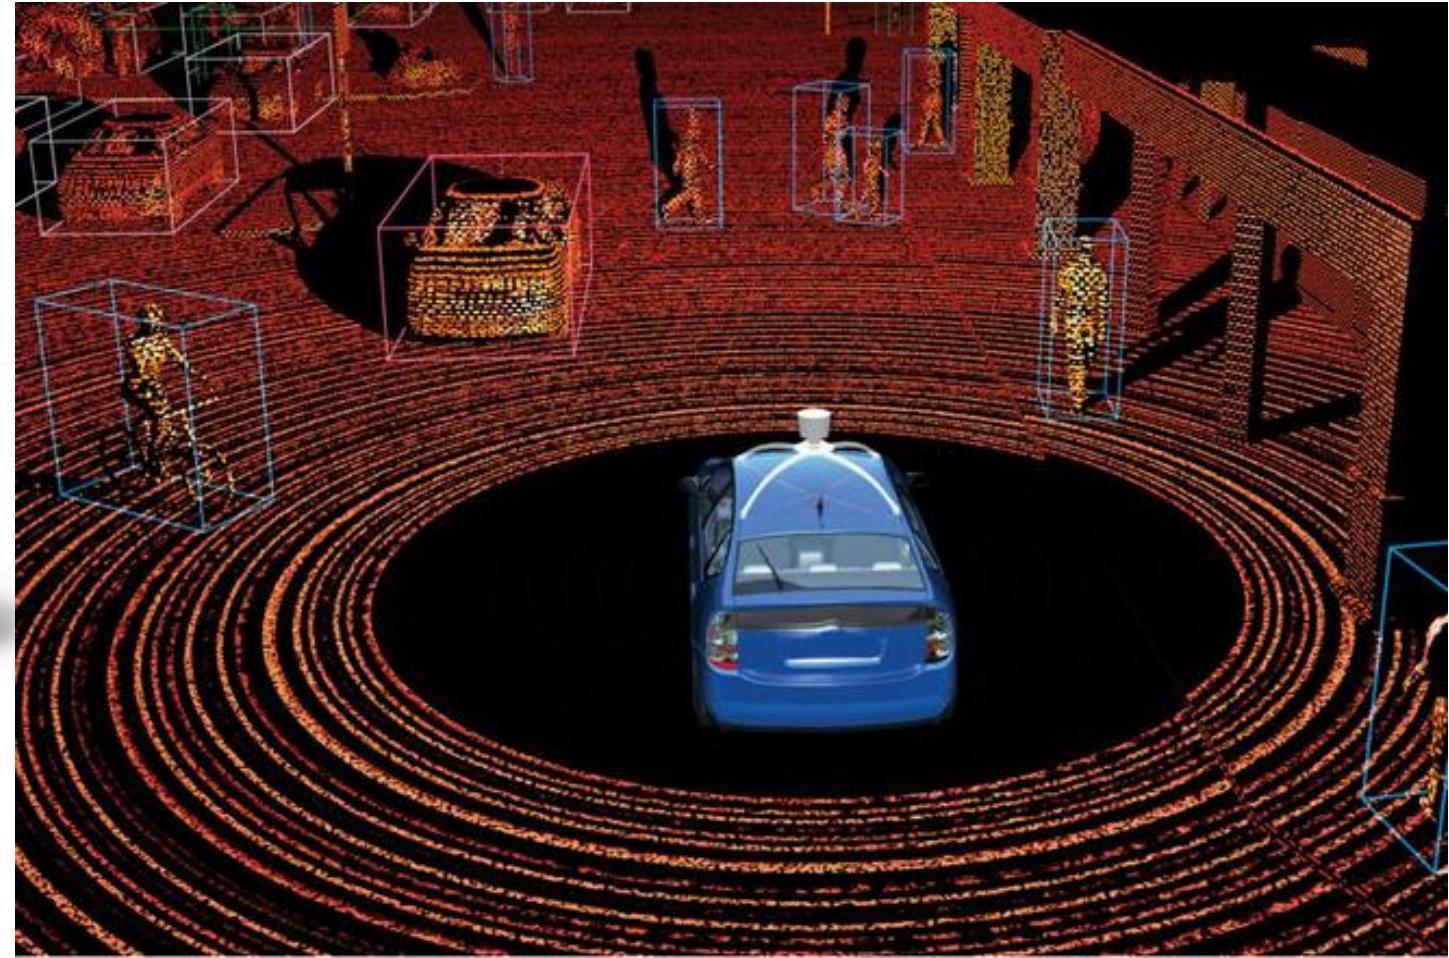
- Inertial sensors



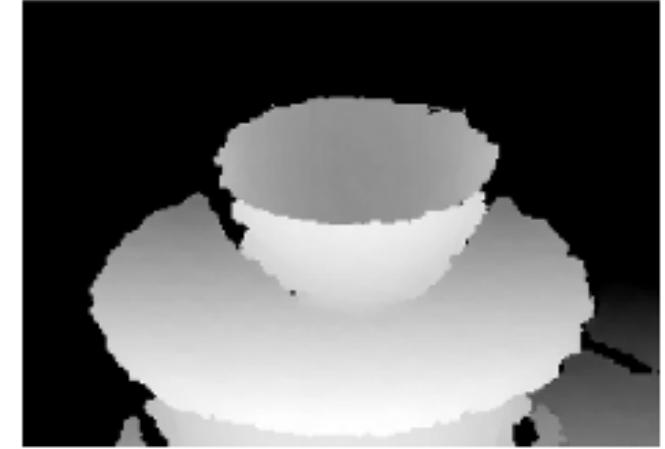
# SLAM: different sensors



# SLAM: different sensors



# SLAM: different sensors



**KINECT**  
for XBOX 360.

# Sensor classes



- Two main types of exteroceptive sensors used in robotics:
  - **Photometric:**  
energy reflection (regular images), energy emission
  - **Geometric:**  
position, surface, volume, etc.

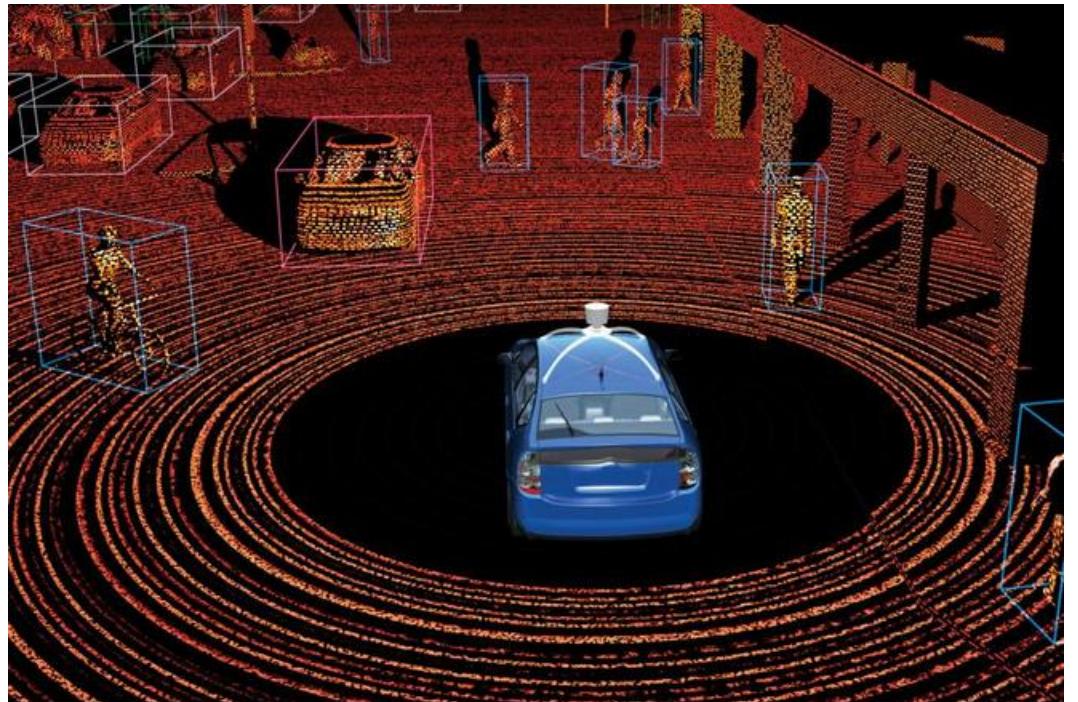
# Sensor classes

- Photometric (visual) sensors
  - Measures photometric effects
  - Allows us to extract features
    - 2D-2D, 2D-3D, or 3D-3D correspondences
  - Allows us to map those features to 3D points (explicitly)
    - Represented either as nodes in a graph, or in a state vector
    - Full graph representations depend on it



# Sensor classes

- Geometric sensors (Laser scanners)
  - Typically: no feature/point correspondences!
    - No appearance information (or at least rarely used)
    - Only few, unstable local invariant keypoints (geometric!)



## Photometric sensor

Pros

- Low cost/weight/energy sensors;
- Rich textures;
- Primary problem: pose estimation

## Geometric sensors

- Direct depth measurements;
- Robust to environment changes;
- Primary problem: map representation

## Photometric sensor

Pros

- Low cost/weight/energy sensors;
- Rich textures;
- Primary problem: pose estimation

## Geometric sensors

- Direct depth measurements;
- Robust to environment changes;
- Primary problem: map representation

Cons

- No depth; scale unobservability
- Dependent on external light source;

- Heavy/expensive sensors;
- Systematic noise;

# Some sensors make things easy



- Laser scanners
  - 2D or 3D
  - directly sense depth!
- Cheap alternative
  - RGB-D camera
- Let's start with those examples!

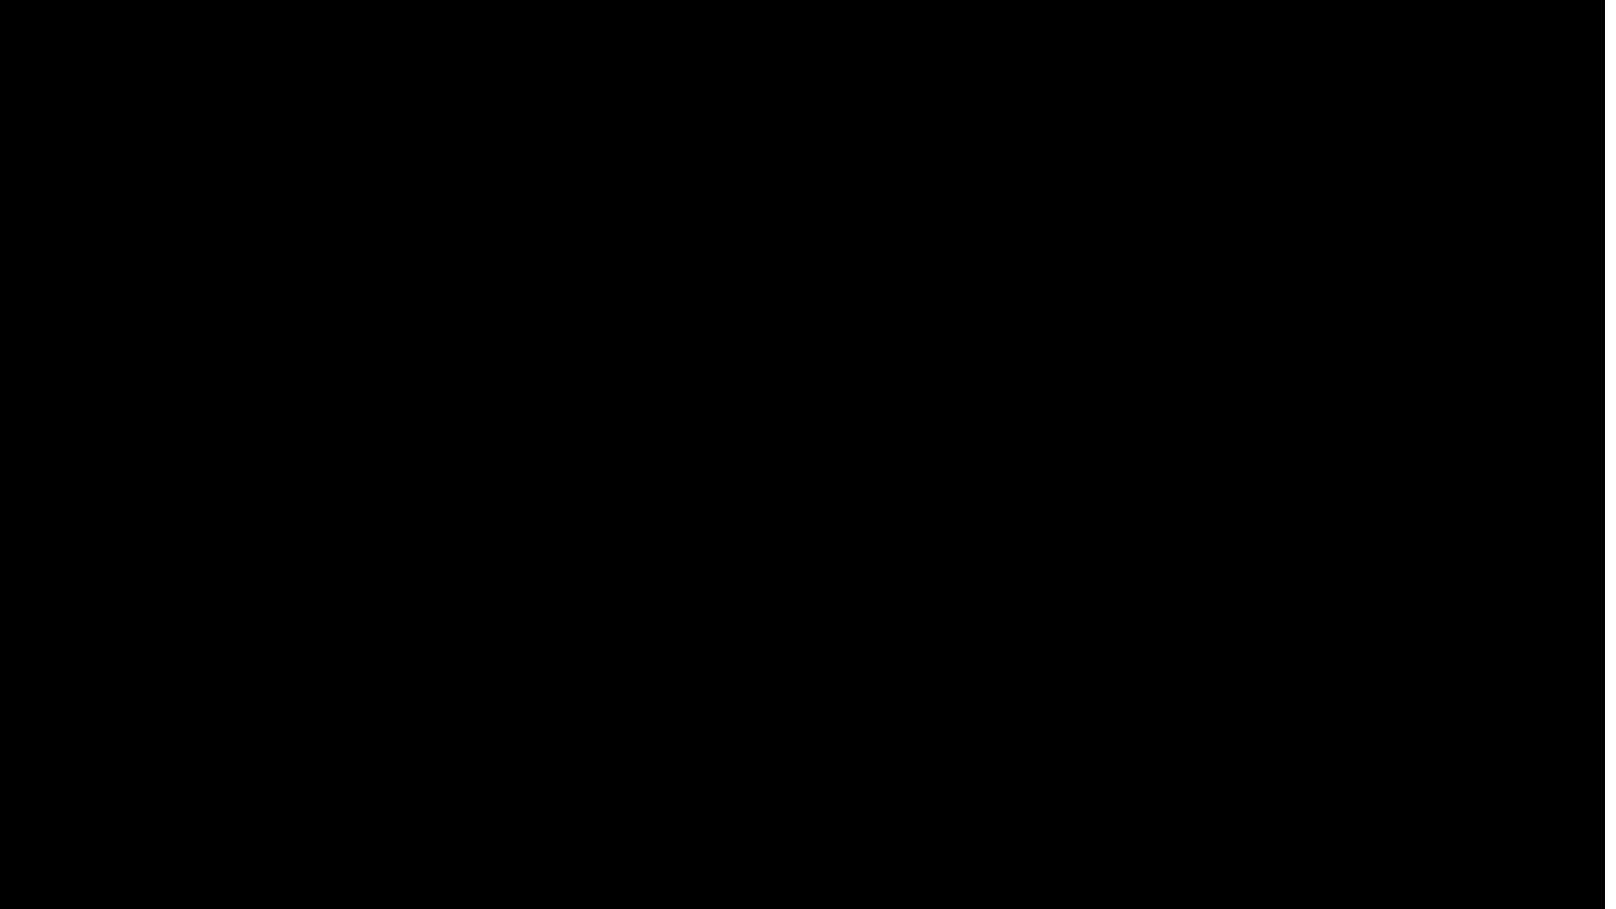


# Example

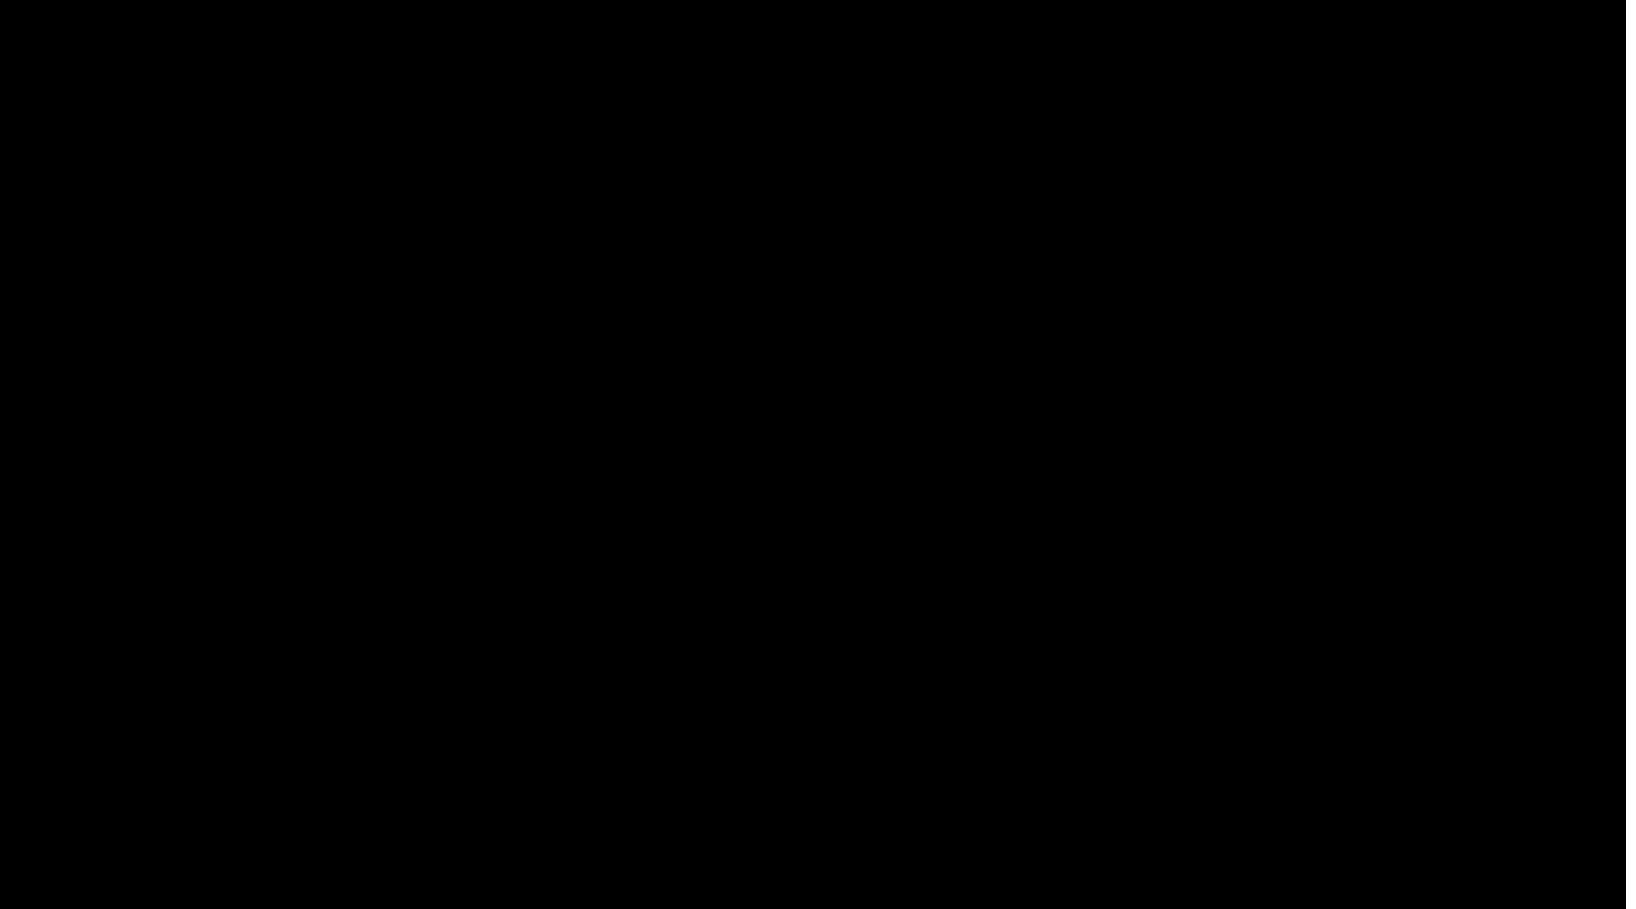
- Robot providing odometry
- Equipped with
  - Planar Laser scanner
  - Pitching planar Laser scanner
  - Velodyne
  - ...



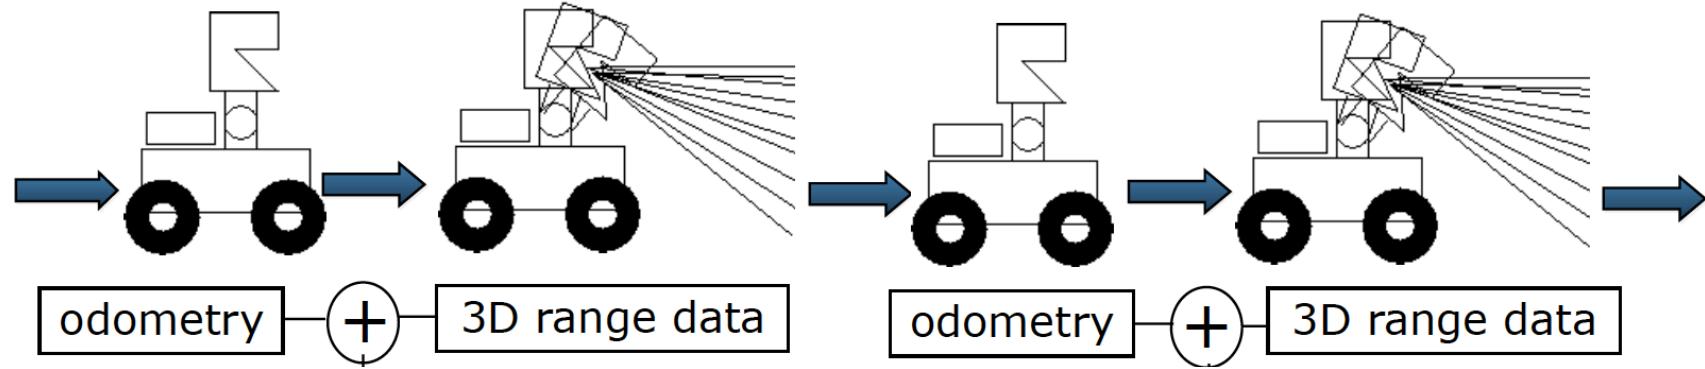
# Example



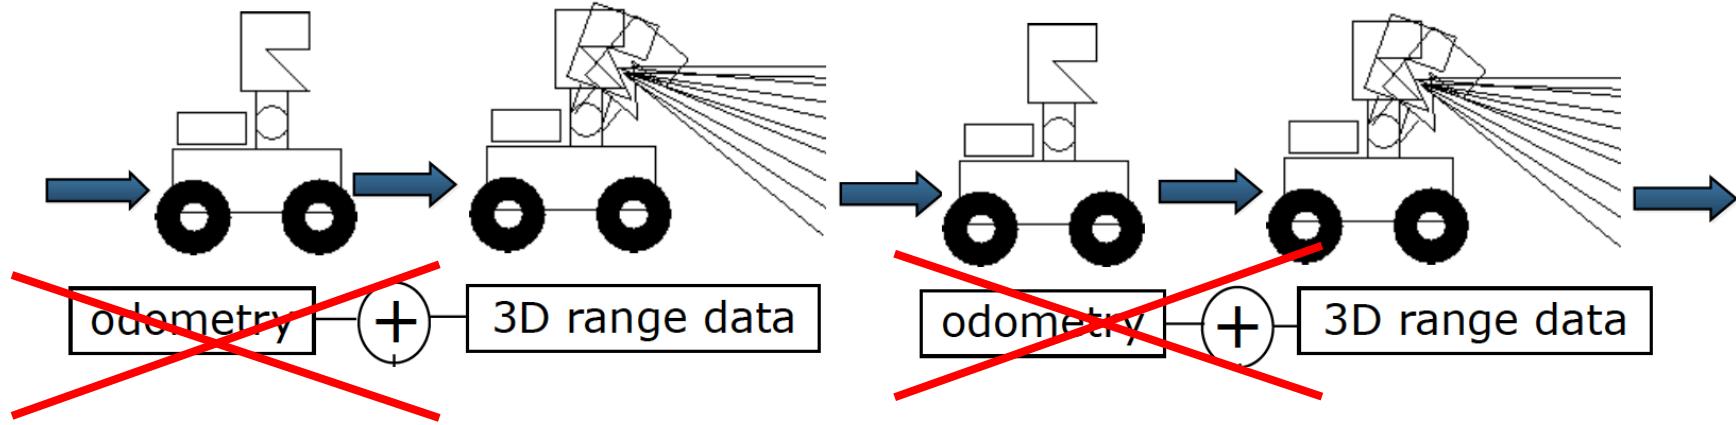
# Also possible in full 3D (or 6D)



# Example

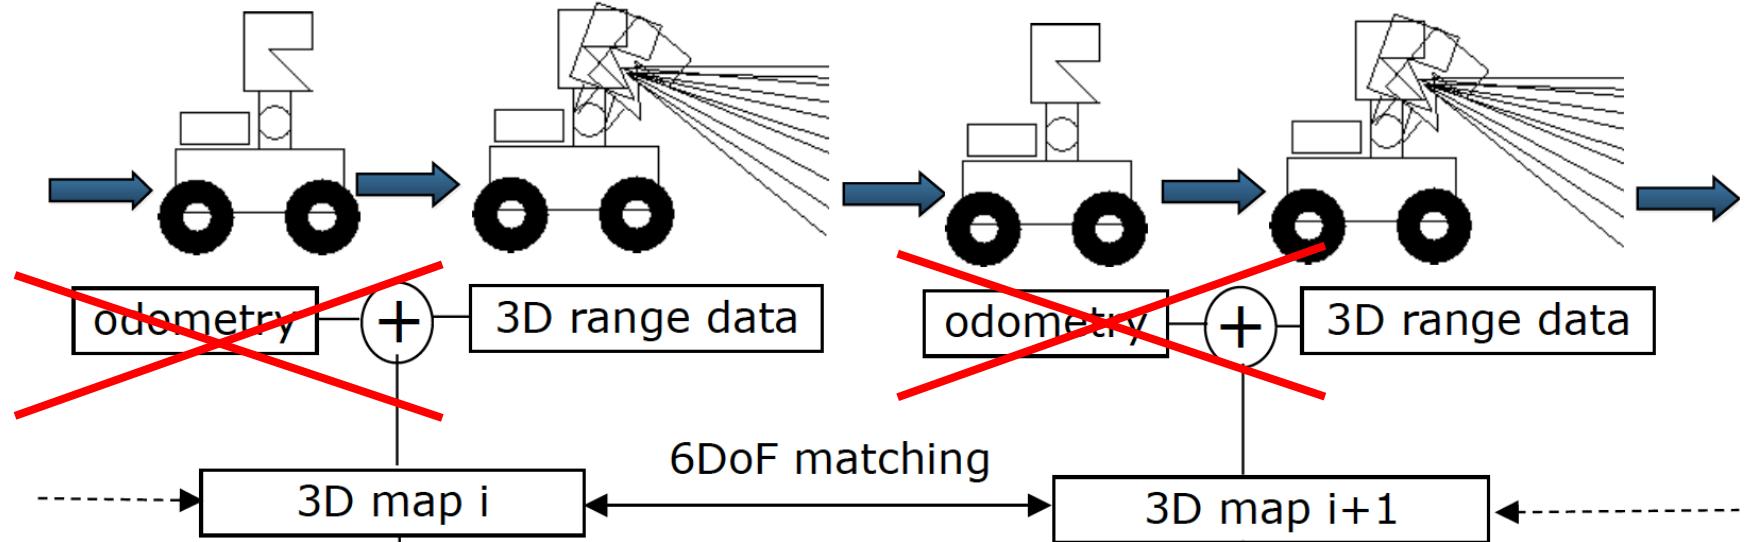


# Example

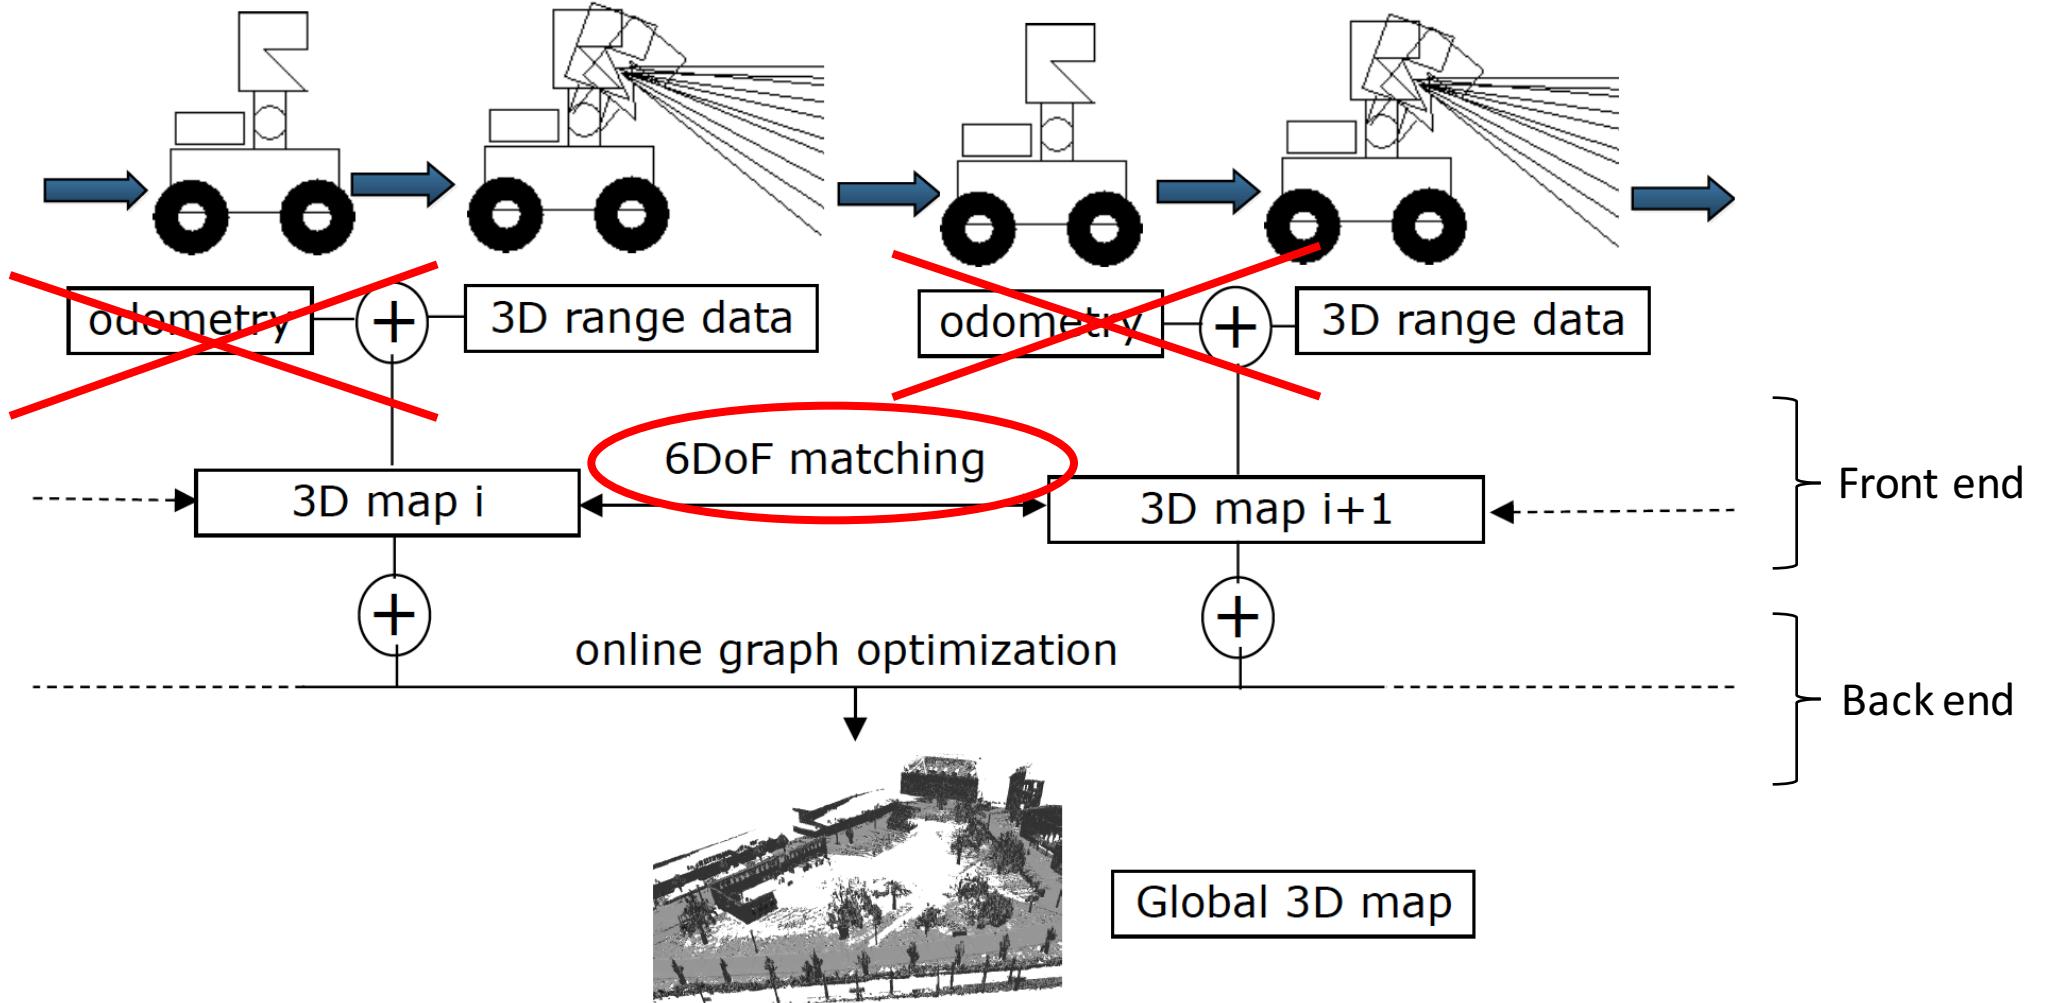


- Odometry is inaccurate
  - Cheap sensors
  - Wheel slippage
  - Difficult terrain

# Example

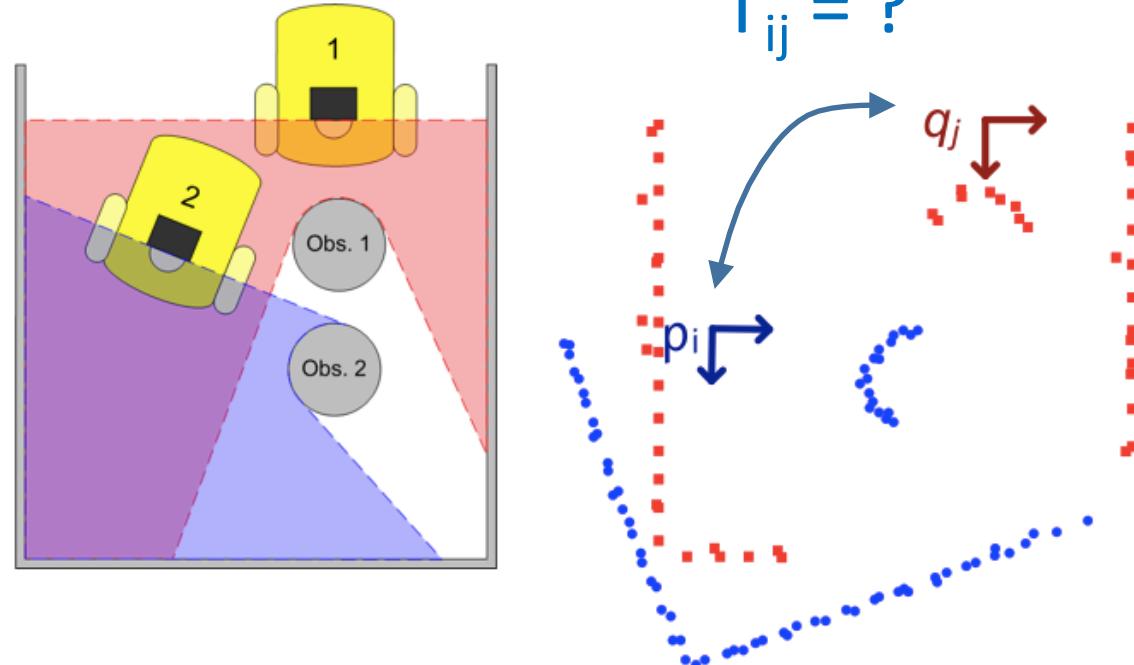


# Example



# SLAM with Lasers

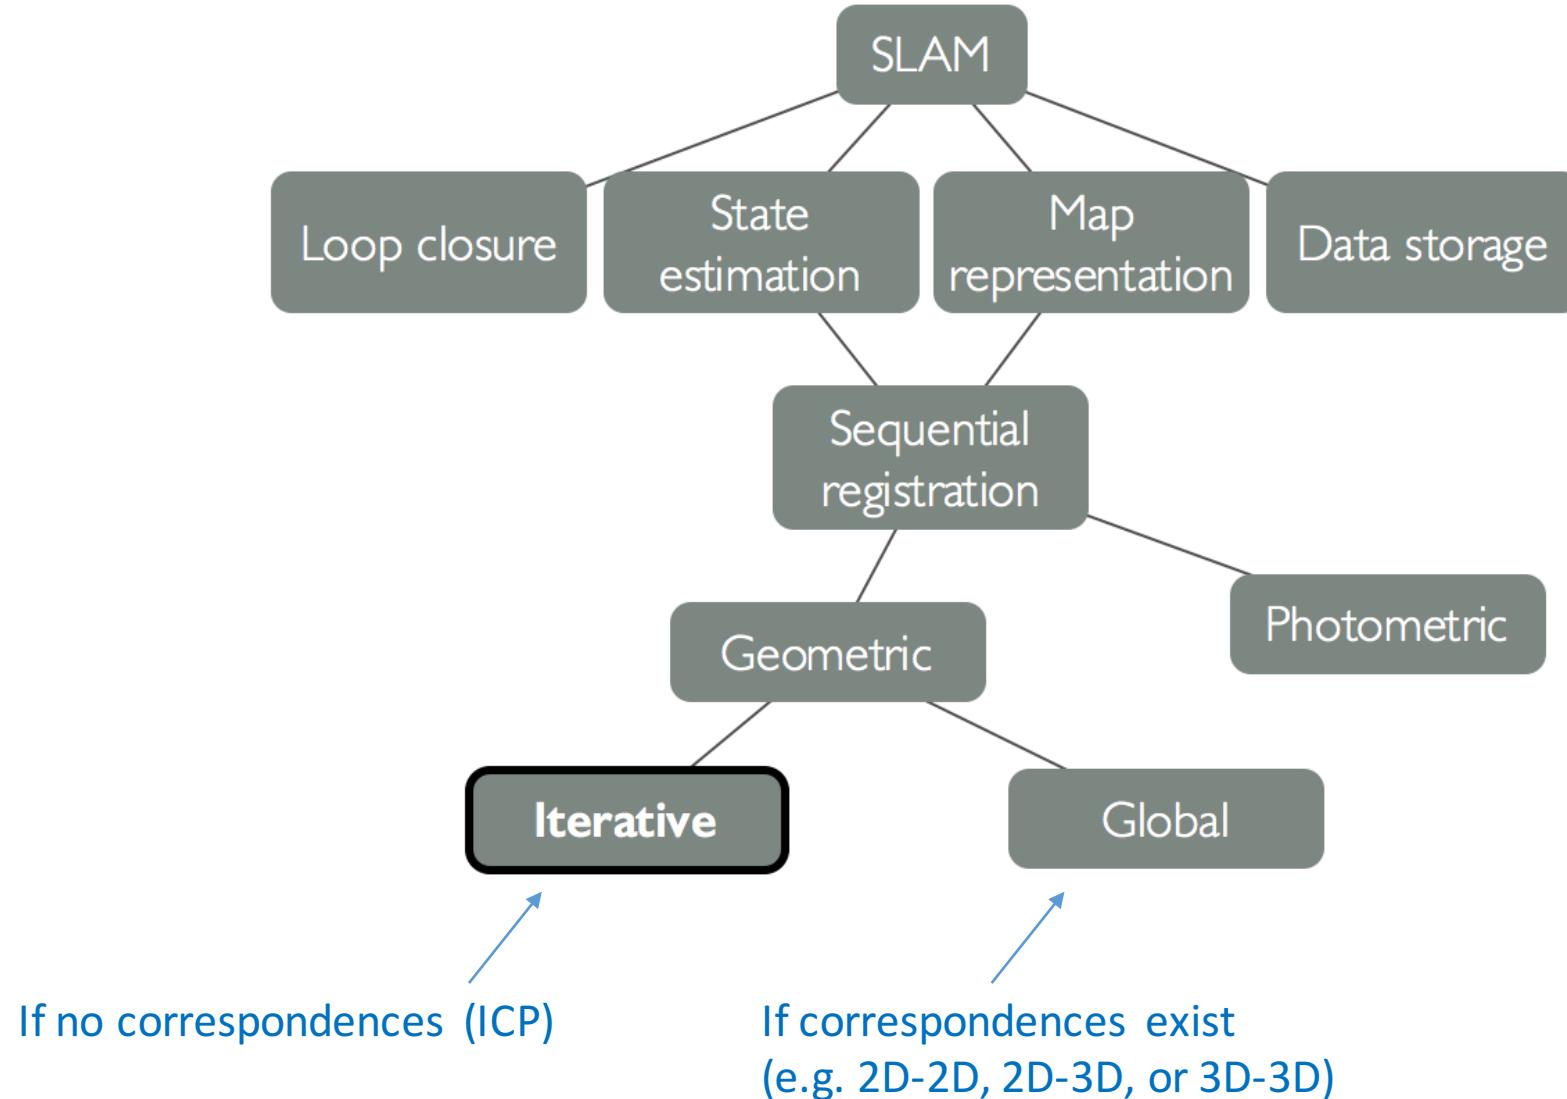
- How to do 6 DoF scan registration for relative localization?



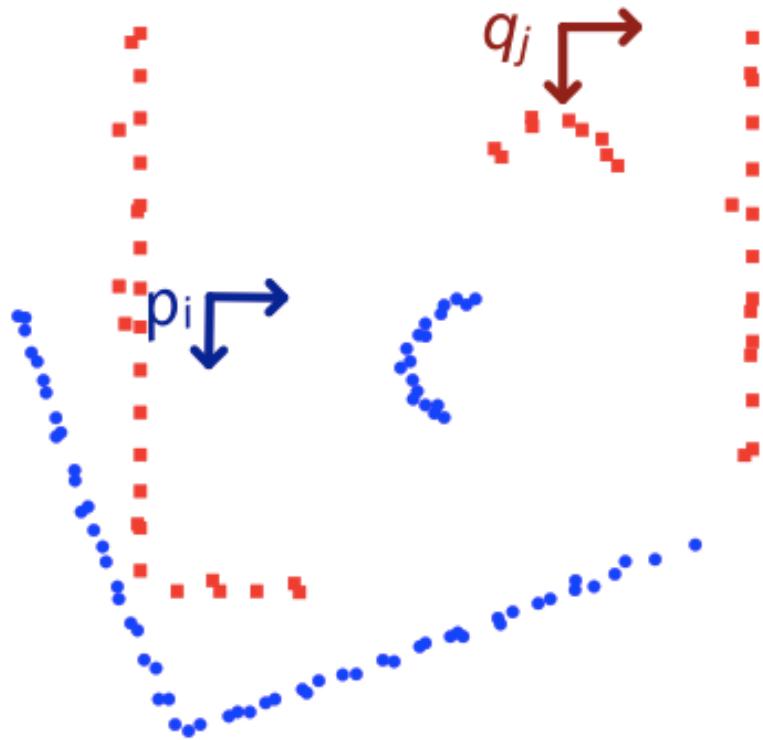
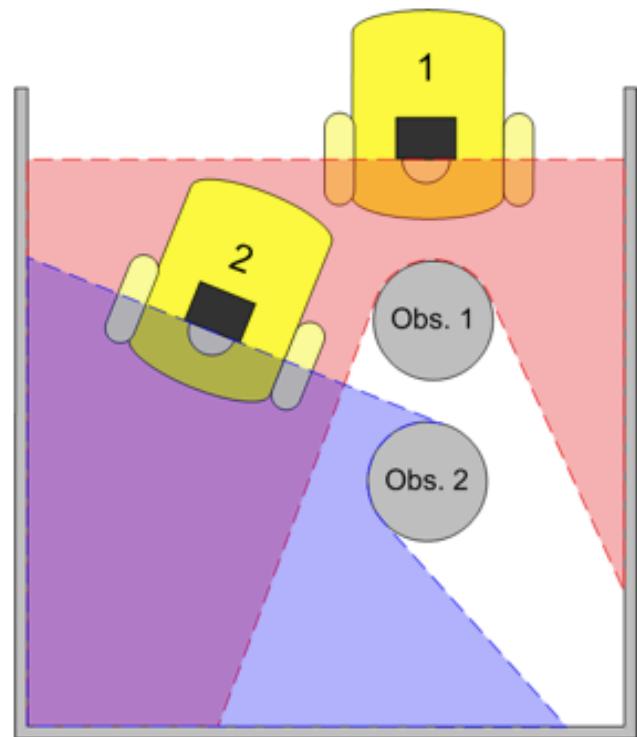


# Iterative Closest Points (ICP)

# Classification of registration methods



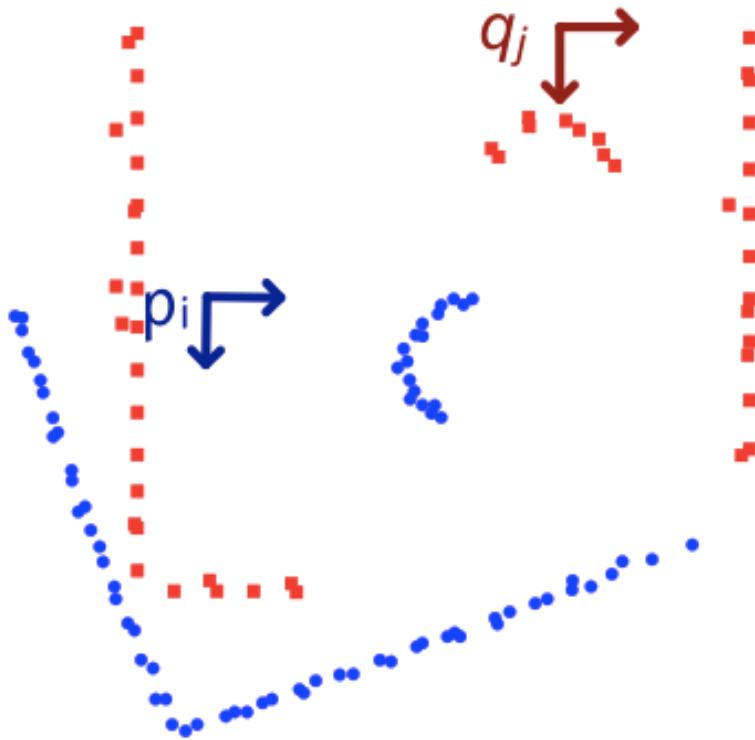
# ICP: Toy example in 2D



# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. Minimization

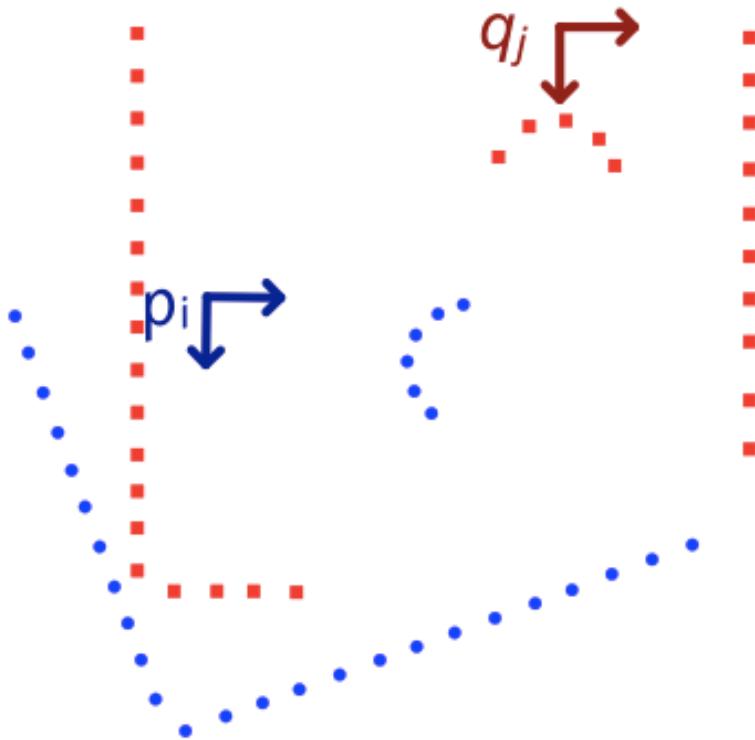
*(Steps defined in Rusinkiewicz 01)*



# ICP: Toy example in 2D



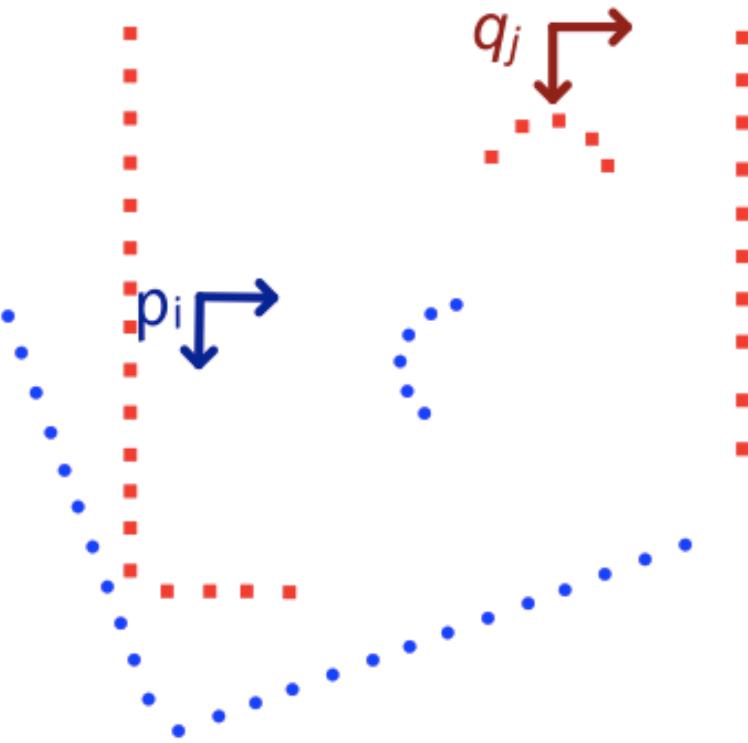
1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. Minimization



# ICP: Toy example in 2D

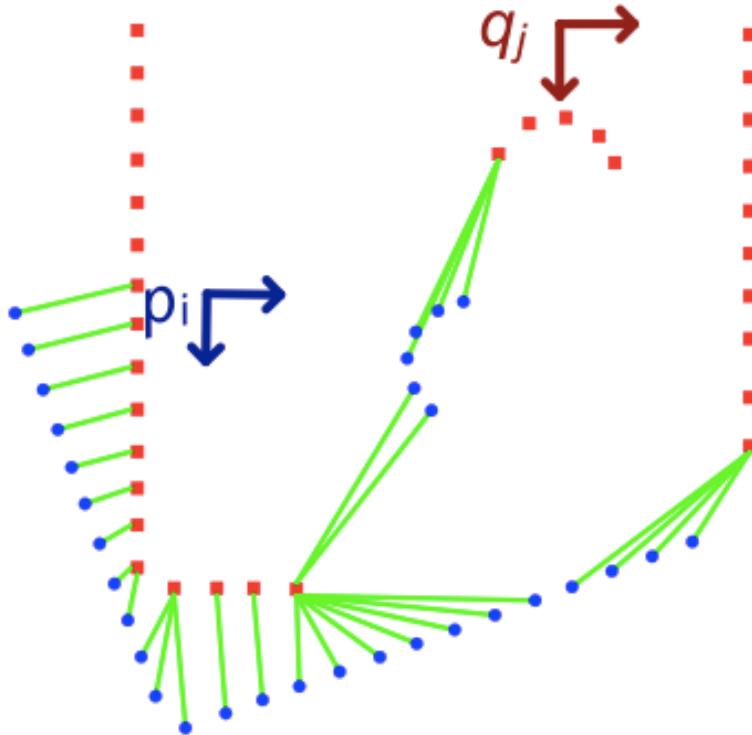


1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. Minimization



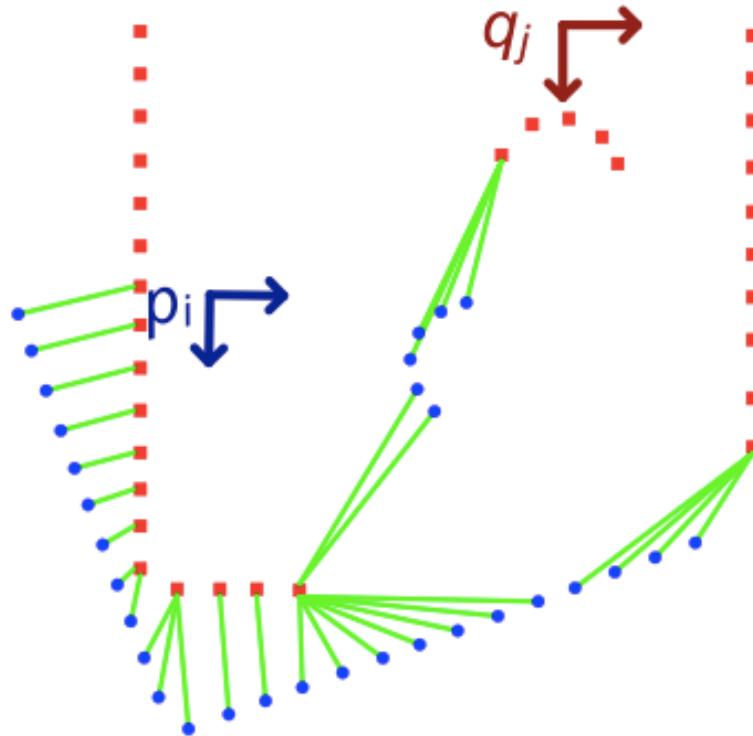
# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. Minimization



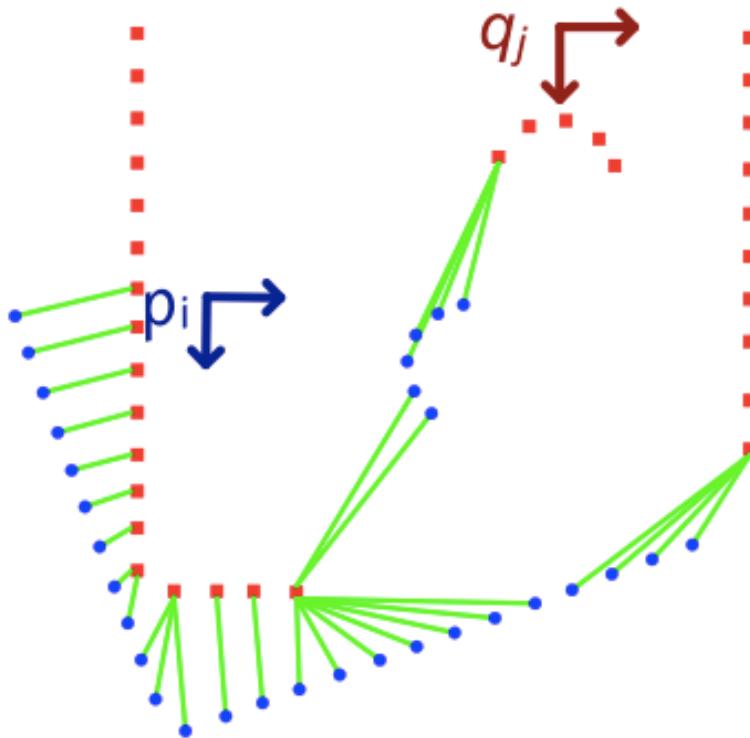
# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. Minimization



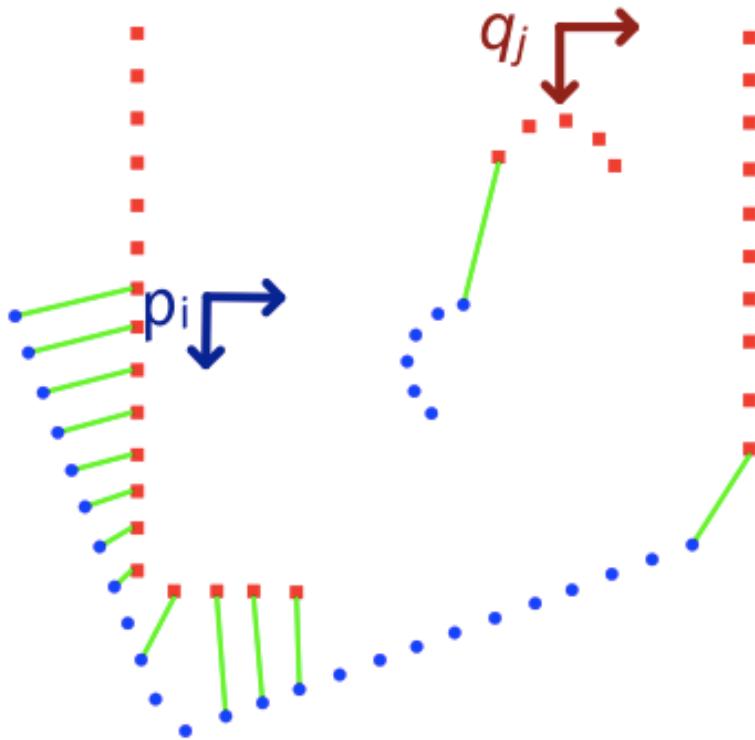
# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. **Rejection**
5. Error
6. Minimization



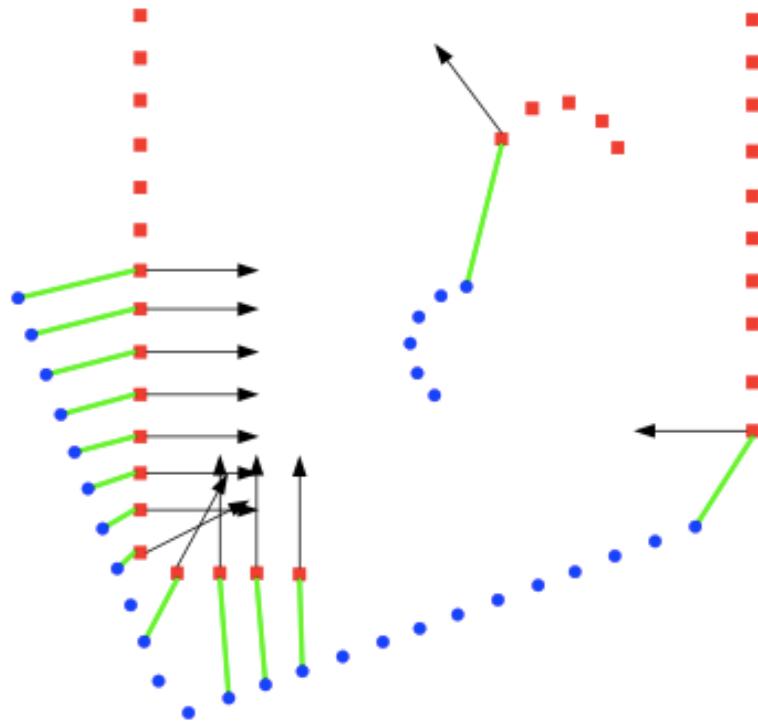
# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. **Rejection**
5. Error
6. Minimization



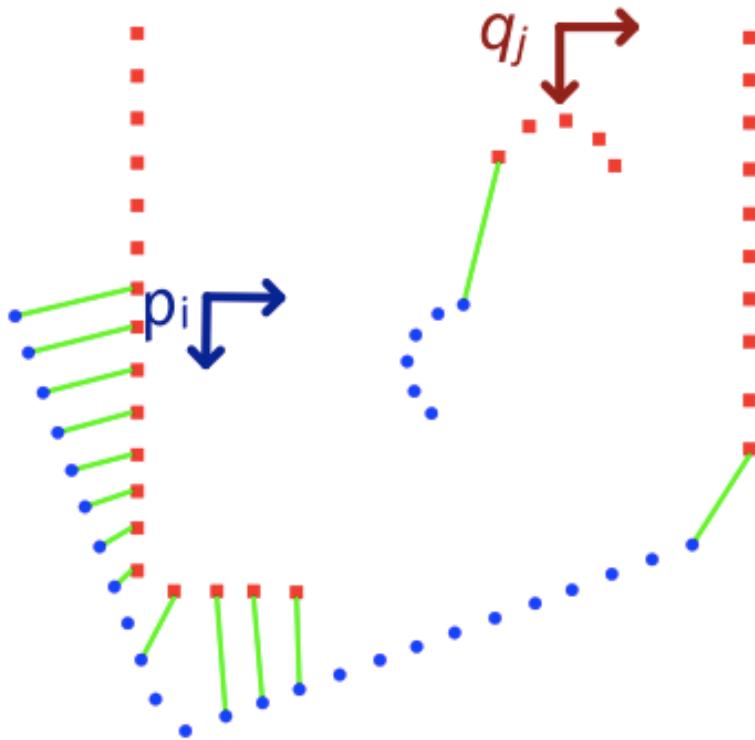
# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. Minimization



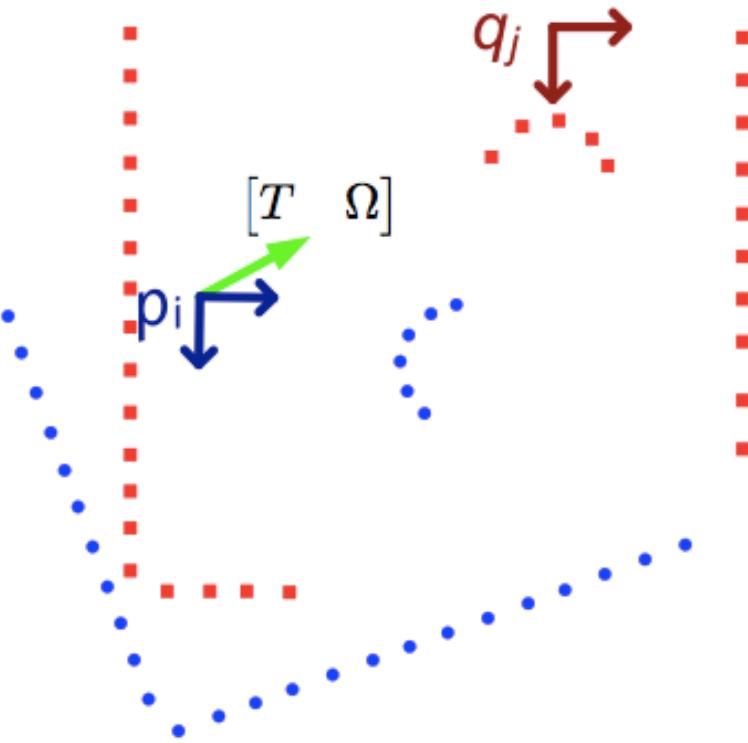
# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. **Minimization**



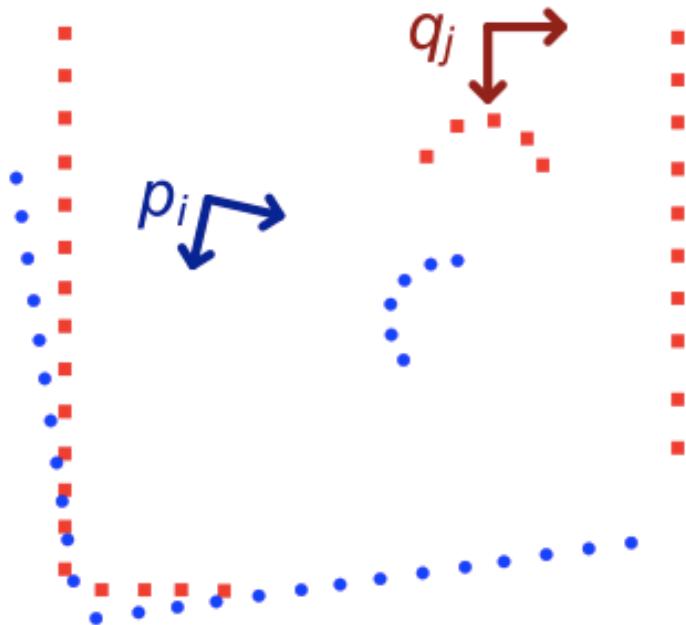
# ICP: Toy example in 2D

1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. **Minimization**



# ICP: Toy example in 2D

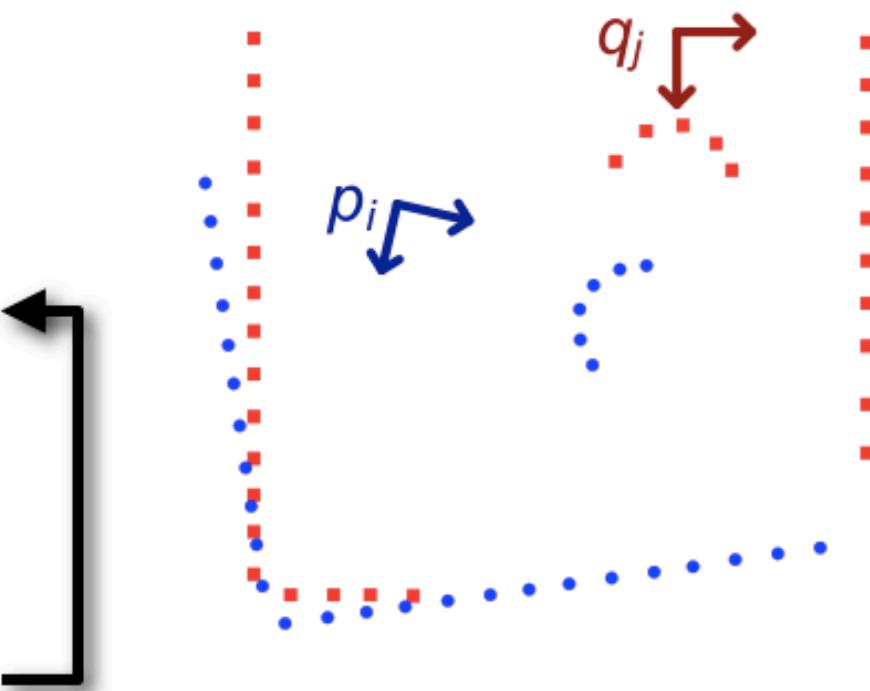
1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. **Minimization**



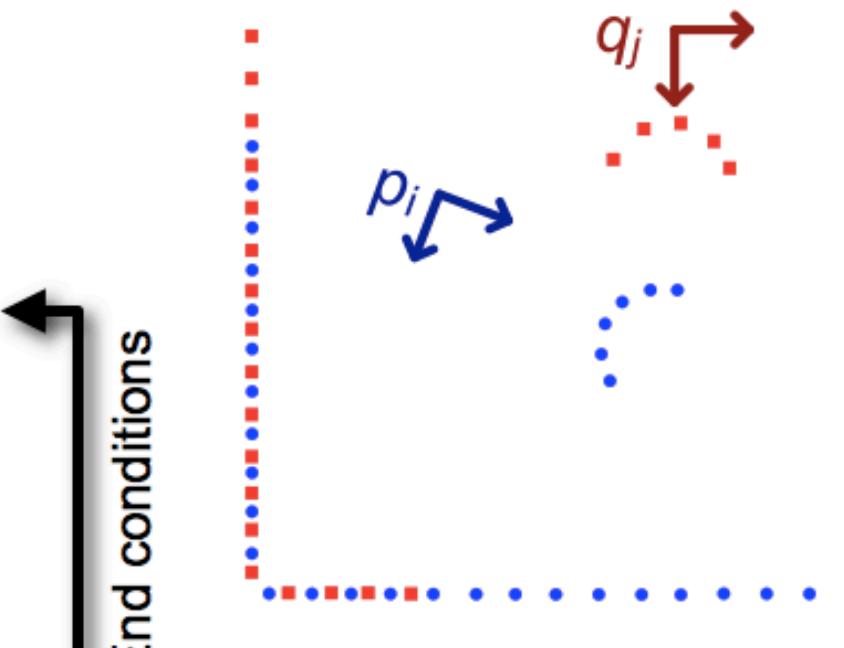
# ICP: Toy example in 2D



1. Preprocessing
2. Matching
3. Weighting
4. Rejection
5. Error
6. Minimization



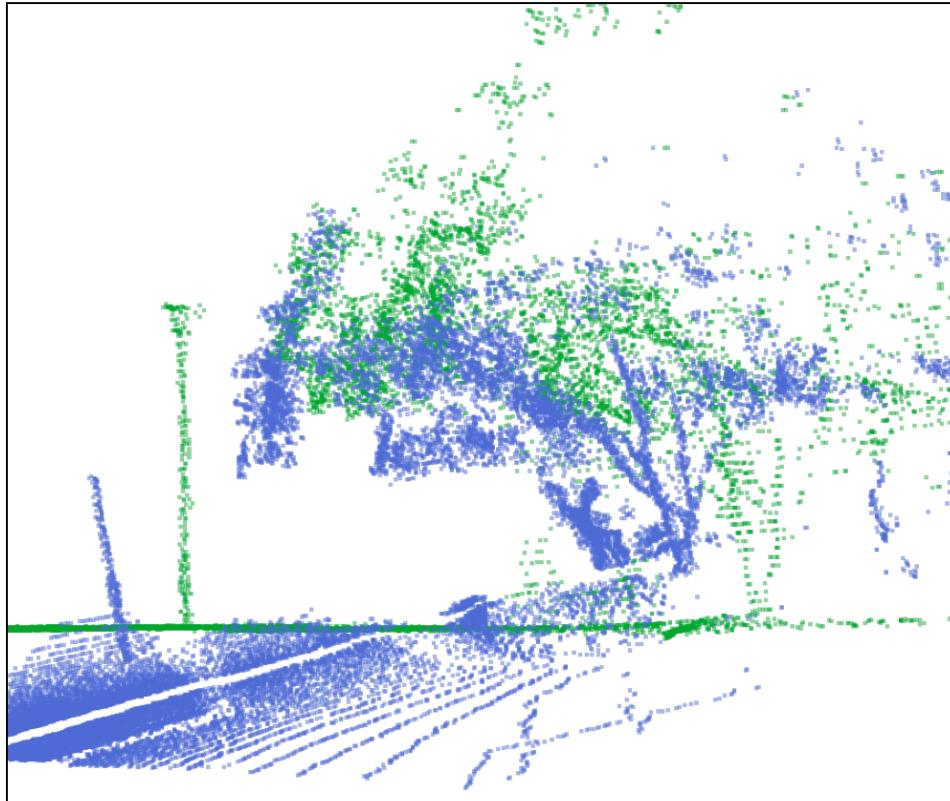
# ICP: Toy example in 2D

1. Preprocessing
  2. Matching
  3. Weighting
  4. Rejection
  5. Error
  6. Minimization
- End conditions
- 

# ICP: Example in 3D

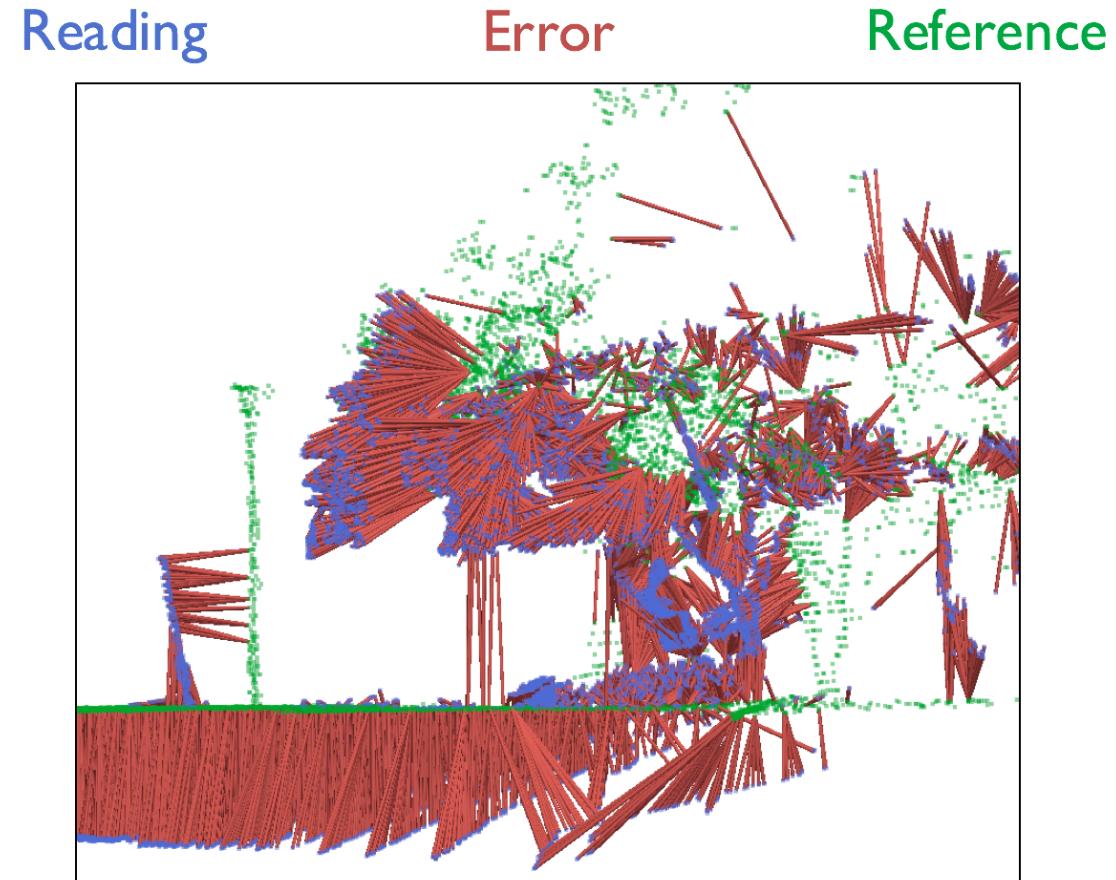


Reading



Reference

# ICP: Example in 3D



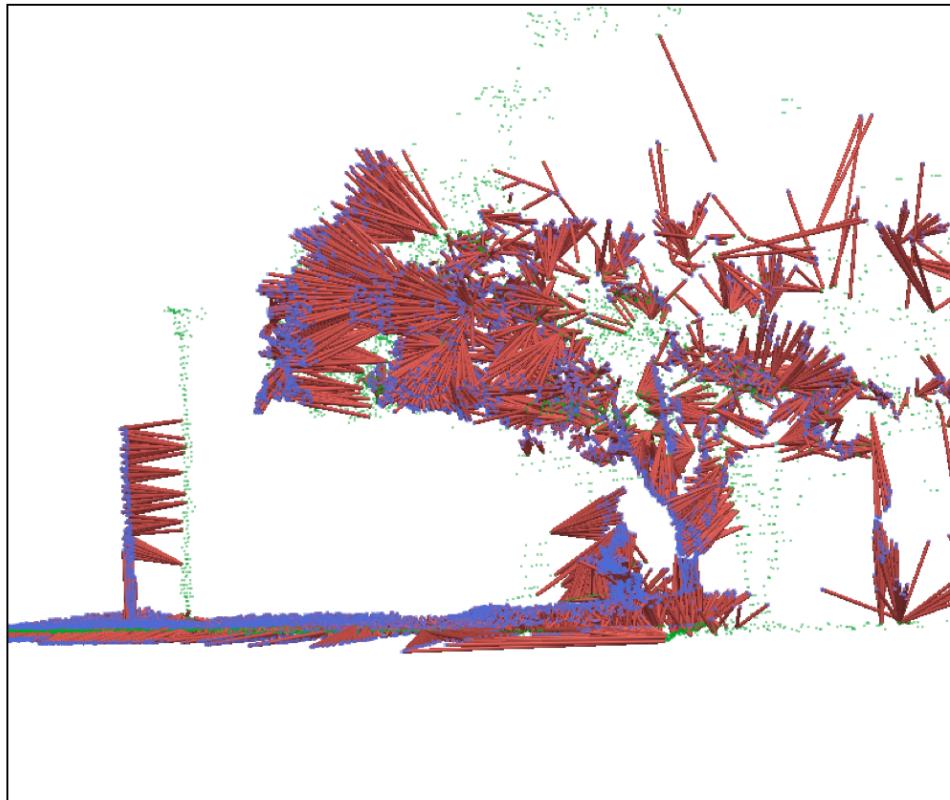
# ICP: Example in 3D



Reading

Error

Reference



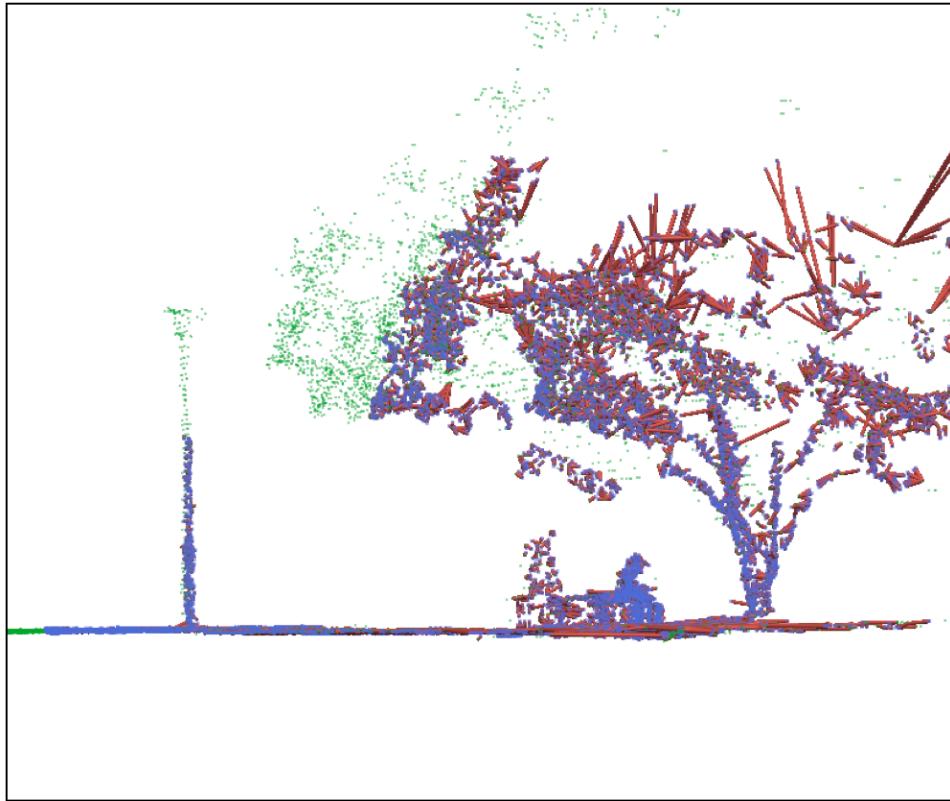
# ICP: Example in 3D



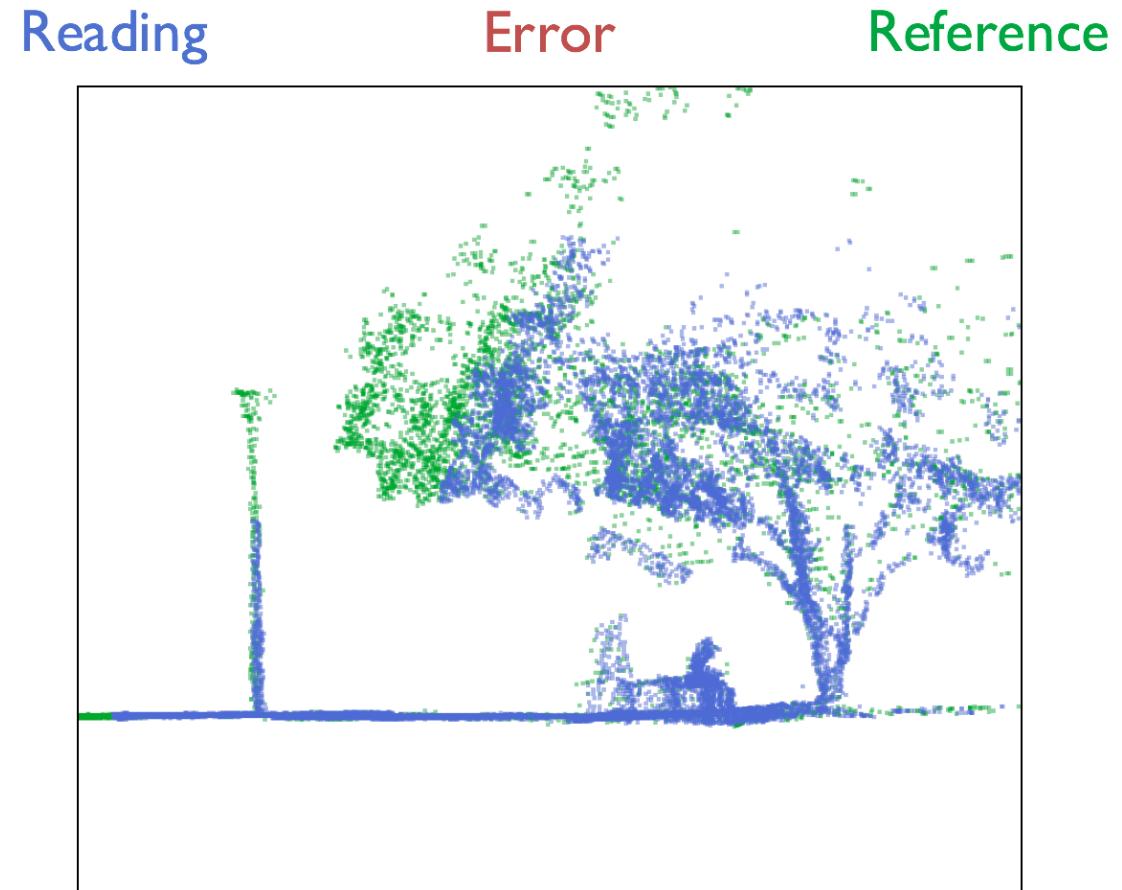
Reading

Error

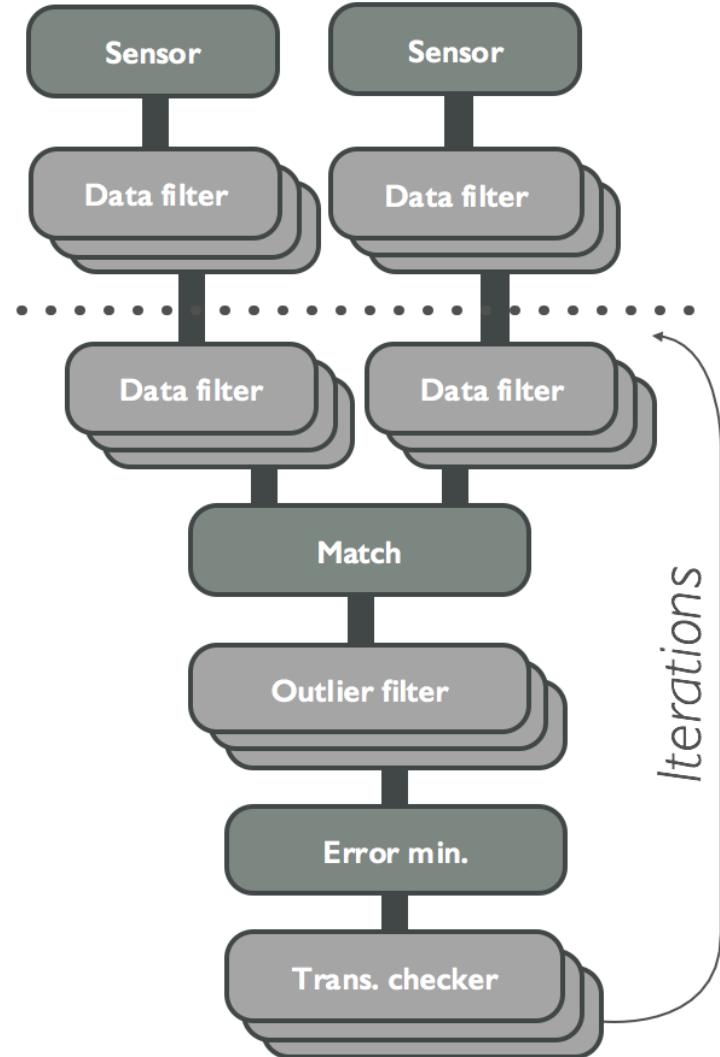
Reference



# ICP: Example in 3D

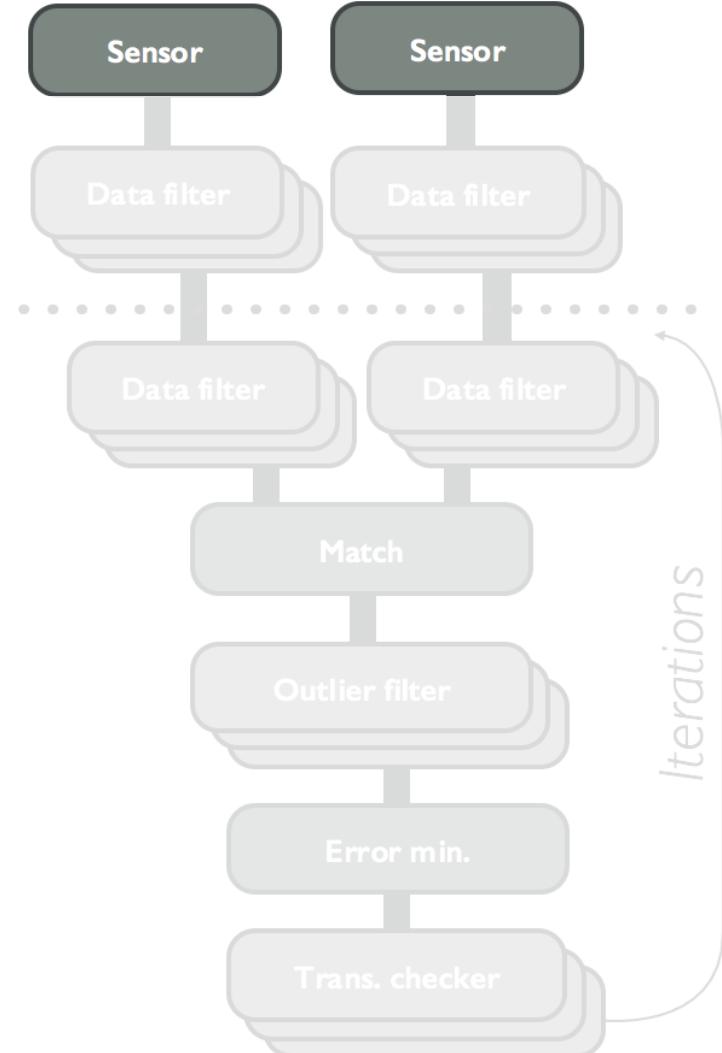


# Overview



- Sensors
- Data filters
- Matching function
- Robust regression
- Minimization
- Convergence detection

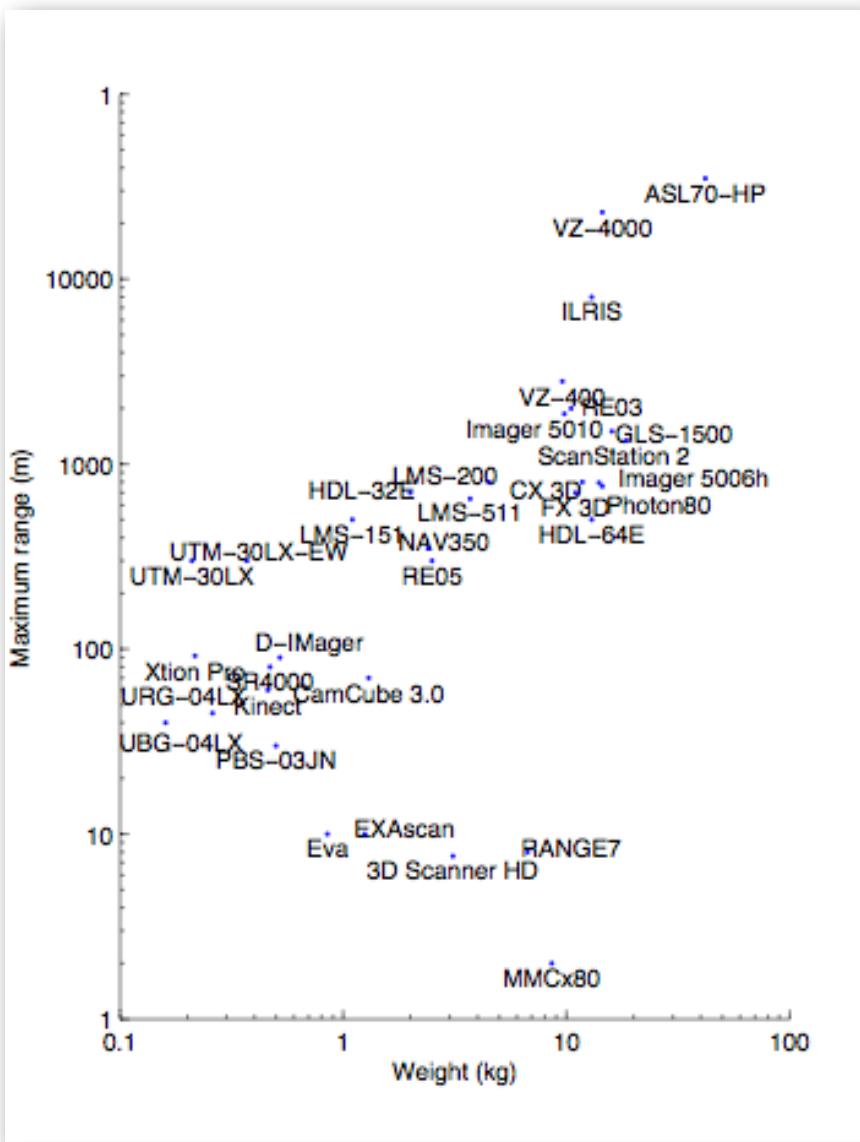
# Sensors



- Characteristics

- Power consumption
- Field of view
- Accuracy/repeatability
- Min/max distance
- Frequency
- Weight
- Volume
- Type of detection
- Wave length
- Etc.

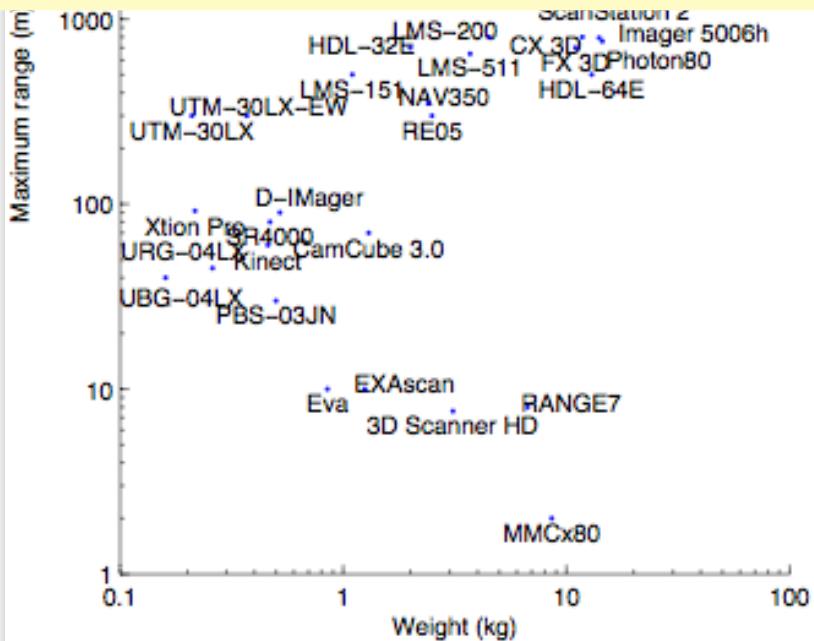
# Available sensors



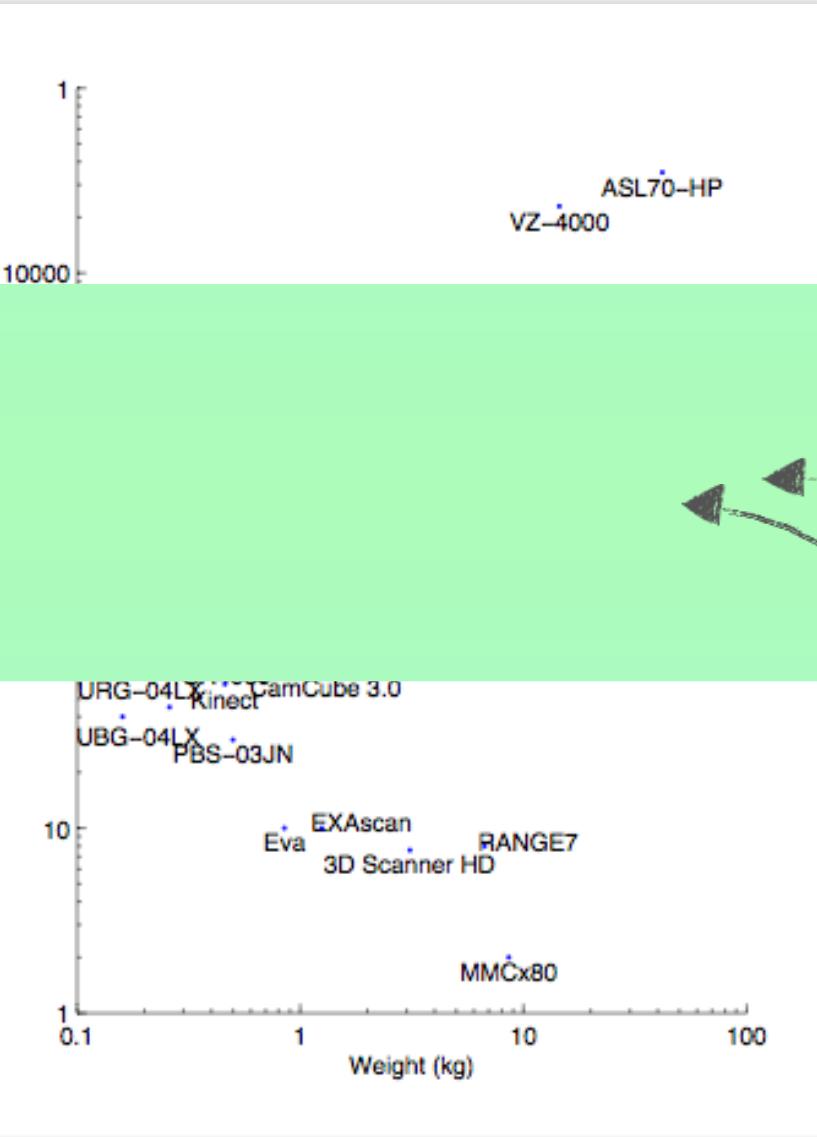
# Available sensors

1

## Airborne Survey



# Available sensors



*Terrestrial Survey*



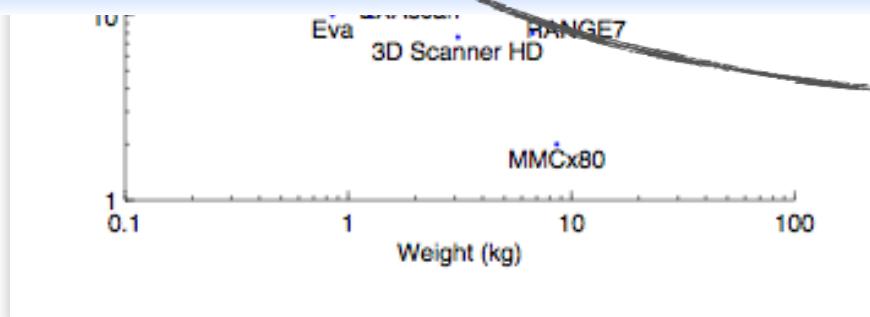
# Available sensors



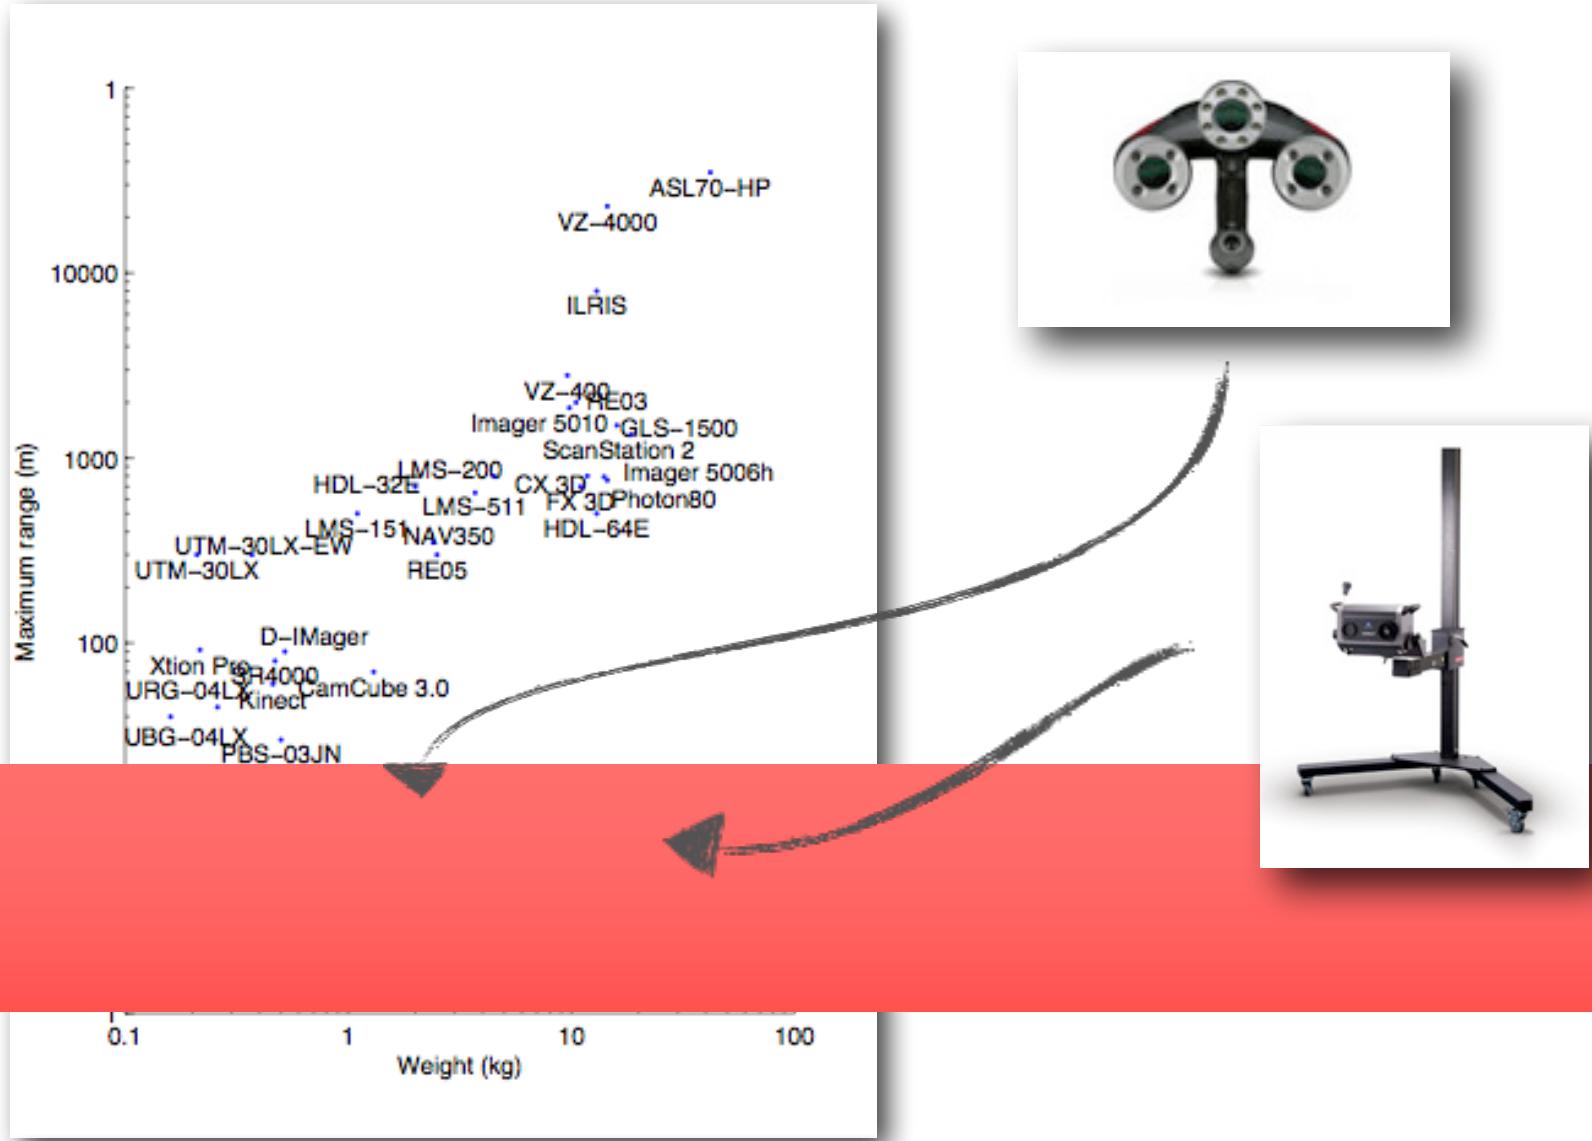
*Safety*

*Robotics*

*Entertainment*

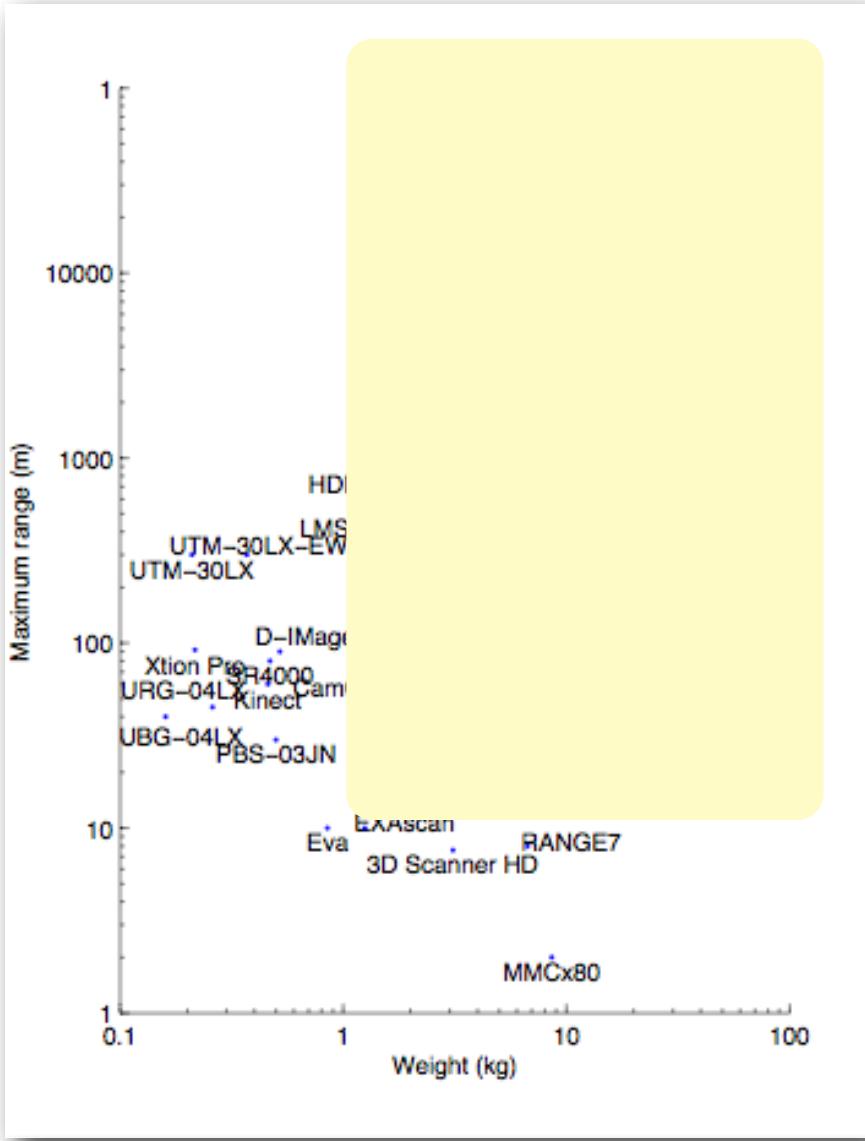


# Available sensors

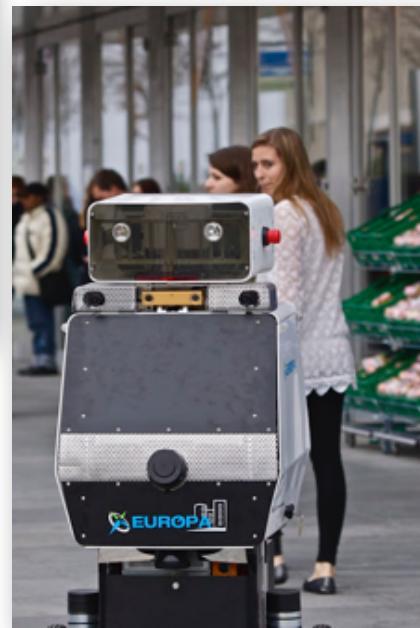
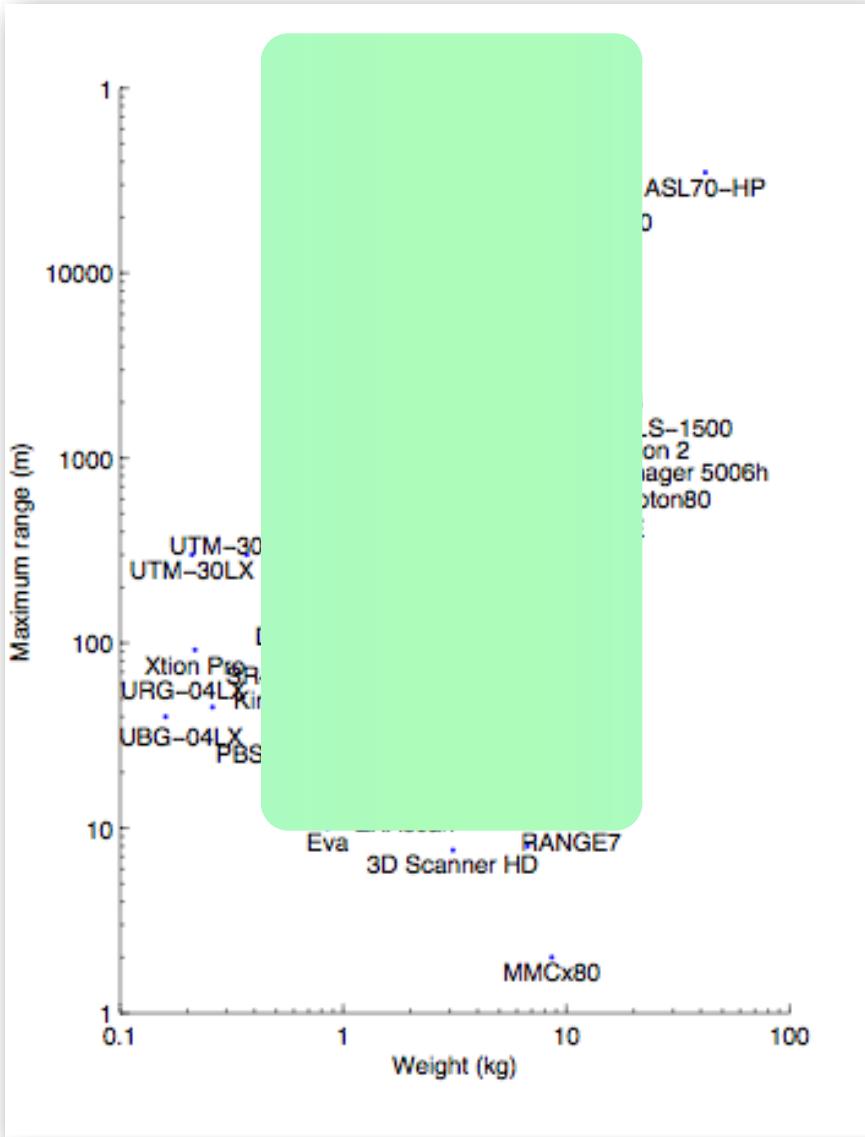


*Industrial Inspection*

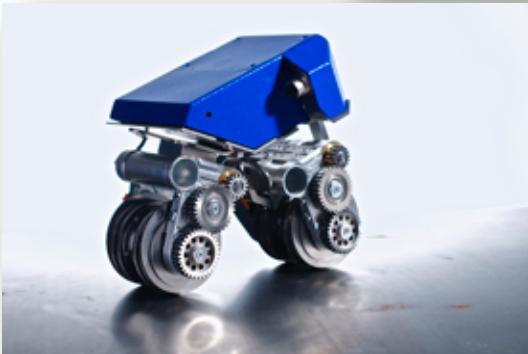
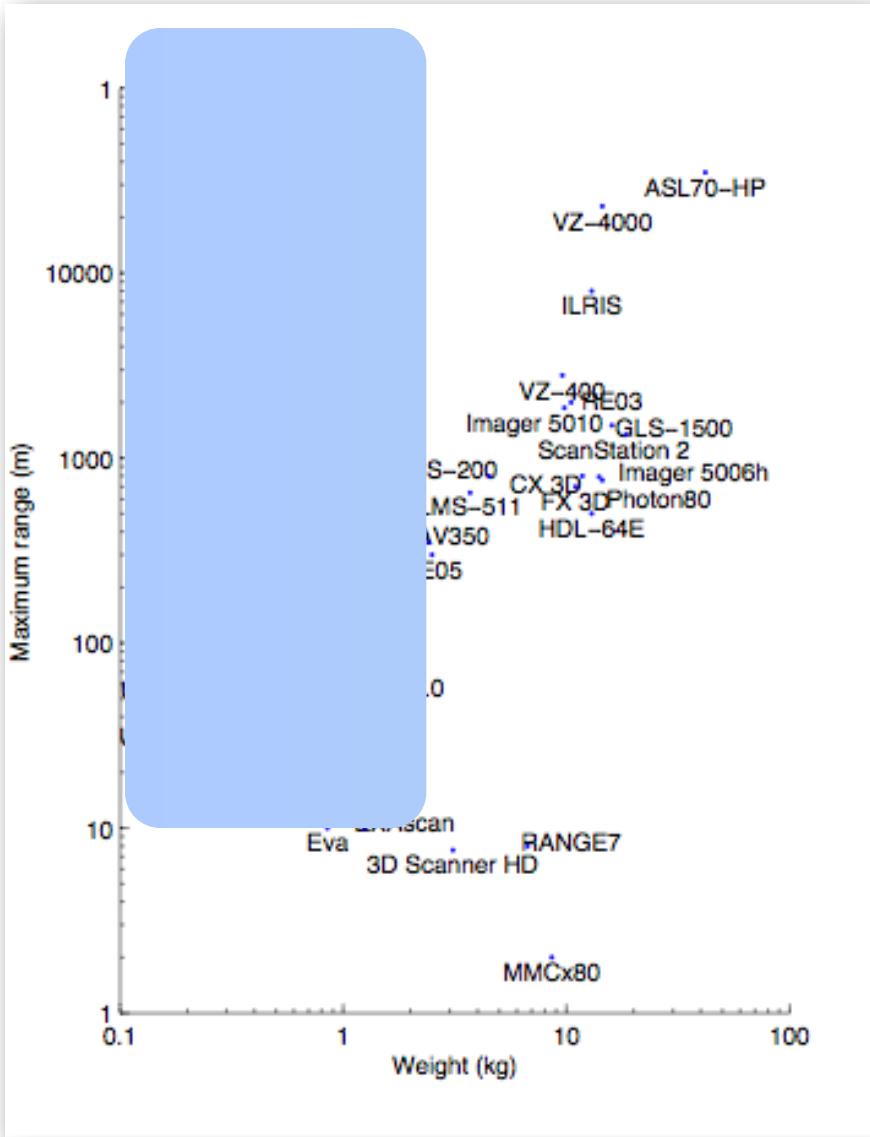
# Available sensors



# Available sensors



# Available sensors

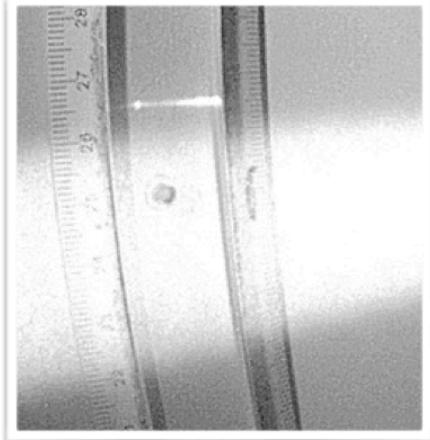


# Sensors: quality criteria

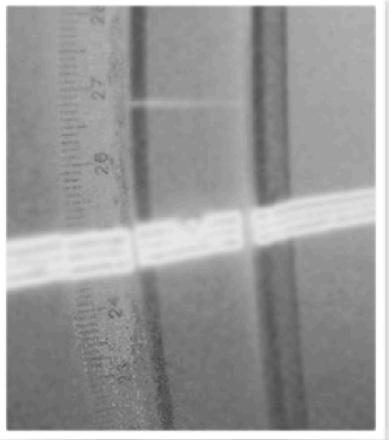


- Uncertainty
  - Laser beam width
  - Depth measurement
  - Types of surfaces
- Impact on registration
  - Mixed pixels
  - Point cloud density
  - Sensor position

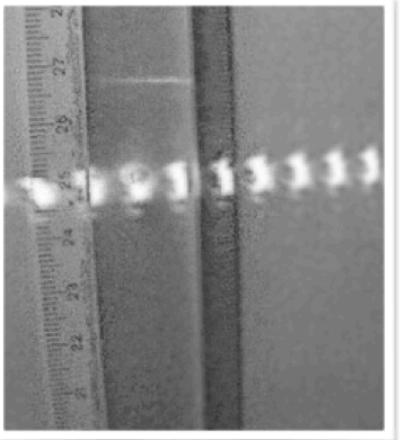
# Sensor uncertainty



LMS-151: **0.83 deg**



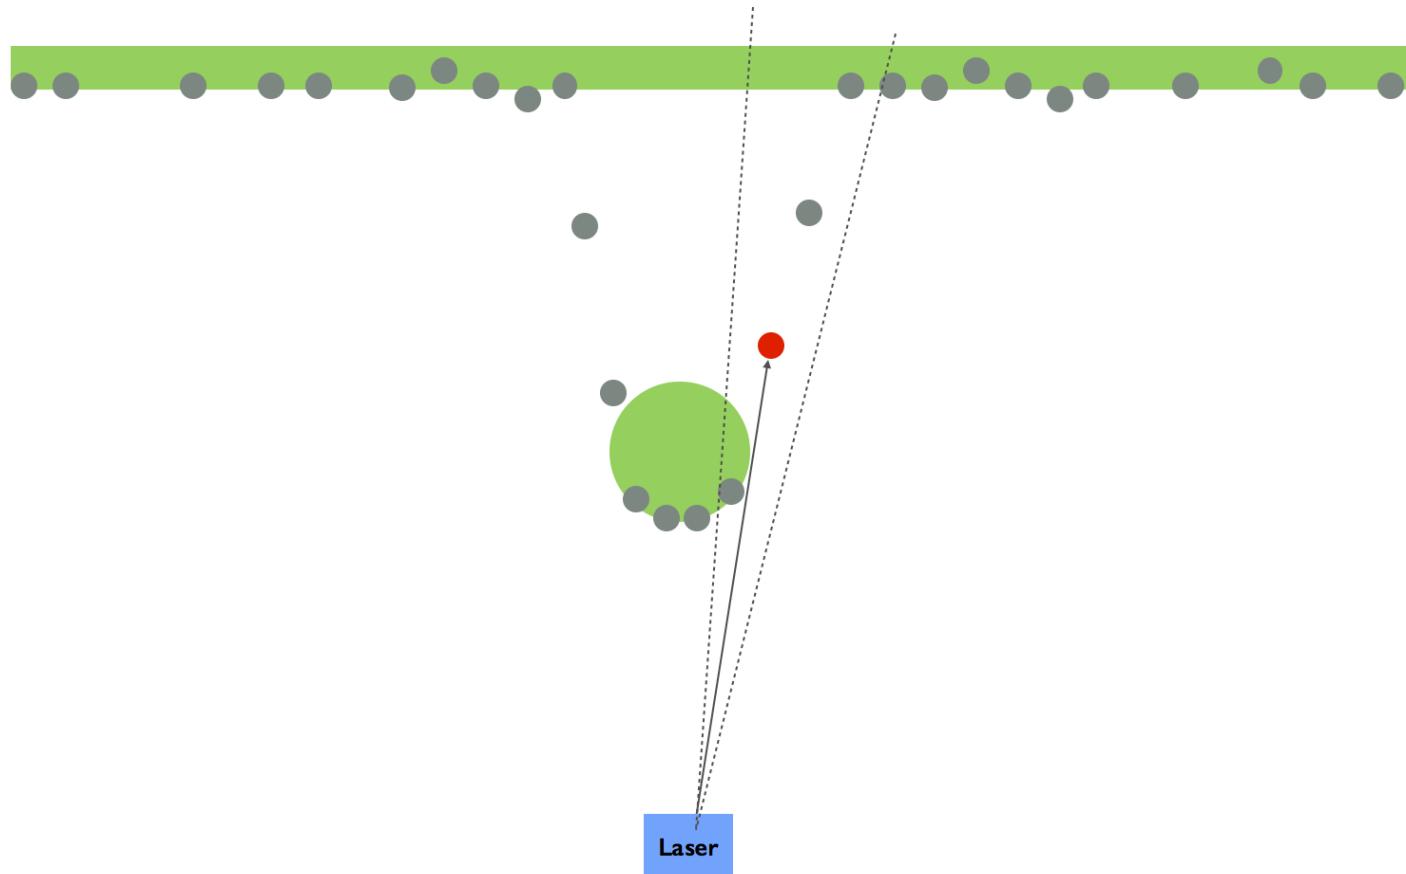
UTM-30LX: **0.13 deg**



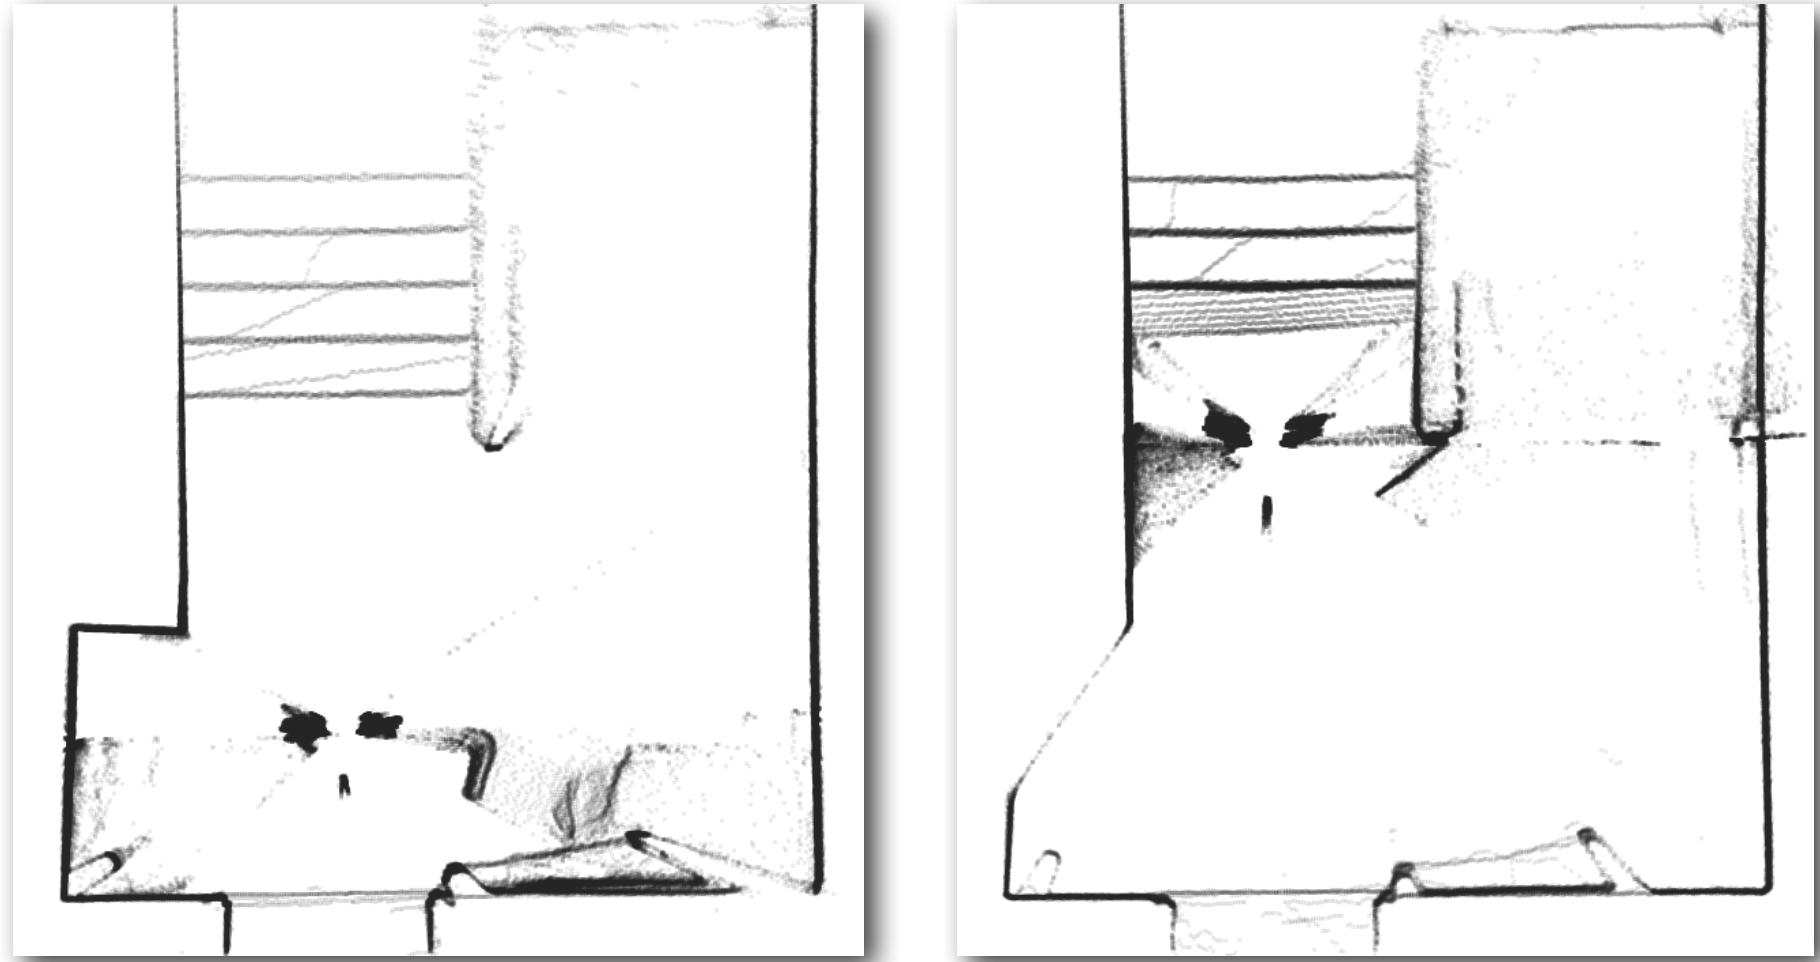
URG-04LX: **0.14 deg**



# Shadow points / mixed pixels



# Shadow points / mixed pixels



Top view of a staircase

# Sensor beam

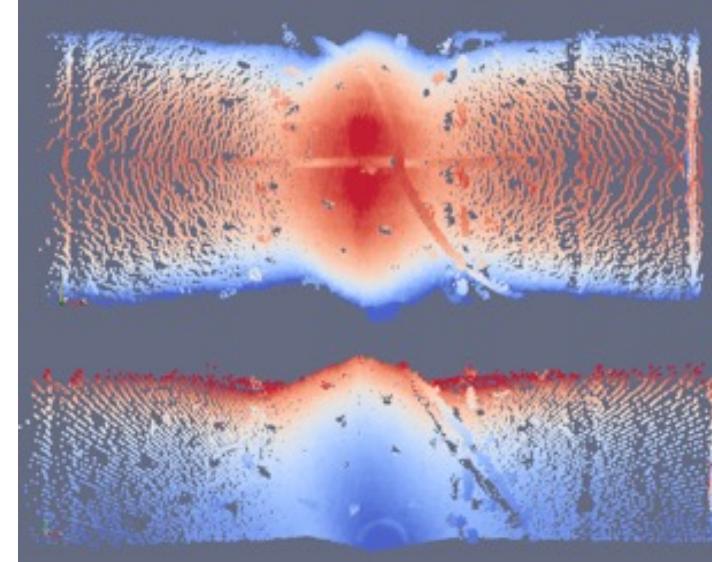


- Impact of large beam radius
  - More mixed pixels (shadow points)
  - More safety (with overlap between beams)

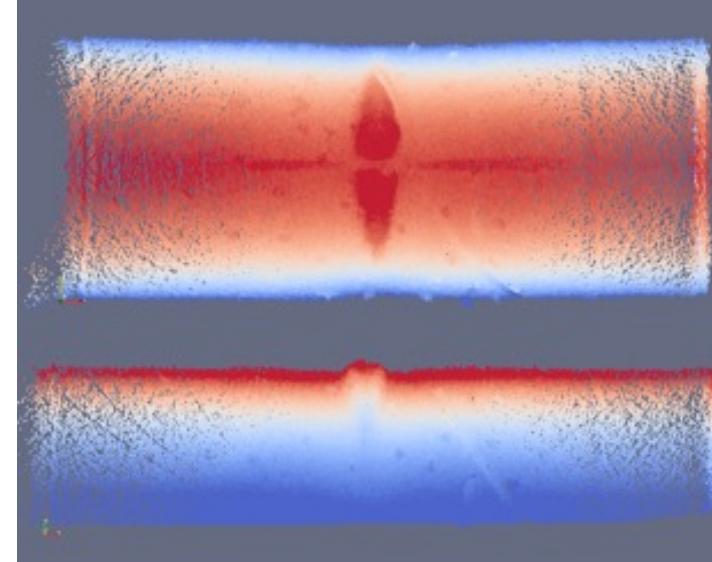
# Types of surface



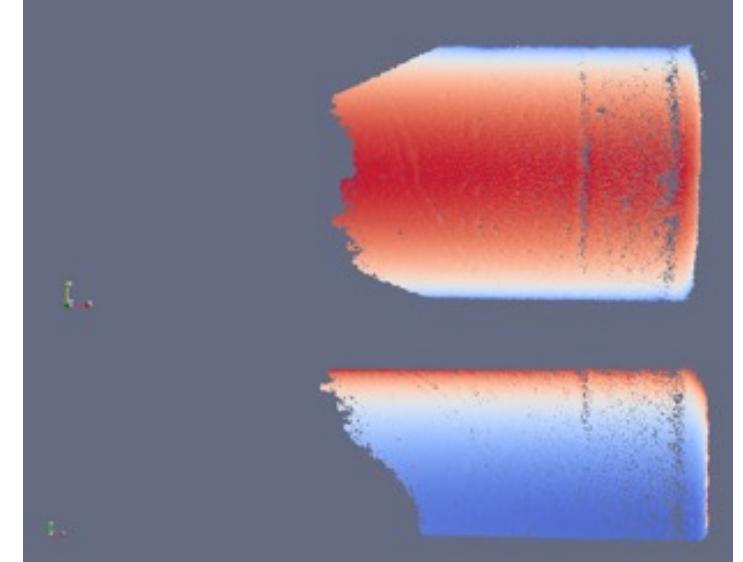
*URG-04LX*



*UTM-30LX*

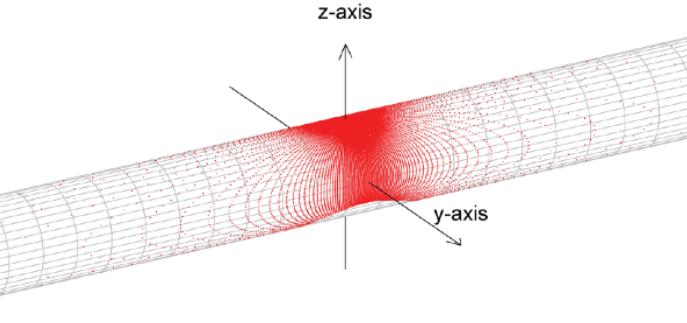
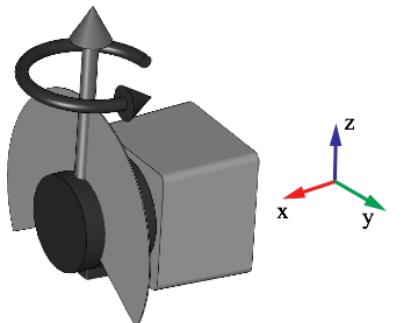
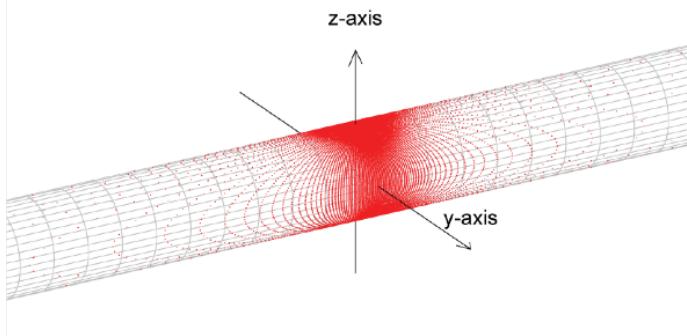
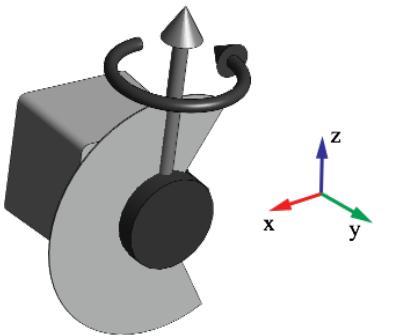
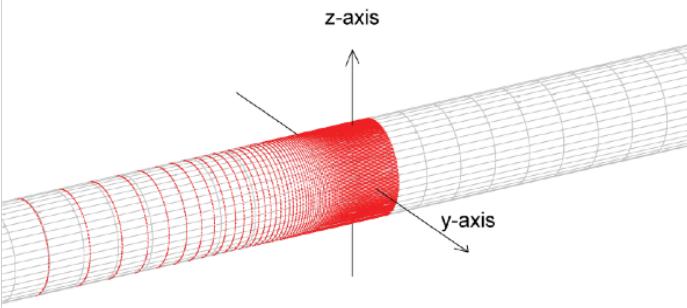
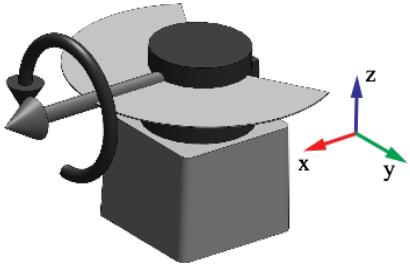
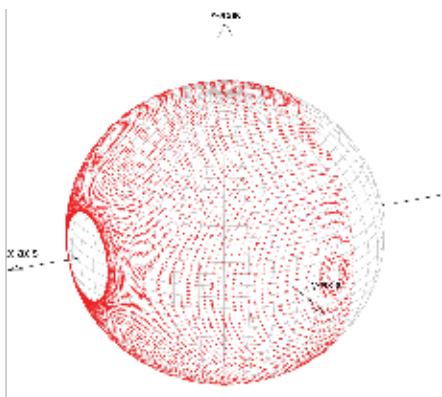
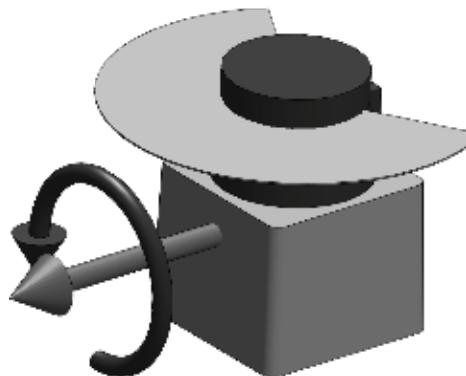
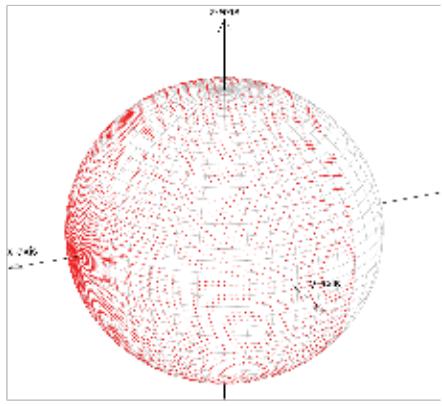
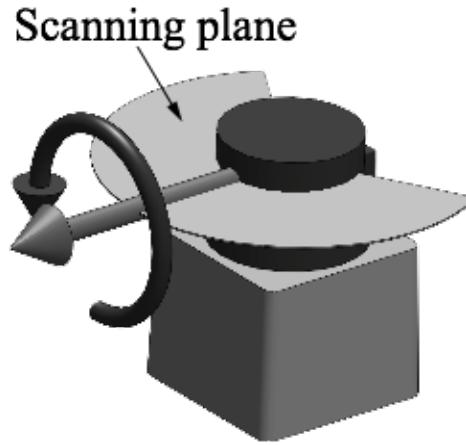


*Kinect*

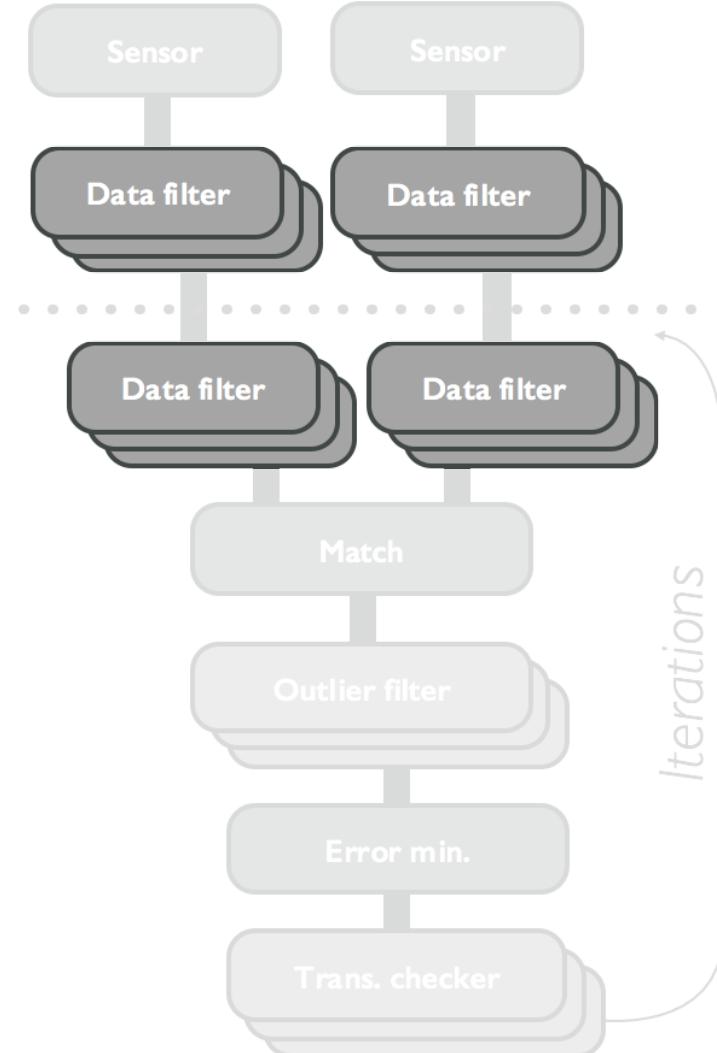


**Aluminum tube**

# Point cloud density



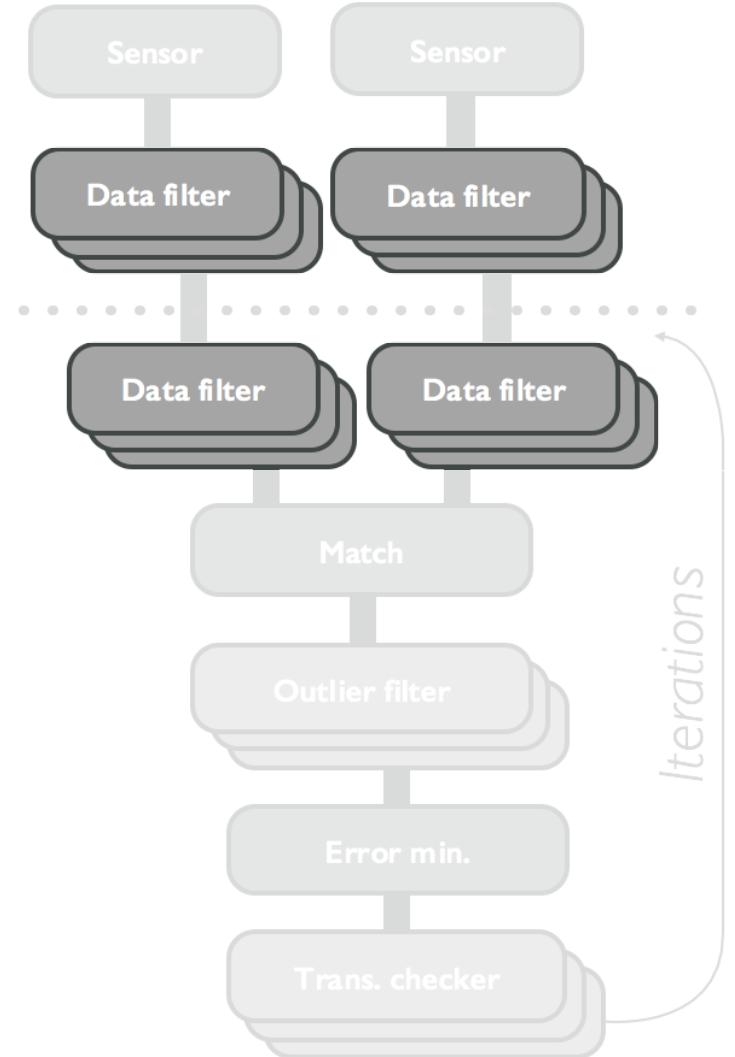
# Data filters



- Main goals

- Filter noise
- Reduce redundant information
- Augment discrepancy of points

# Data filters

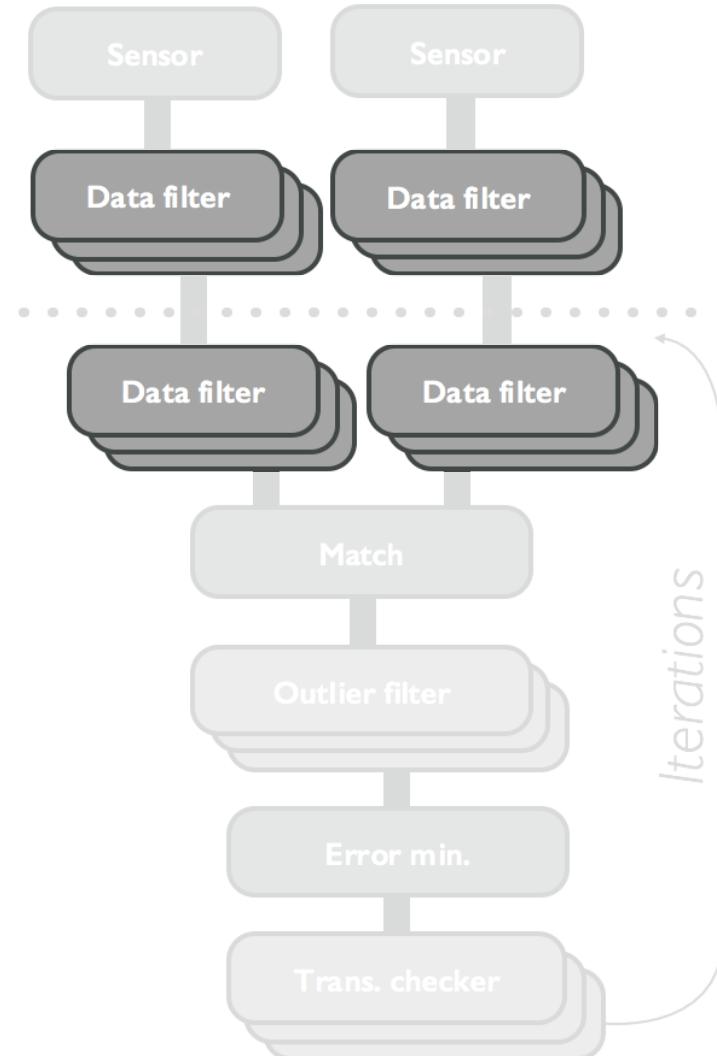


Number of points  $N$

Dimension K

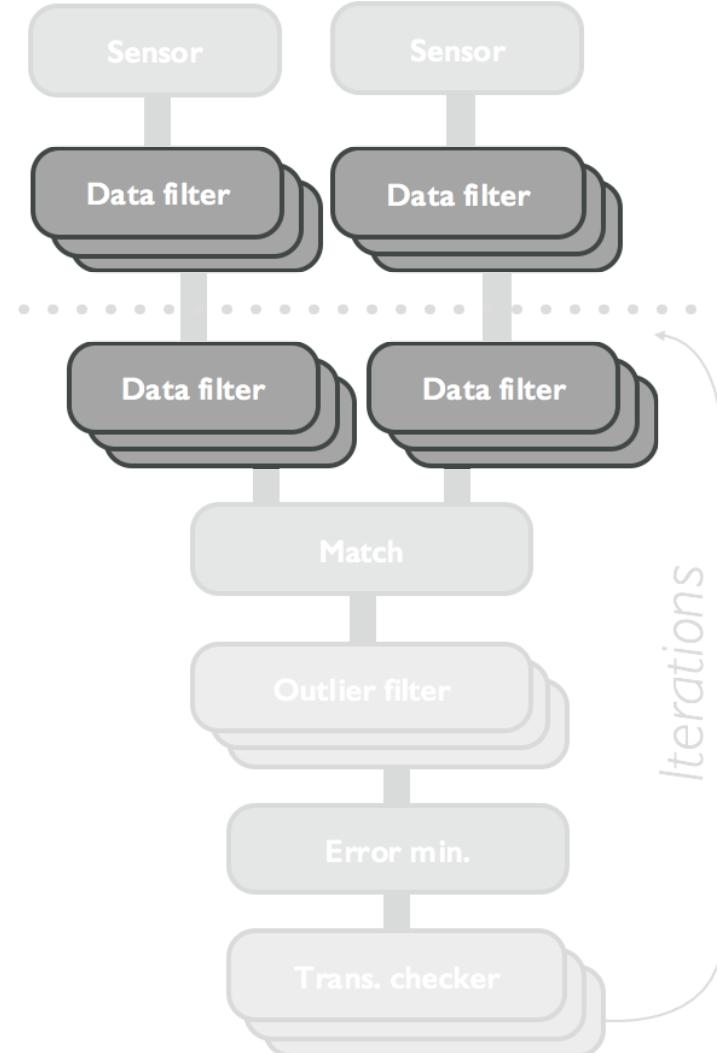
$$\begin{bmatrix} x_1 & x_2 & x_3 & \cdots & x_n \\ y_1 & y_2 & y_3 & \cdots & y_n \\ z_1 & z_2 & z_3 & \cdots & z_n \end{bmatrix}$$
$$\begin{bmatrix} r_1 & r_2 & r_3 & \cdots & r_n \\ g_1 & g_2 & g_3 & \cdots & g_n \\ b_1 & b_2 & b_3 & \cdots & b_n \end{bmatrix}$$

# Data filters



- Main goals:
    - Filter noise (reduce  $N$ )
    - Reduce redundant information (reduce  $N$ )
    - Augment discrepancy of points (augment  $K$ )

# Data filters



- Examples:

- Point reduction
- Sensor noise weight
- Surface extraction

# Point reduction

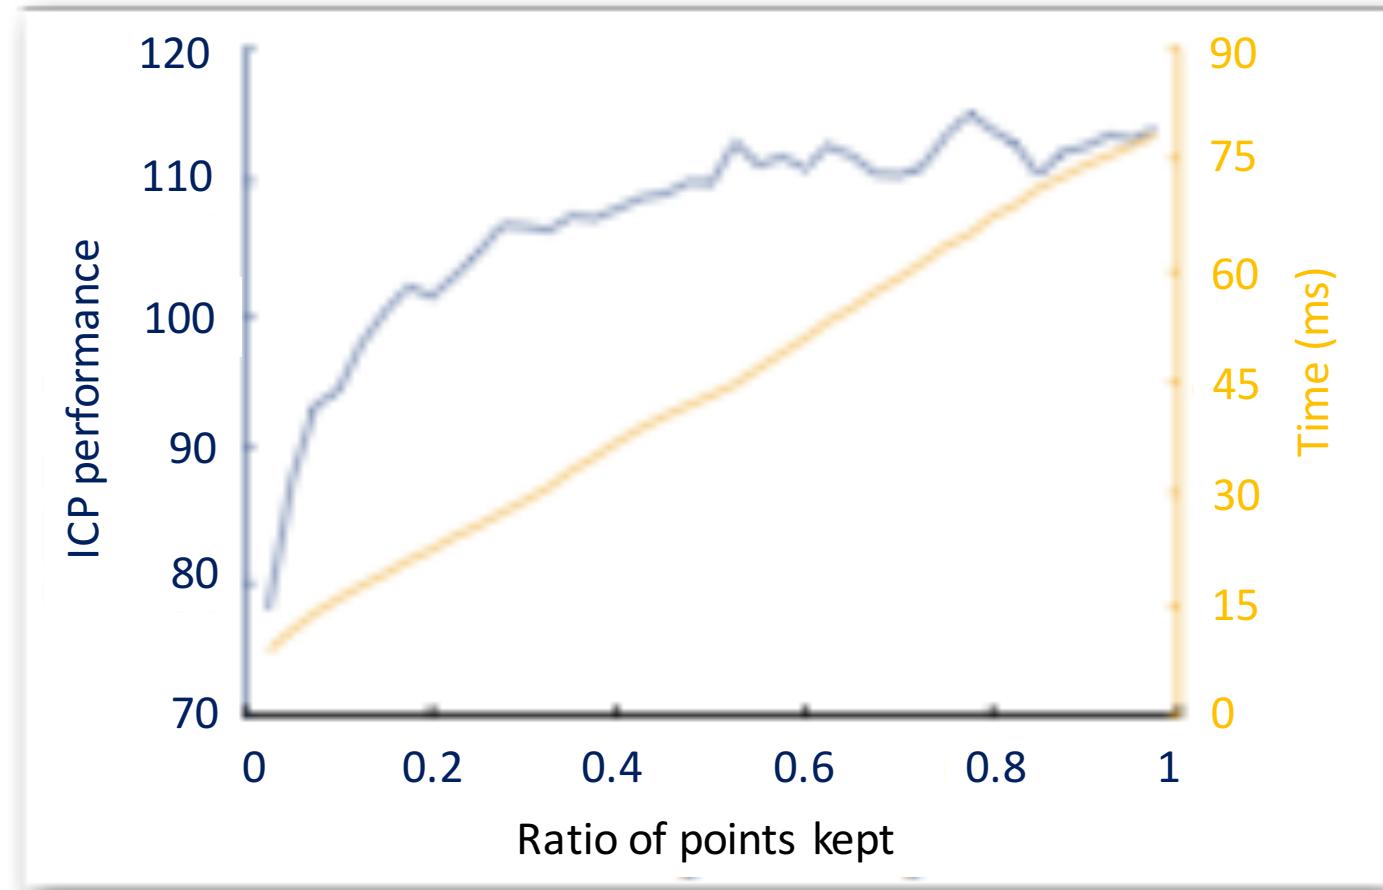


- Fixed-step subsampling
  - Very fast
  - May generate artifact
- Random sub-sampling
  - Fast
  - More control on reduction factor
- Density
  - Slow (needs a kD-tree)
  - Reduce minimization local minima
- Geometric sampling [Gelfand et al. 2003]
  - Very slow
  - Reduce more minimization local minima

# Point reduction

- Random subsampling

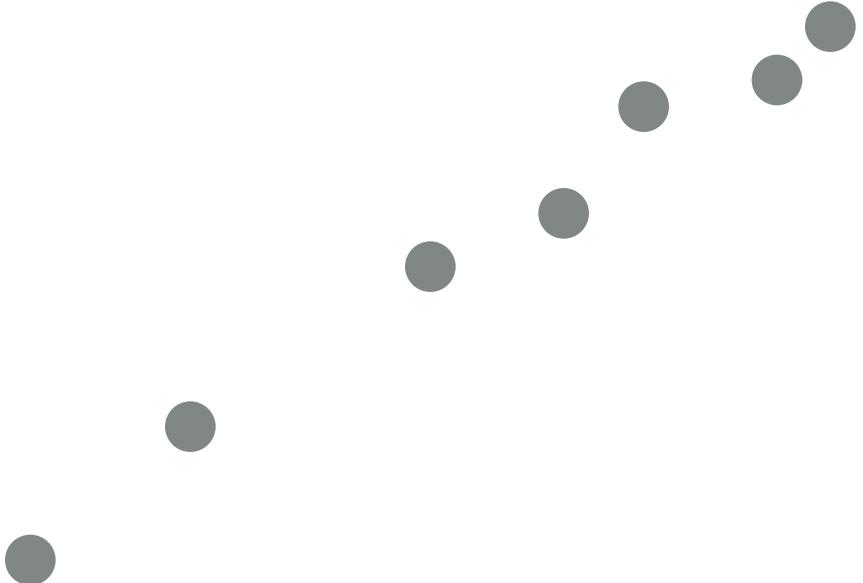
$N = 40,230$



# Surface extraction



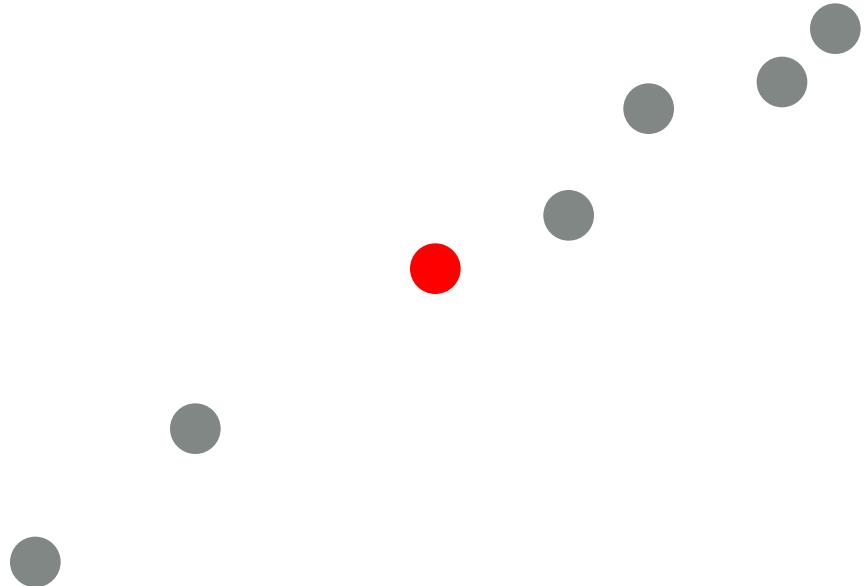
- For each point



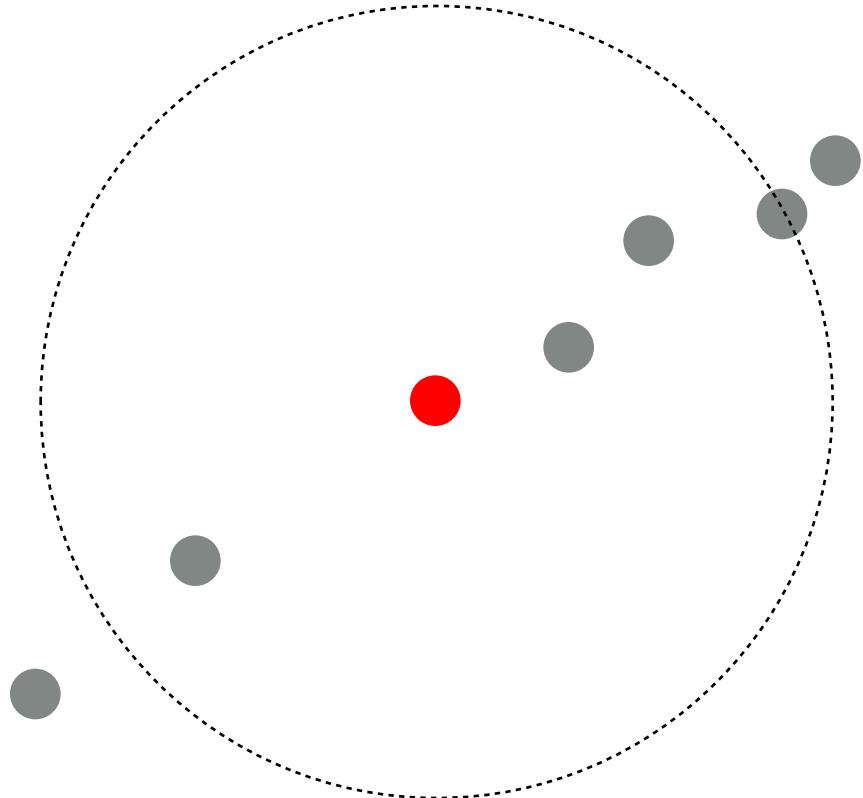
# Surface extraction



- For each point

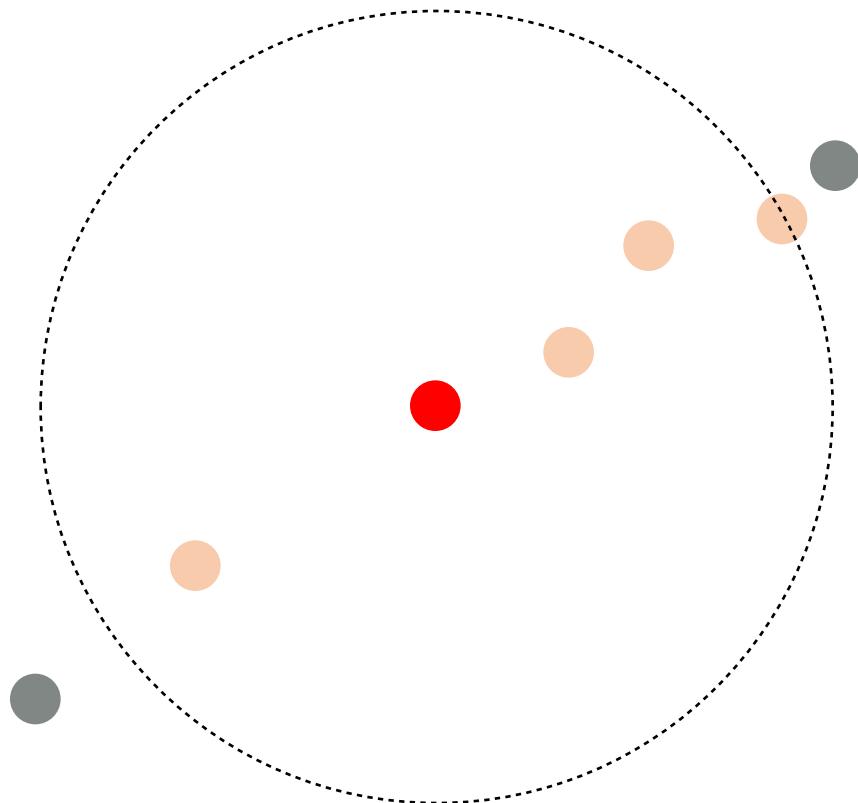


# Surface extraction



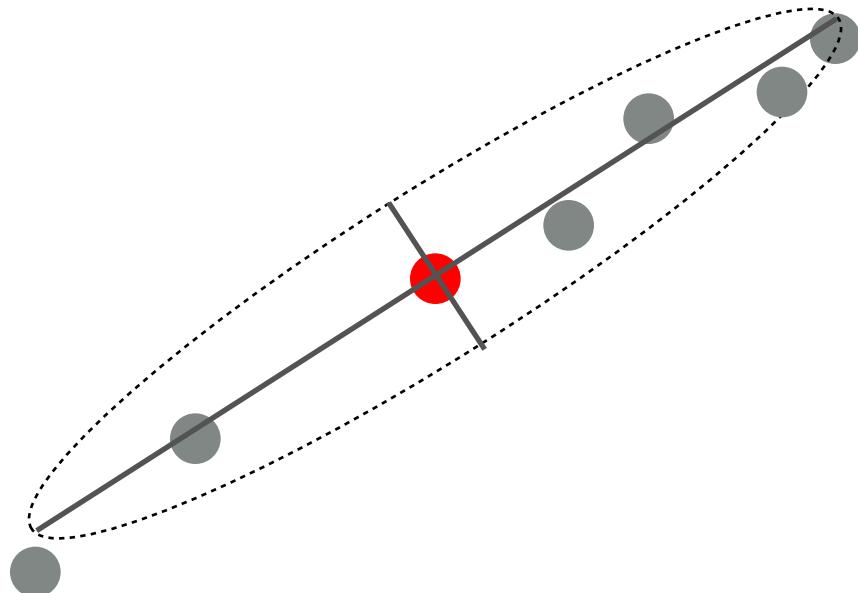
- For each point
  - Define neighborhood

# Surface extraction



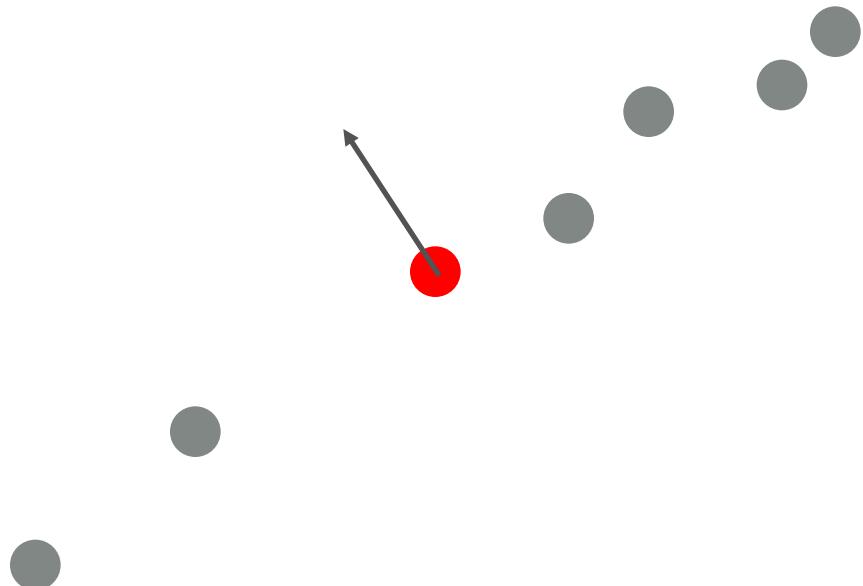
- For each point
  - Define neighborhood
  - Subtract “center”

# Surface extraction



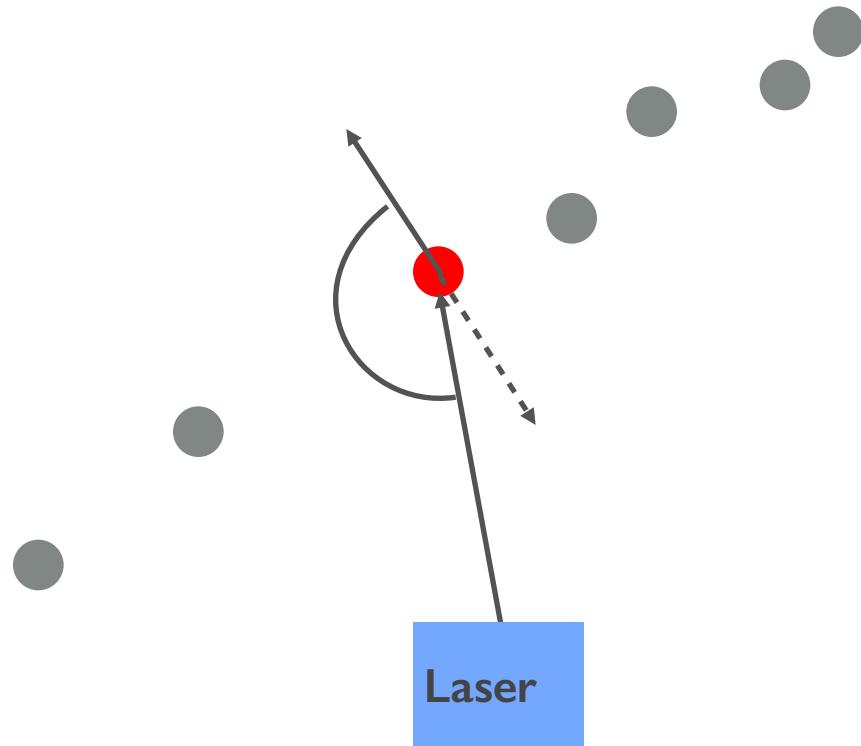
- For each point
  - Define neighborhood
  - Subtract “center”
  - Eigenvalue decomposition

# Surface extraction



- For each point
  - Define neighborhood
  - Subtract “center”
  - Eigenvalue decomposition
  - Select the eigenvector associated with the smallest Eigenvalue

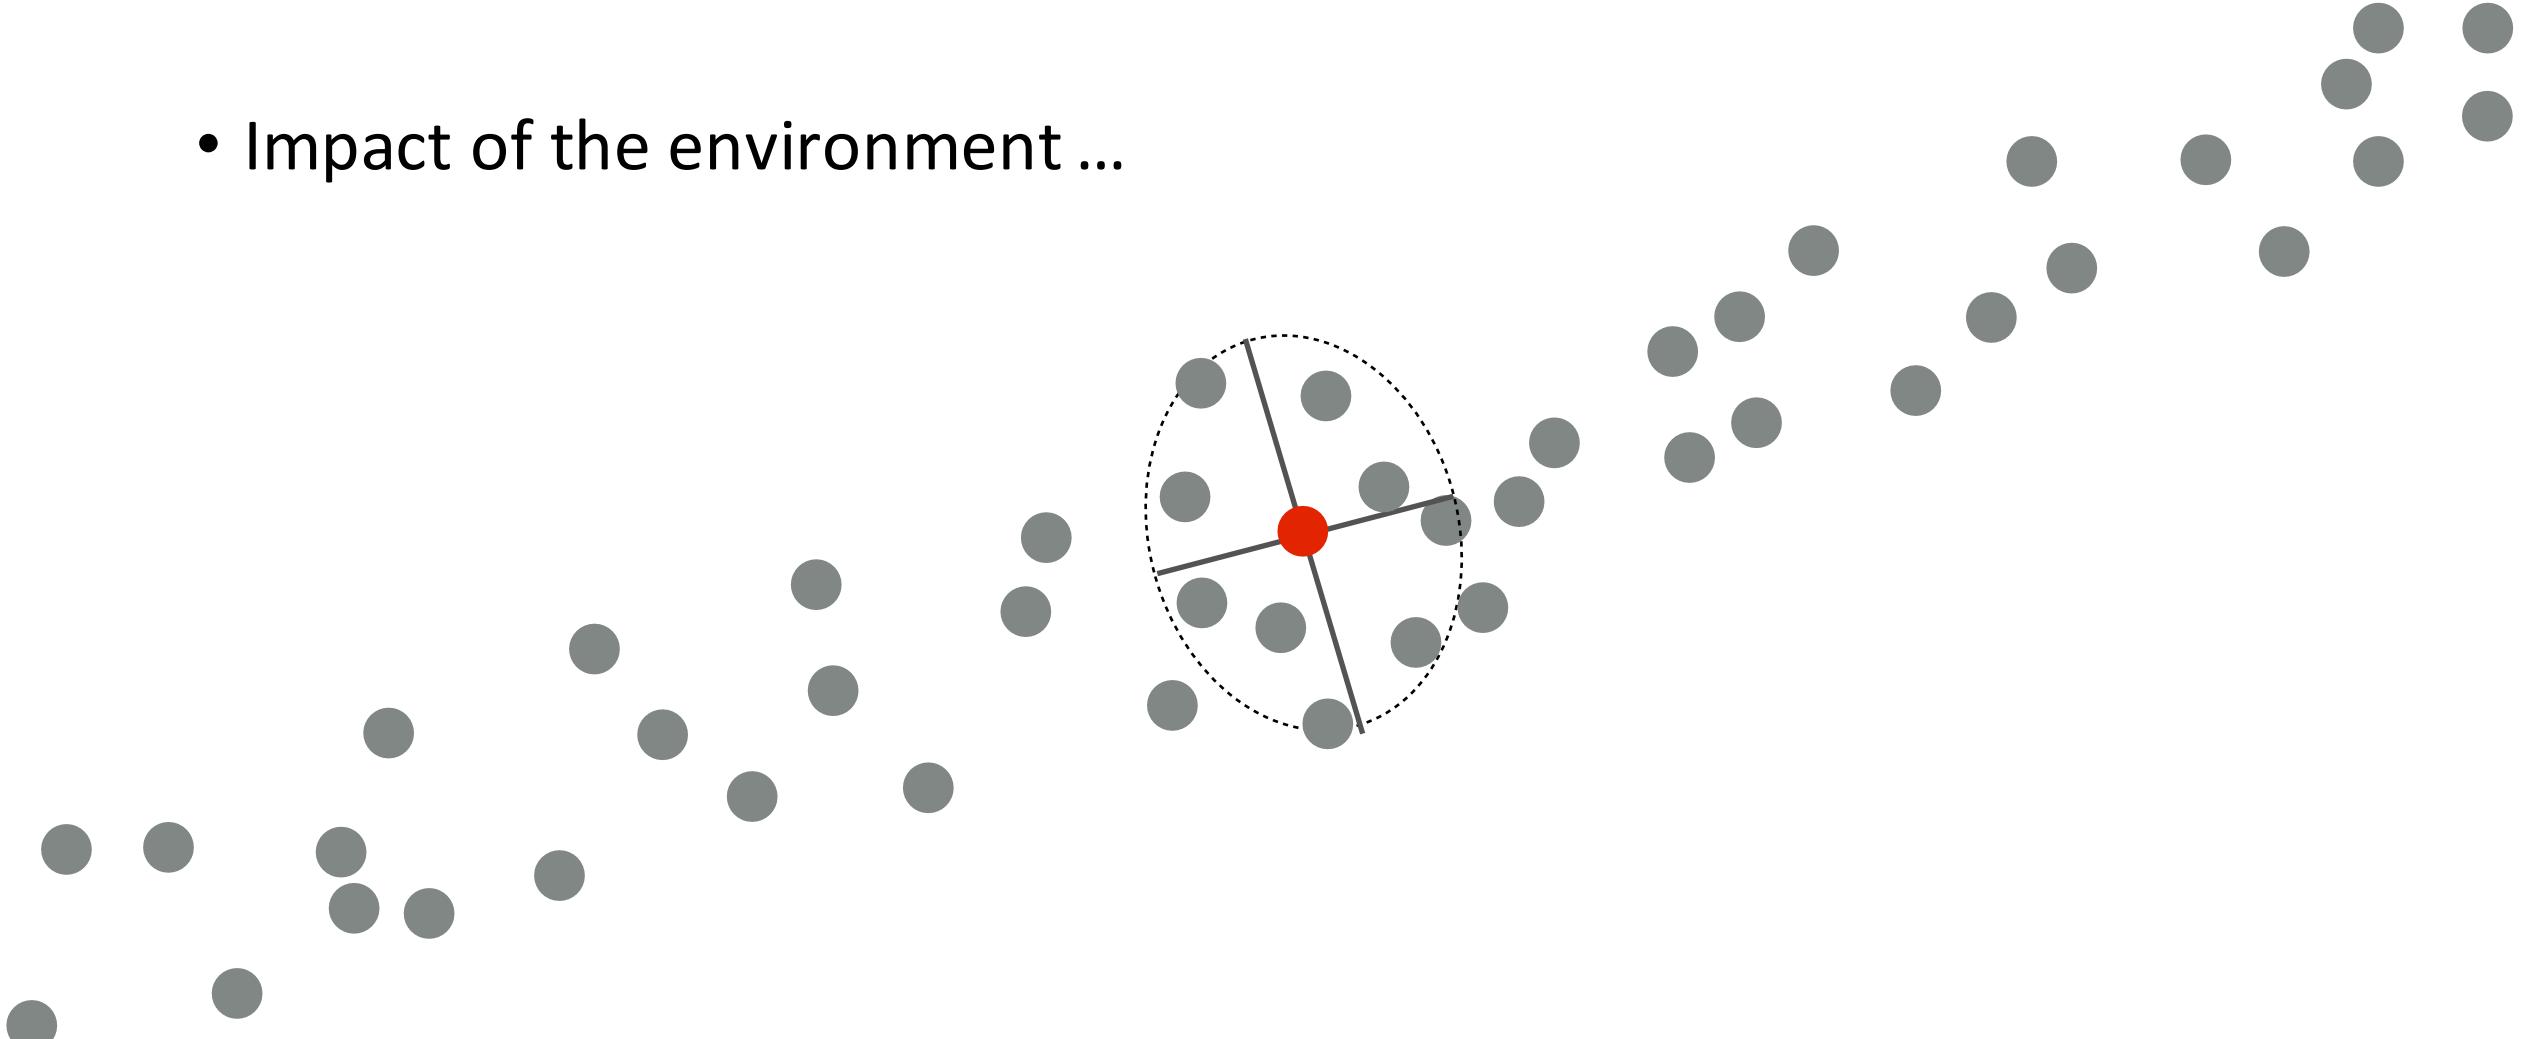
# Surface extraction



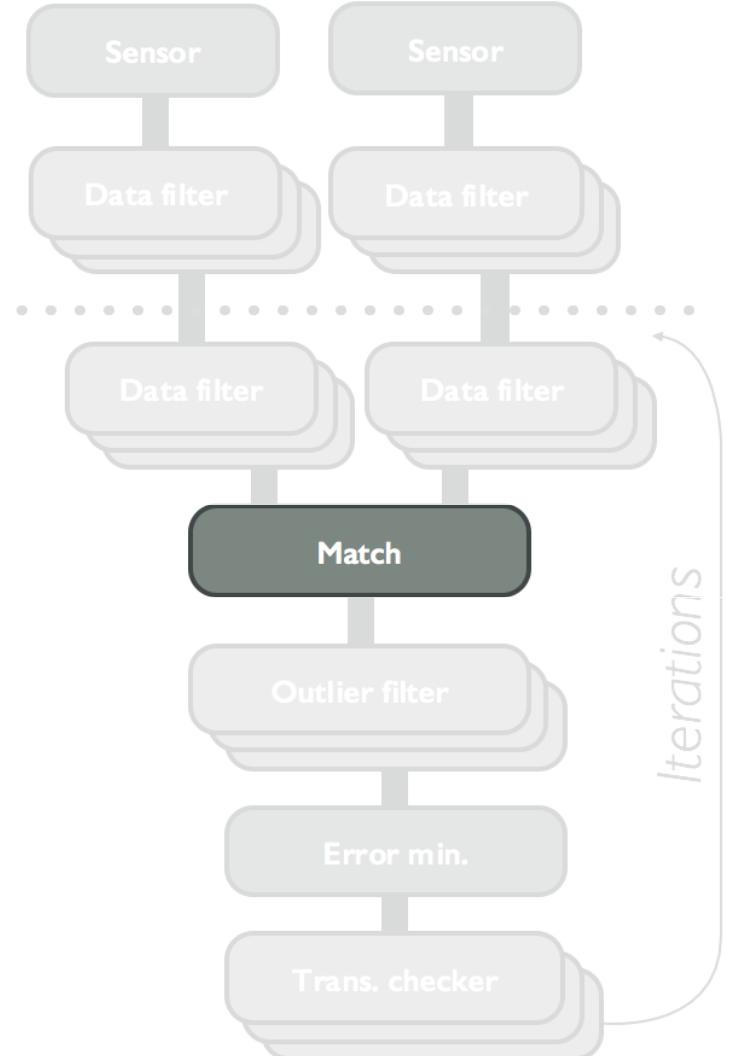
- For each point
  - Define neighborhood
  - Subtract “center”
  - Eigenvalue decomposition
  - Select the eigenvector associated with the smallest Eigenvalue
- Bonus:
  - Reorient the vector

# Surface extraction

- Impact of the environment ...



# Matching



- Main goal
- Augment robustness to misalignment

# Matching



- Match direction:
  - From reading to reference
  - From reference to reading
  - Both direction

# Matching



- Match direction:
  - From reading to reference
  - From reference to reading
  - Both direction
- Number of matches:
  - One-to-one (typical)
  - One-to-many  
(slower but more robust)

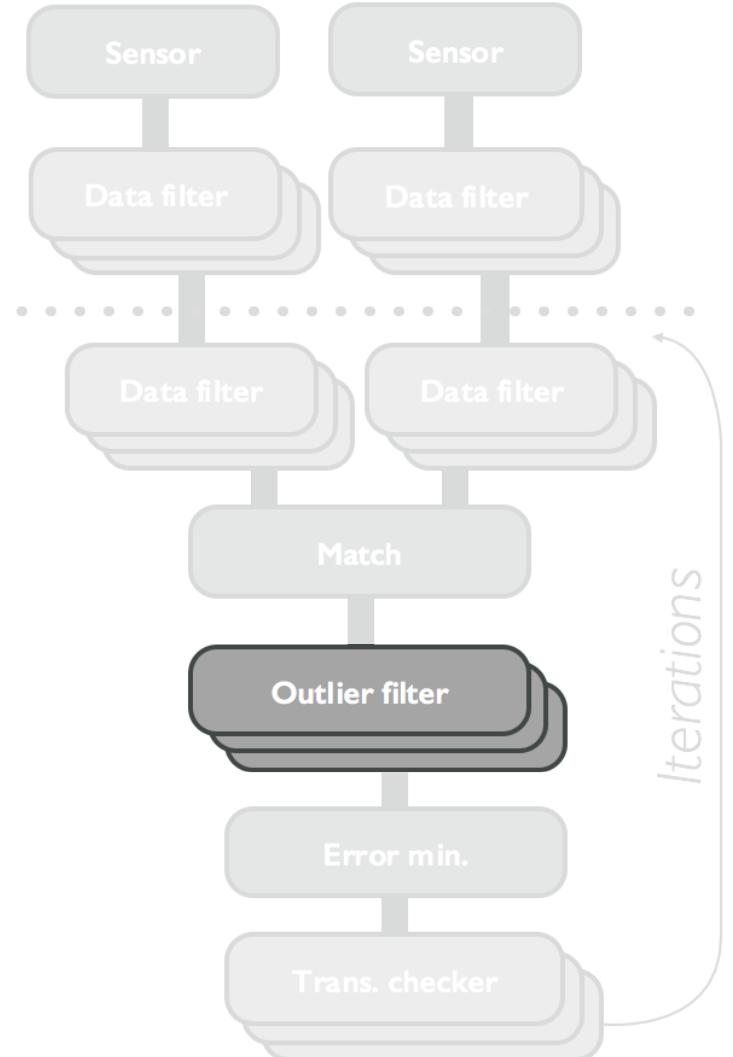
# Matching



- Match direction:
  - From reading to reference
  - From reference to reading
  - Both direction
- Number of matches:
  - One-to-one (typical)
  - One-to-many  
(slower but more robust)
- Metric
  - Euclidean distance on points
  - Euclidean distance on points and descriptors
  - Euclidean distance on descriptors
  - Mahalanobis distance
  - Probabilistic distance (entropy)
  - Others

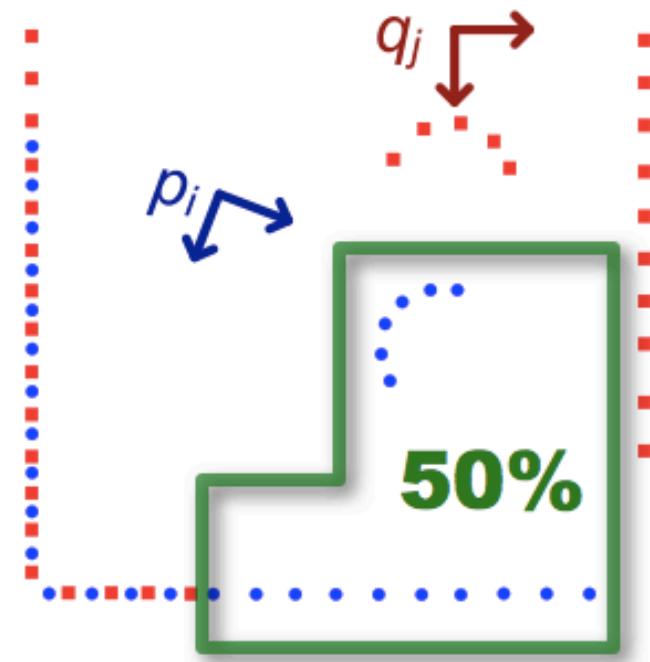
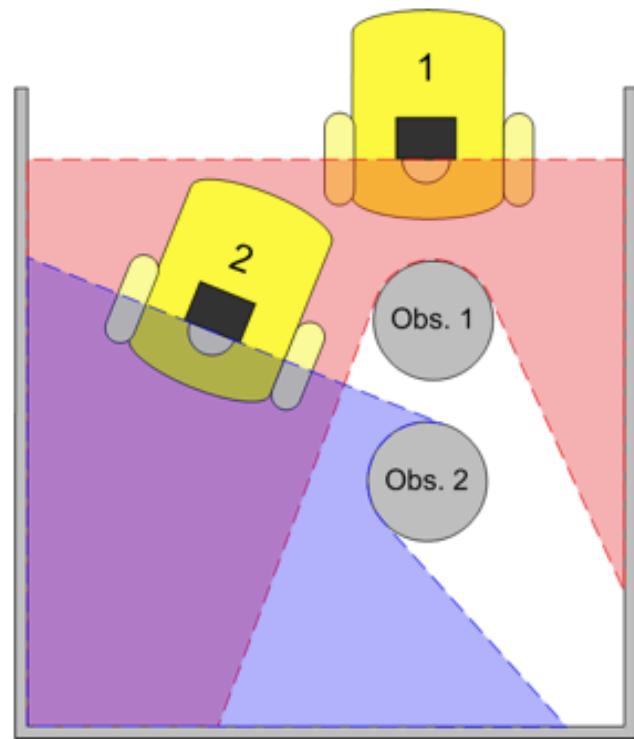
↑  
kD-tree  
↓

# Robust regression



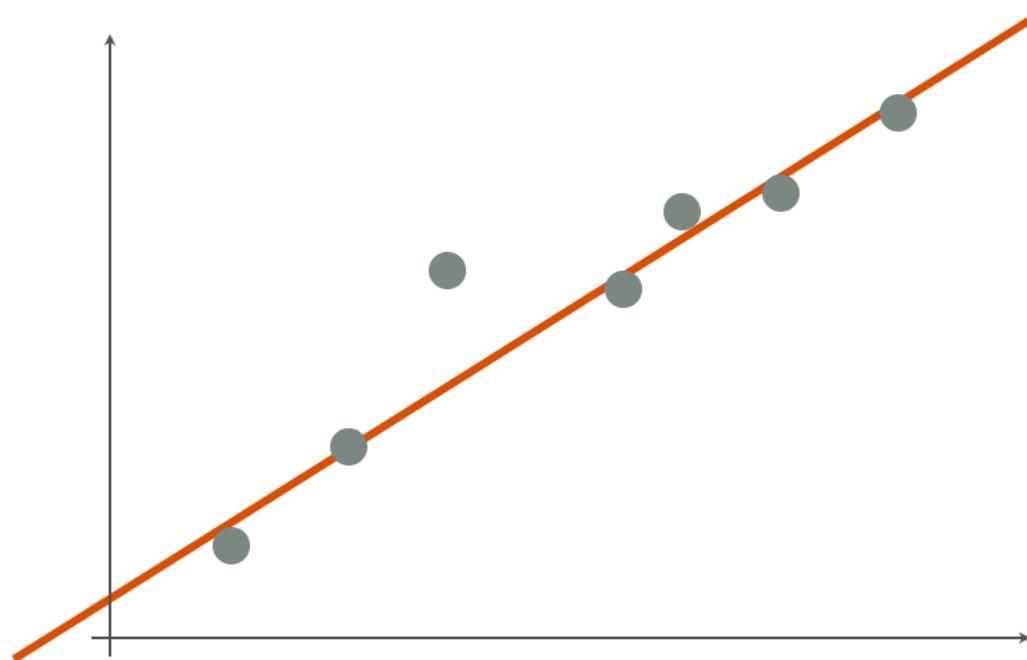
- Main goal
  - Augment robustness to overlap
  - Augment robustness to noise
  - Augment robustness to dynamic elements

# Robustness in case of partial overlaps



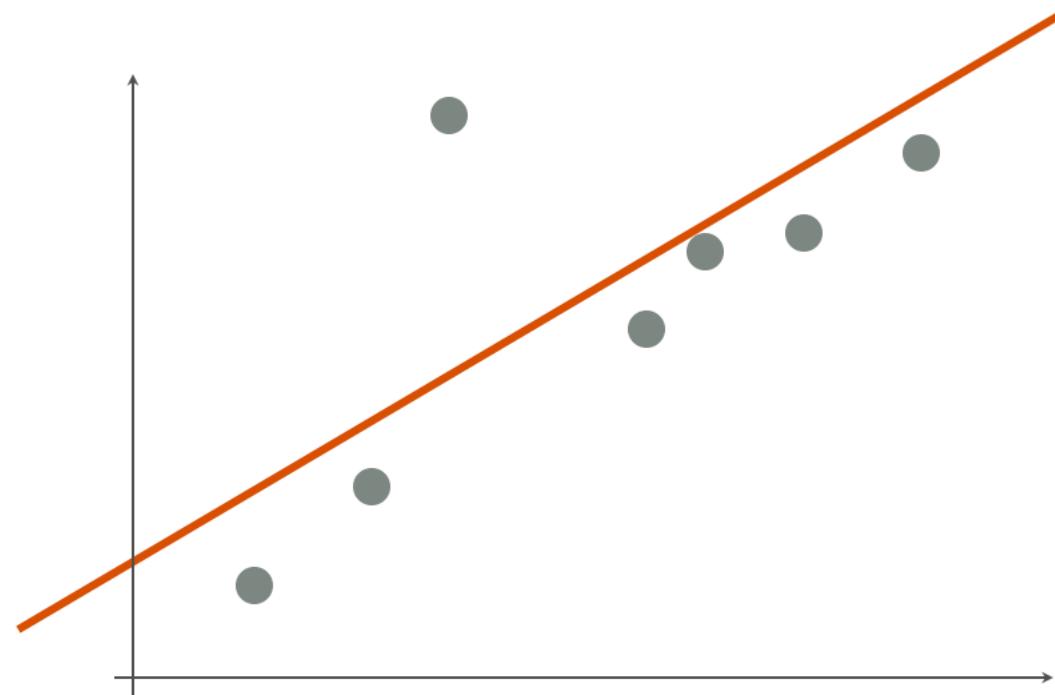
# Robust regression

Why mean  
least-squared is  
sensible to outlier?



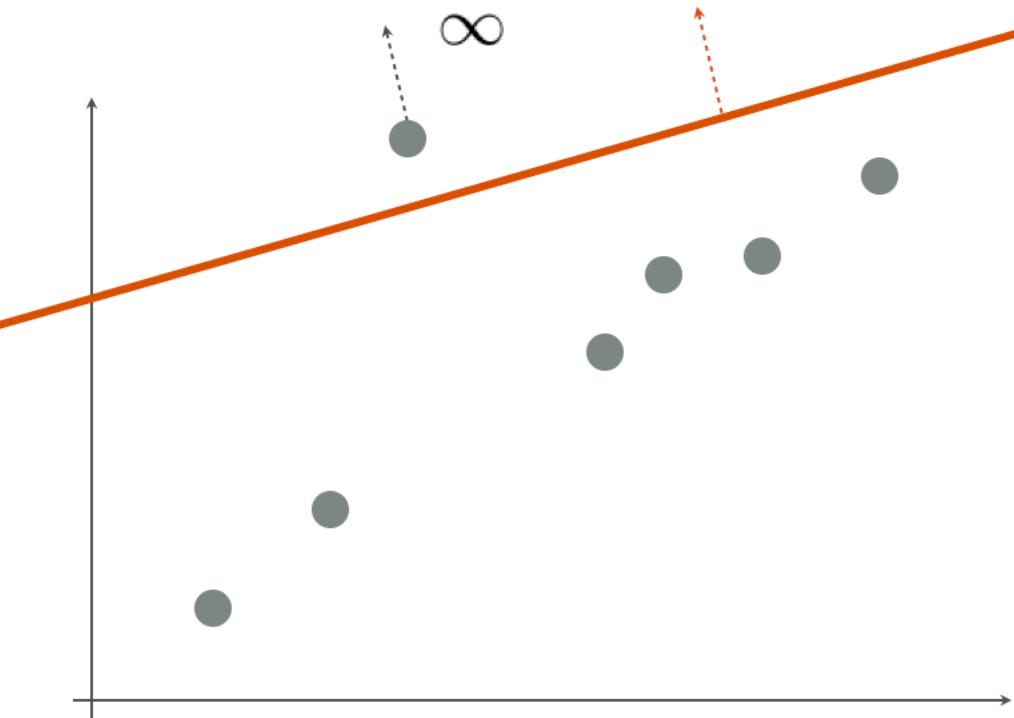
# Robust regression

Why mean  
least-squared is  
sensible to outlier?



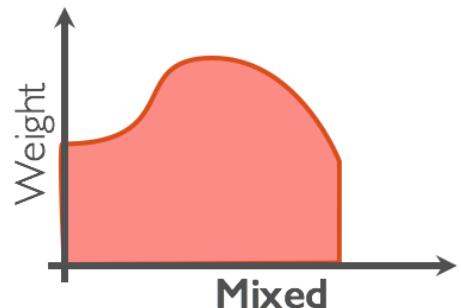
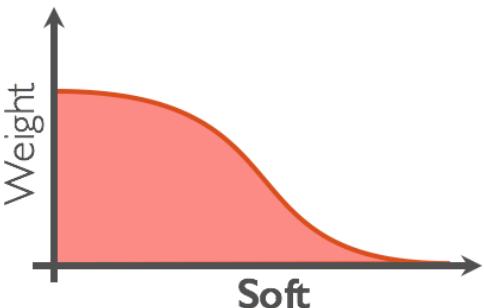
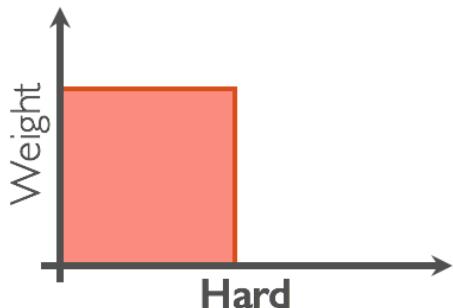
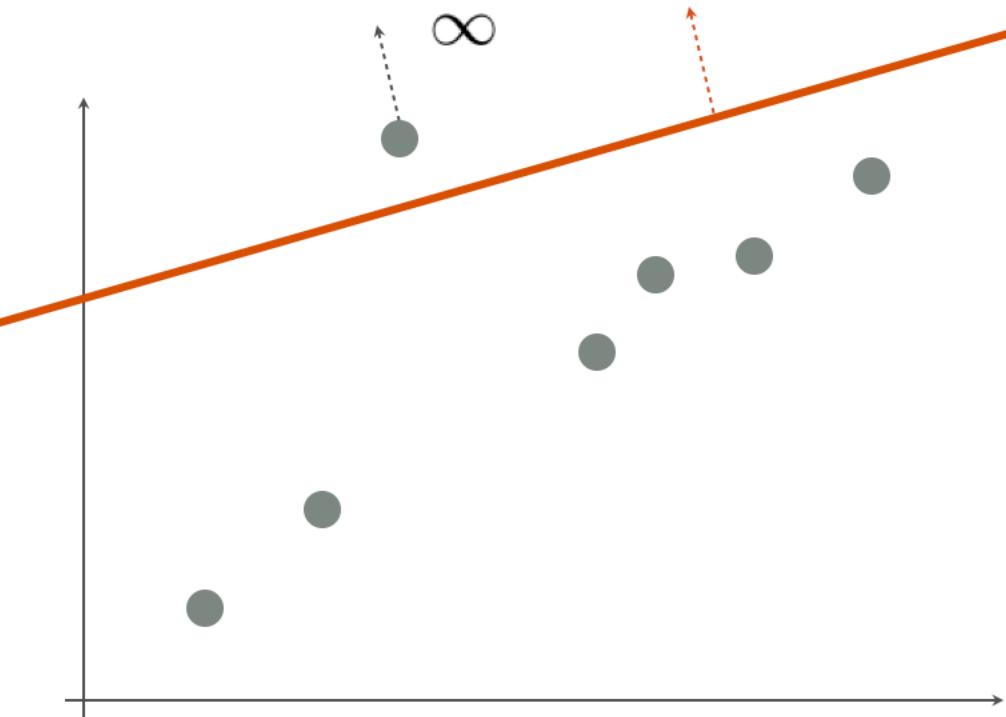
# Robust regression

Why mean  
least-squared is  
sensible to outlier?



# Robust regression

Why mean least-squared is sensible to outlier?



# Rejection



- Fix

$$d_i < d_{max}$$

- Zhang
- Mean
- Median
- Trimmed

# Rejection

- Fix
- **Zhang** (Zhang 94)

$$d_i < \begin{cases} \mu + 3\sigma, & \text{if } 0 \leq \mu < \eta \\ \mu + 2\sigma, & \text{if } \eta < \mu \leq 3\eta \\ \mu + \sigma, & \text{if } 3\eta < \mu \leq 6\eta \\ median, & \text{otherwise} \end{cases}$$

- Mean
- Median
- Trimmed

# Rejection



- Fix
- Zhang
- **Mean** (*Druon 06*)

$$d_i < \mu + \sigma$$

- Median
- Trimmed

# Rejection

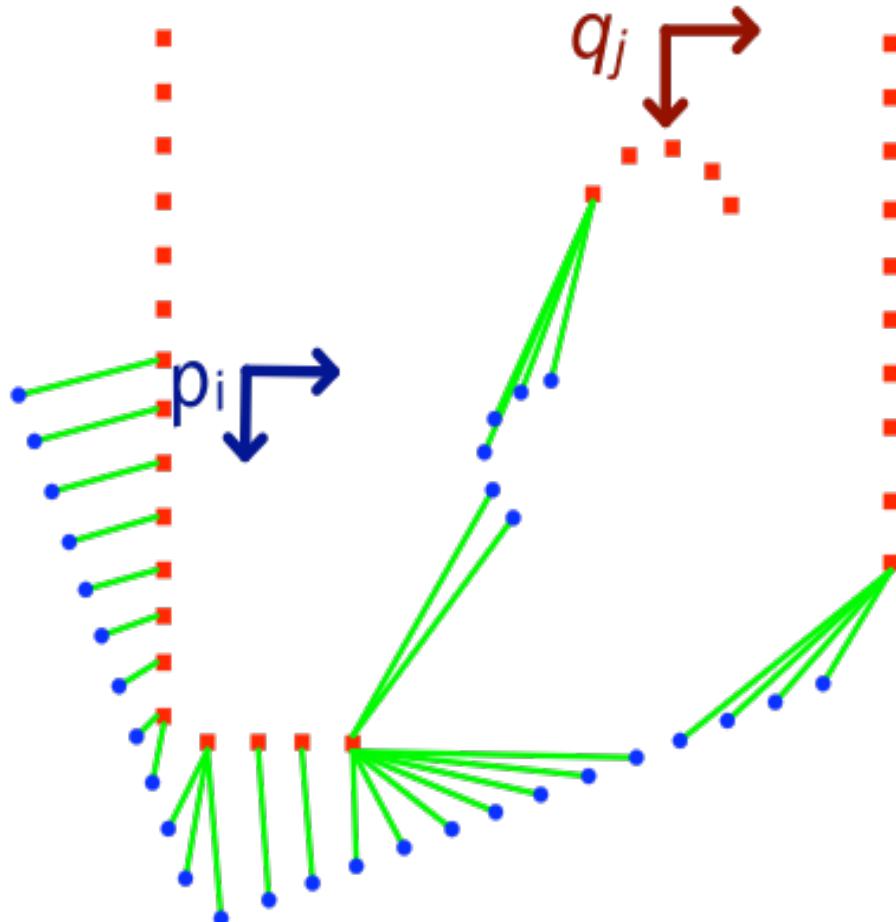


- Fix
- Zhang
- Mean
- **Median** (*Diebel 04*)
- Trimmed

$$d_i < 3 \times median$$

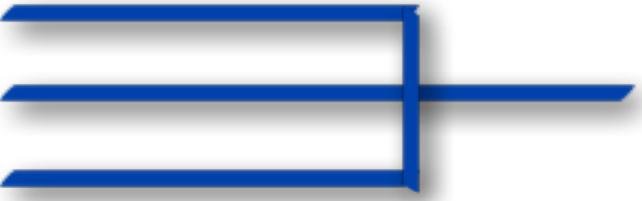
# Rejection

- Fix
- Zhang (*Zhang 94*)
- Mean (*Druon 06*)
- Median (*Diebel 04*)
- Trimmed (*Chetverikov 02*)
  - *remove “x percent” of worst matches*

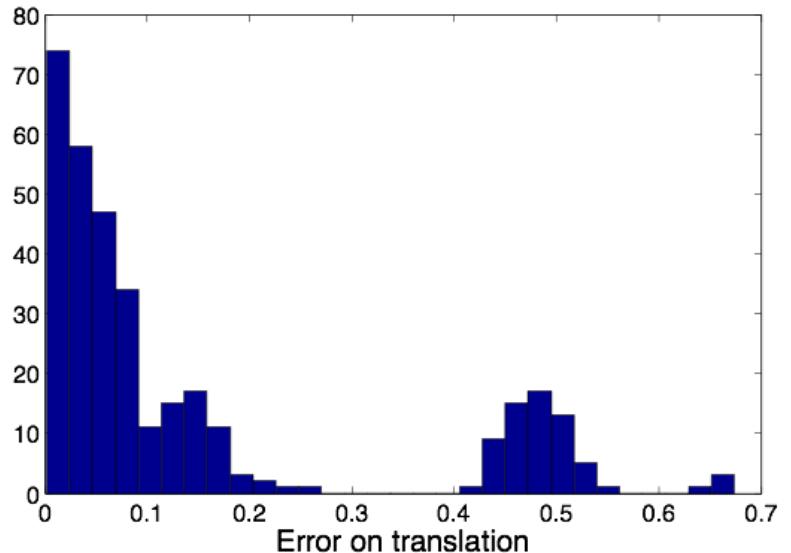
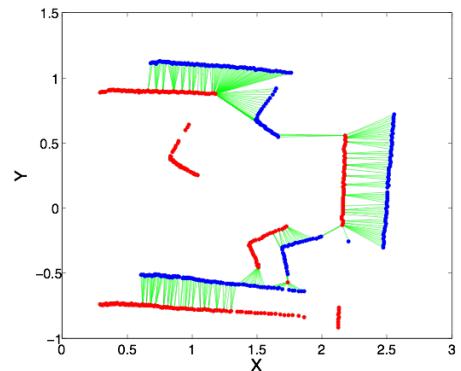


# Rejection

- Fix
- Zhang (*Zhang 94*)
- Mean (*Druon 06*)
- Median (*Diebel 04*)
- Trimmed (*Chetverikov 02*)

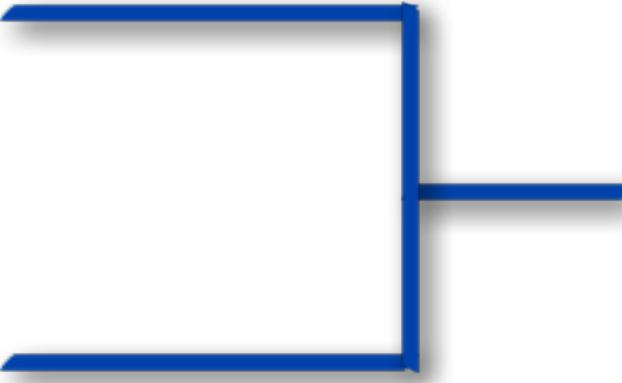


Assume unimodal distribution

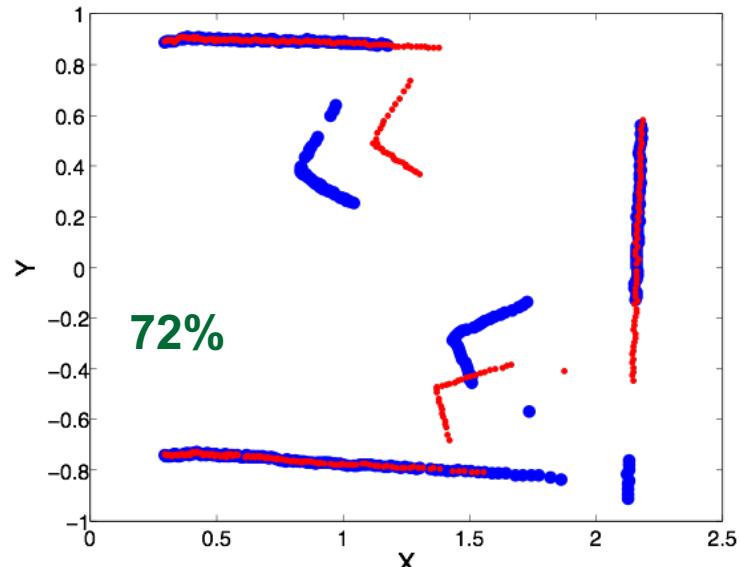
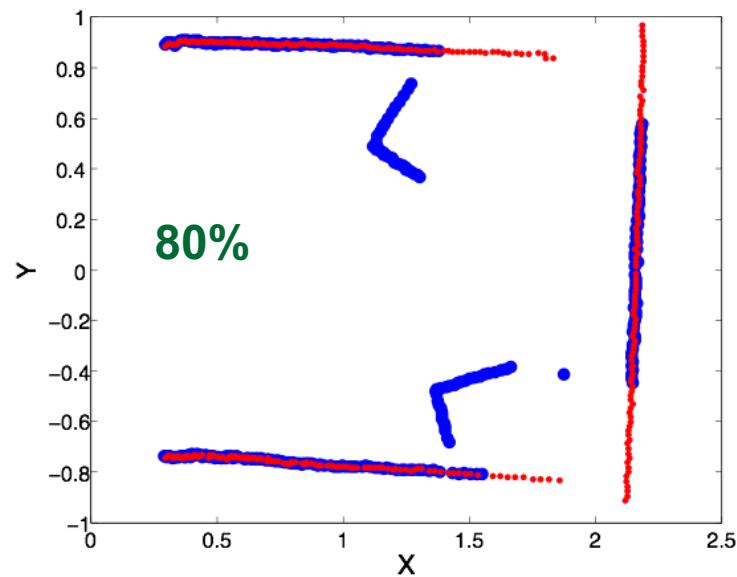


# Rejection

- Fix
- Zhang (*Zhang 94*)
- Mean (*Druon 06*)
- Median (*Diebel 04*)
- Trimmed (*Chetverikov 02*)



Assume few variation in overlapping ratio



# Rejection



- Fix
- Zhang (*Zhang 94*)
- Mean (*Druon 06*)
- Median (*Diebel 04*)
- Trimmed (*Chetverikov 02*)

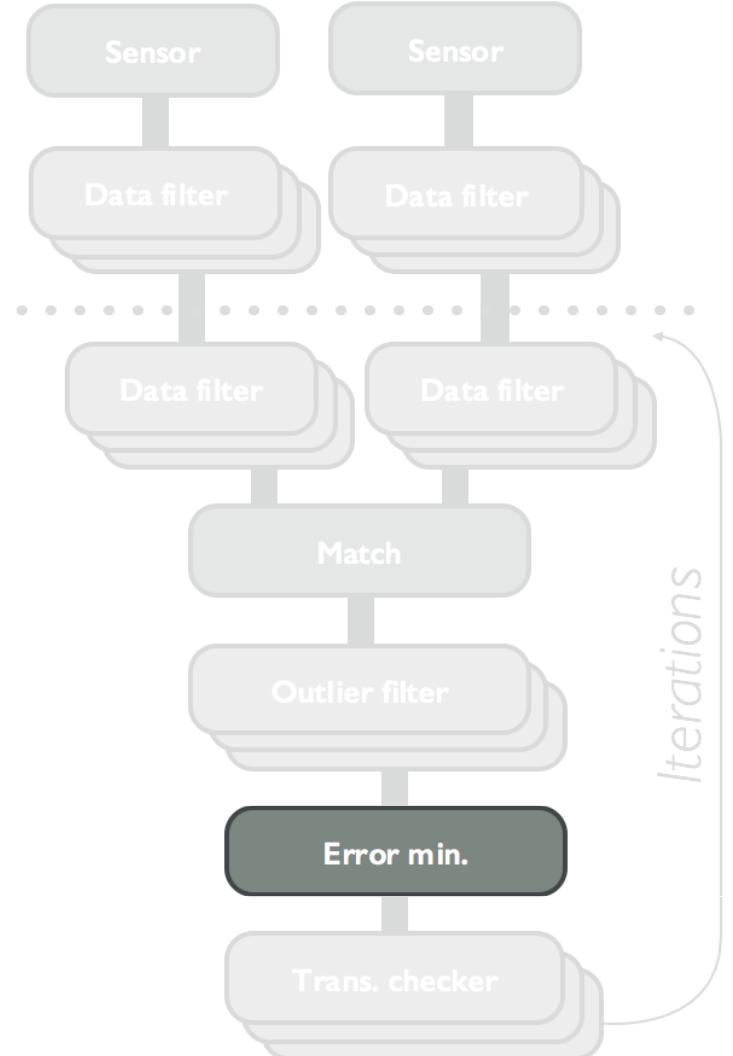


# Rejection



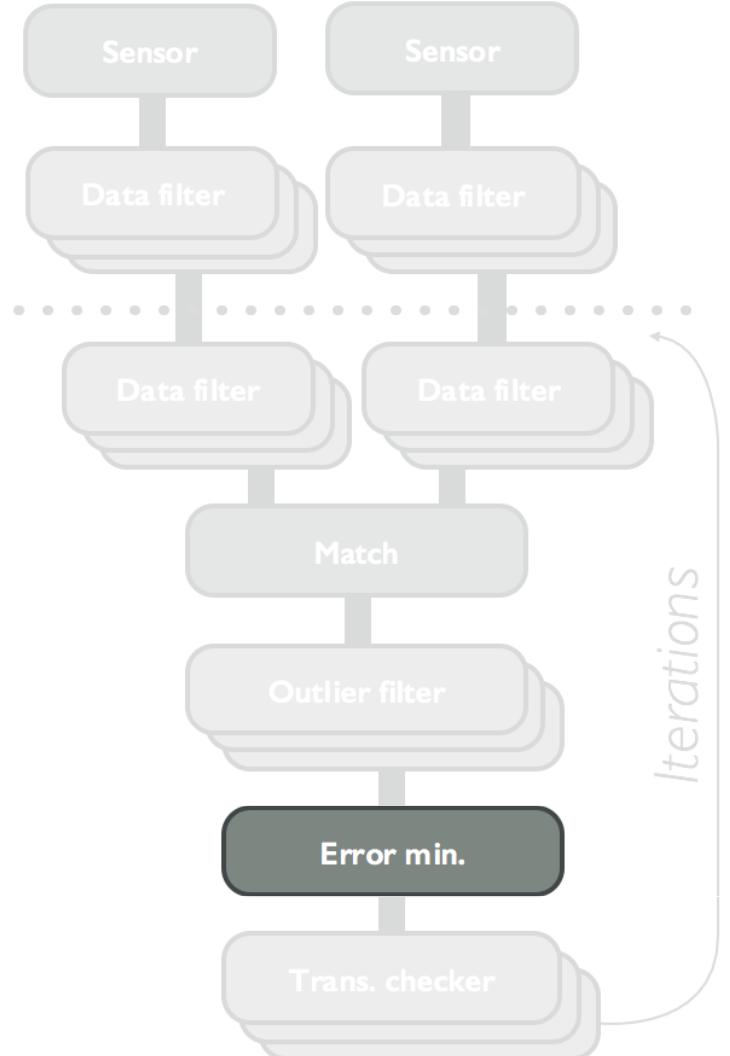
- Fix
- Zhang (*Zhang 94*)
- Mean (*Druon 06*)
- Median (*Diebel 04*)
- Trimmed (*Chetverikov 02*)
  
- Probabilistic (*Feldmar 96*) (*Philips 07*)
- RANSAC (*Arun 87*)

# Minimization



- Main goal
  - Rapid convergence
  - Well constrained

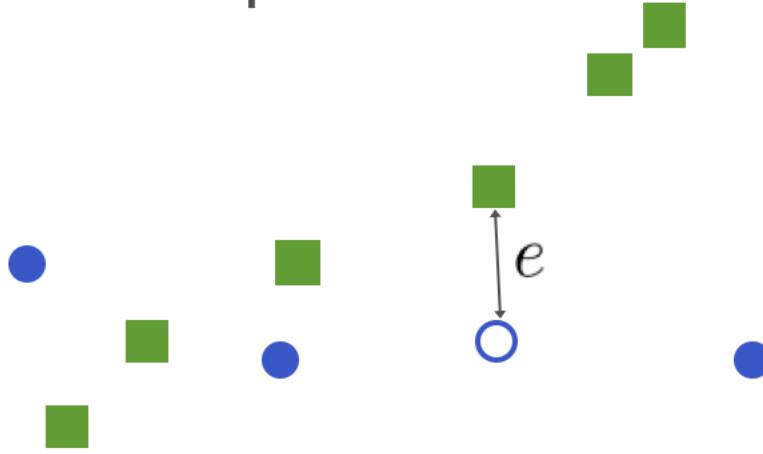
# Minimization



- Error metrics
  - Point-to-point
  - Point-to-plane
- and many more ...

# Minimization

Point-to-point



$$\sum_{i=0}^N \|\hat{T} * p_i - q_i\|^2$$

Besl and McKay [1992]

Walker et al. [1991] Horn [1987] Arun et al. [1987]

# Procrustes alignment



- Arun's method
  - Let  $\mathbf{p}_i$  be the points in scan 1 and  $\mathbf{q}_i$  be the points in scan 2
  - Compute the centers of both point clouds:

$$\bar{\mathbf{p}} = \frac{1}{n} \sum_{i=1}^n \mathbf{p}_i \quad \bar{\mathbf{q}} = \frac{1}{n} \sum_{i=1}^n \mathbf{q}_i$$

- Compute the matrix

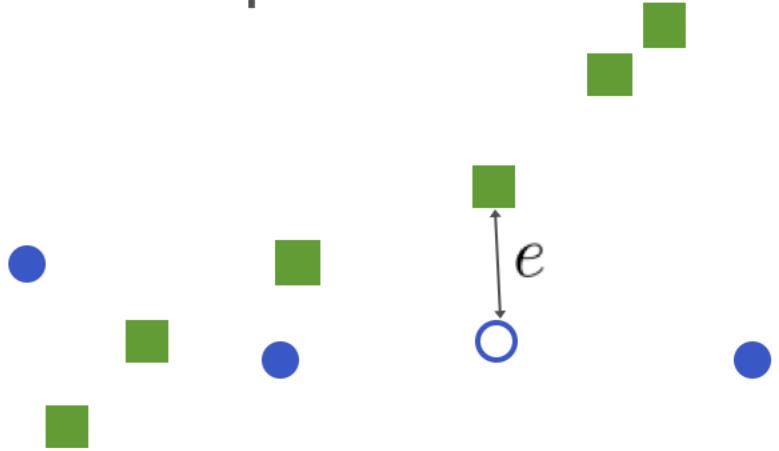
$$\mathbf{Q} = \sum_{i=1}^n (\mathbf{q}_i - \bar{\mathbf{q}})(\mathbf{p}_i - \bar{\mathbf{p}})^t$$

- Rotation  
Translation

$$\begin{aligned}\mathbf{R} &= \mathbf{V}\mathbf{U}^t \text{ where } SVD(\mathbf{Q}) = \mathbf{U}\Sigma\mathbf{V}^* \\ \mathbf{t} &= \bar{\mathbf{p}} - \mathbf{R}\bar{\mathbf{q}}\end{aligned}$$

# Minimization

Point-to-point

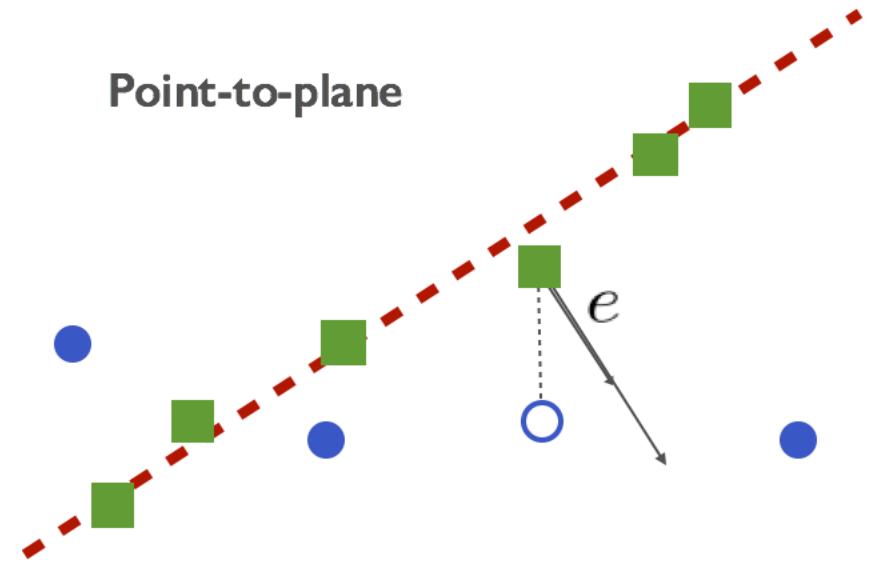


$$\sum_{i=0}^N \|\hat{T} * p_i - q_i\|^2$$

Besl and McKay [1992]

Walker et al. [1991] Horn [1987] Arun et al. [1987]

Point-to-plane

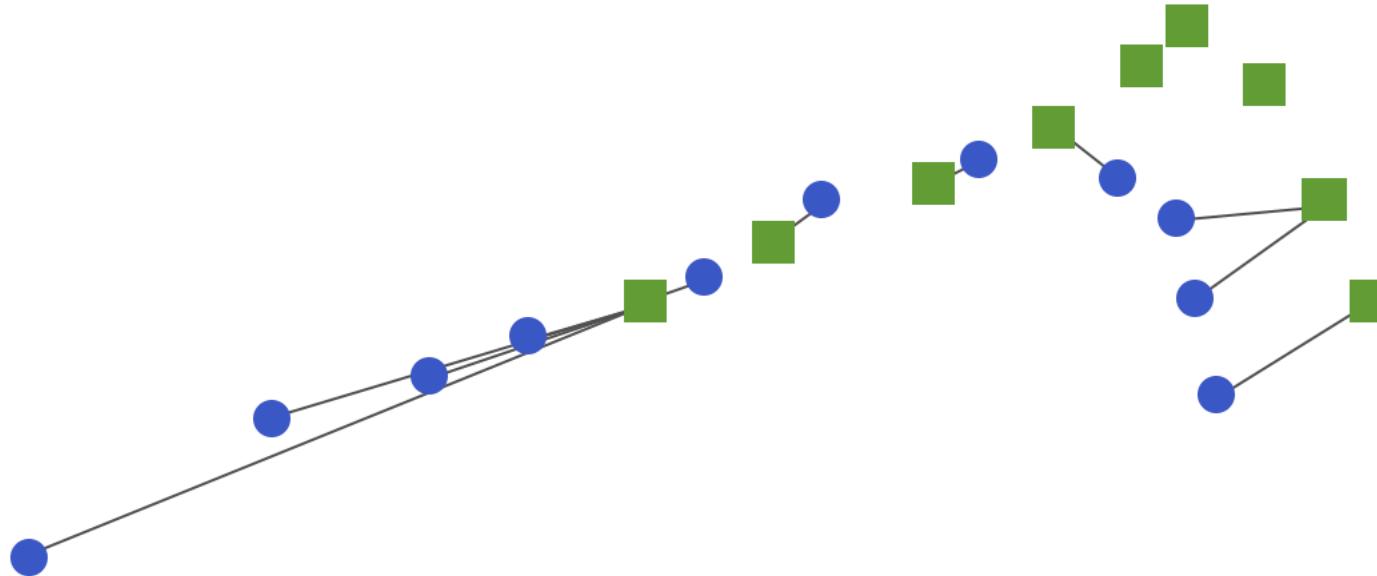


$$\sum_{i=0}^N \|(\hat{T} * p_i - q_i) \cdot n_i\|^2$$

Chen and Medioni [1992]

# Minimization

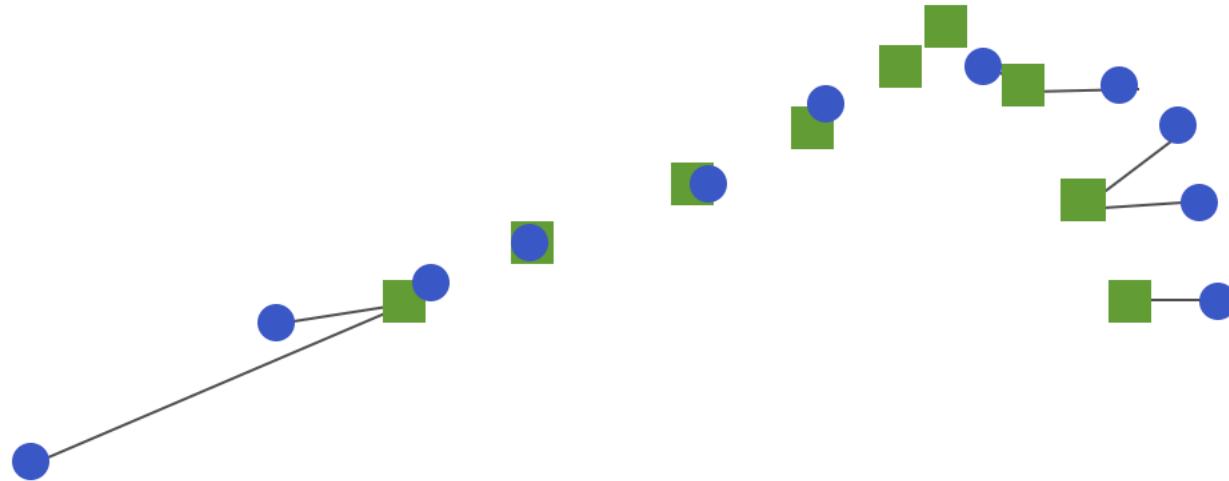
- Sliding situation



# Minimization

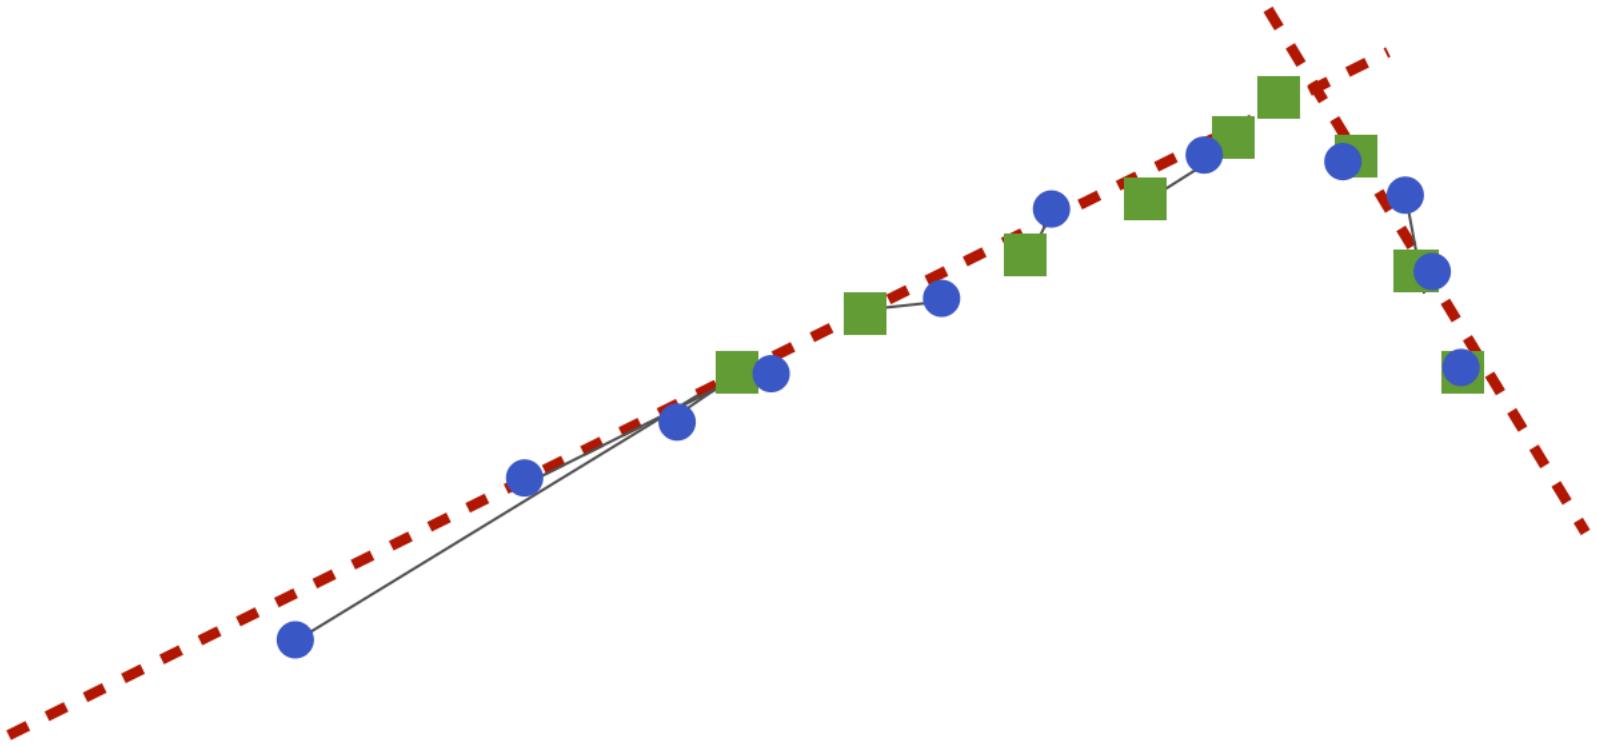


- Point-to-point distance

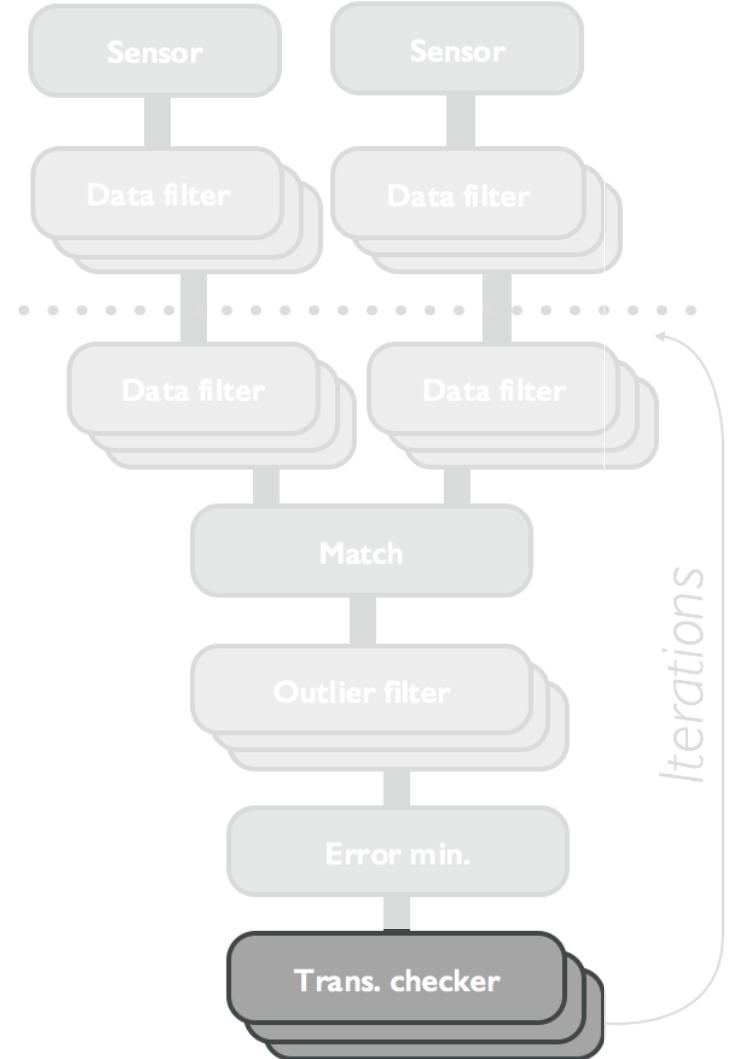


# Minimization

- Point-to-plane distance



# Convergence detection



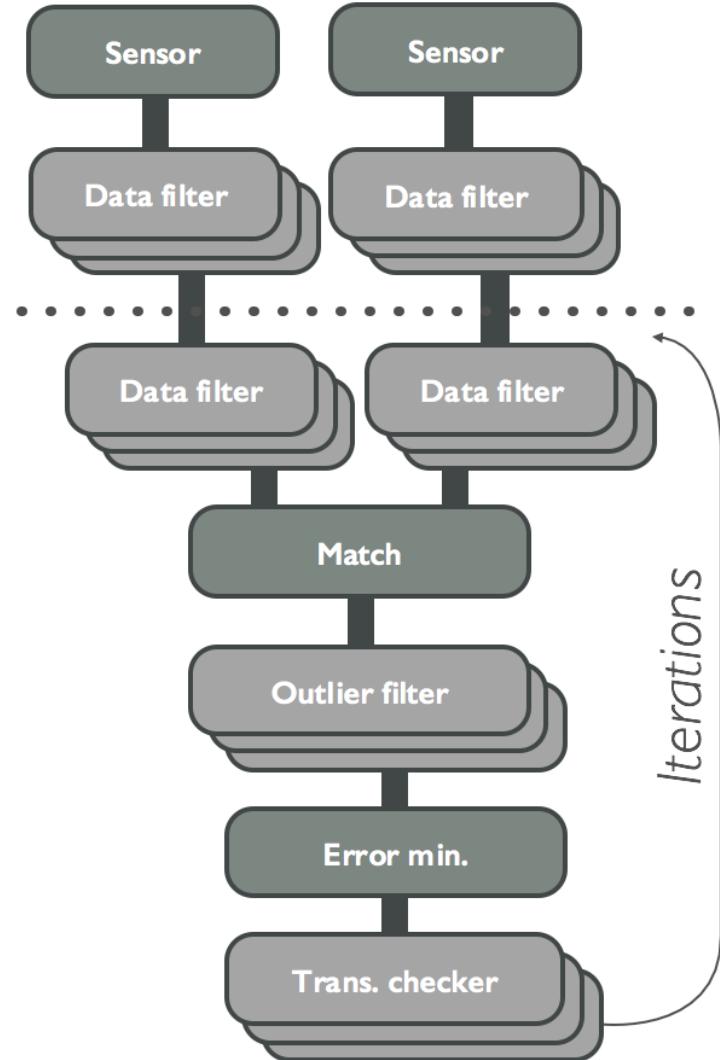
- Main goal:
- Iteration early out

# Convergence detection



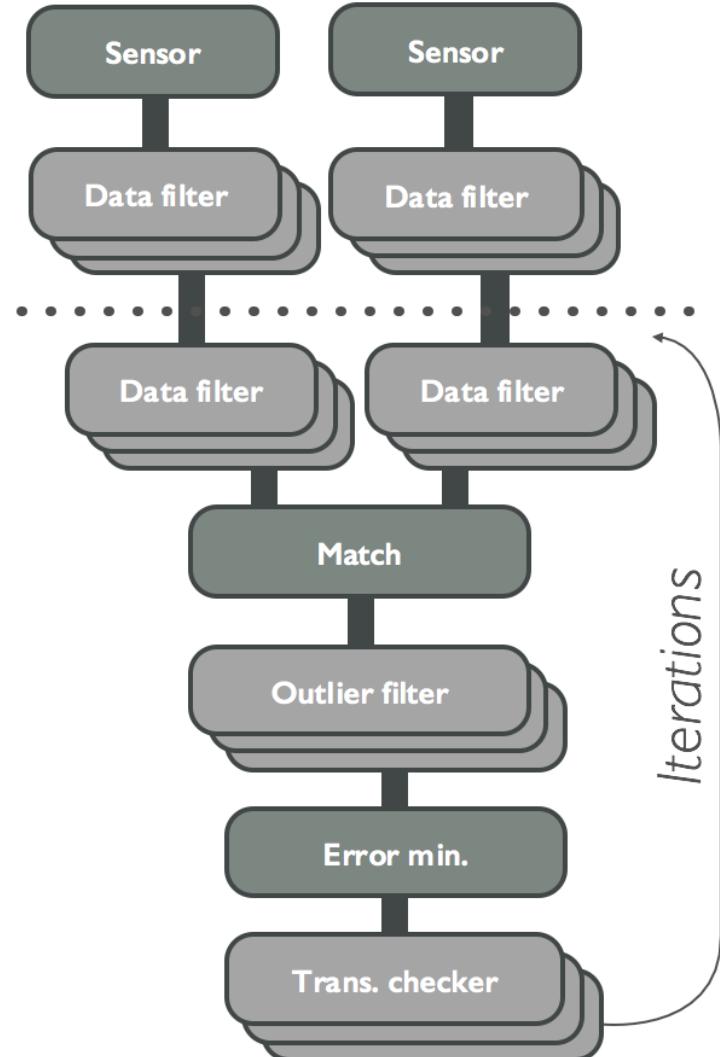
- Maximum number of iterations
  - Blind
- Absolute motion
  - Limit based on expected sensor noise
  - Detect divergence
  - Deals with rotation and translation separately
- Motion difference between iterations
  - Detect convergence up to wanted noise
  - Deals with rotation and translation separately

# Summary



- Accuracy vs. Speed!

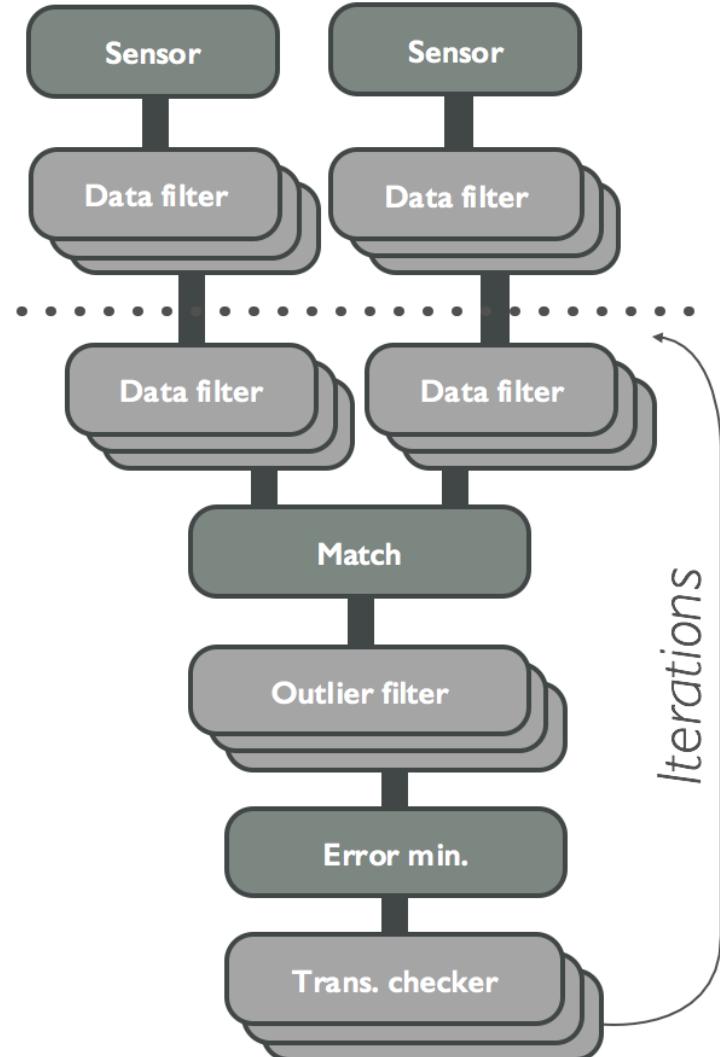
# Summary



## • Augment Accuracy

- Keep more points
- Augment the number of matches
- Adjust your expected overlap
- Verify that normal vectors are well extracted
- Add more iterations

# Summary



- **Augment Speed**

- Reduce number of points as soon as possible, compute only the minimum
- Approximative search (kD-tree)
- Point-to-plane takes less iterations
- Relax targeted values